




Probability and Recurrences

Admin

- **Assignment 8** is out due next Wed; based on this week of lectures
 - **Lots of new definitions**---helpful to review notes and examples from class before/while attempting HW
- Assignment 9 will be the last one you turn in 
- Assignment 10 will consists of practice problems
 - Help prepare you for finals; review all topics
- Find out about CS courses next year: **Preregistration Info Session at 3.15pm**
- **Health day** next week 

Recap: Last Lecture

- Expectation, law of total expectation and linearity of expectation
- Common tricks to compute expectation:
 - Break down into indicator RVs, or random variables we already know how to compute expectation of and use LoE
 - Condition random variable on an event and form a recurrence
- n Bernoulli trials when probability of success is p
 - Expectation of number of successes?
 - Expected number of trials until first success?
- n th Harmonic number?
 - $1 + 1/2 + 1/3 + \dots + 1/n = \Theta(\log n)$





Coupon/Pokemon Collector Problem



Gotta' Catch 'Em All


- Suppose there are n different types of Pokemon cards
- In each trial we purchase a pack that contains **one** Pokemon card
- We repeat until we have at least one of each type of card, how many packs does it take in expectation to collect all?
- Let X be the r.v. equal to the number of packs bought until you first have a card of each type.
- **Goal**: compute $E[X]$
- We break X into smaller random variables
- **Idea**: we make progress every time we get a card we don't already have




Pokemon Collector Problem

- Let X_i denote the "length of the i th phase", that is, the number of packs bought during the i th phase (i th phase ends as soon as we see the i th distinct card)

- Thus, $X = \sum_{i=1}^n X_i$





- Each phase can be thought of as flipping a biased coin until we see a head, where seeing a head = getting a new card



Pokemon Collector Problem

- $E[X_i]$ is the expected number of coin flips until success (expectation of a geometric r.v.)

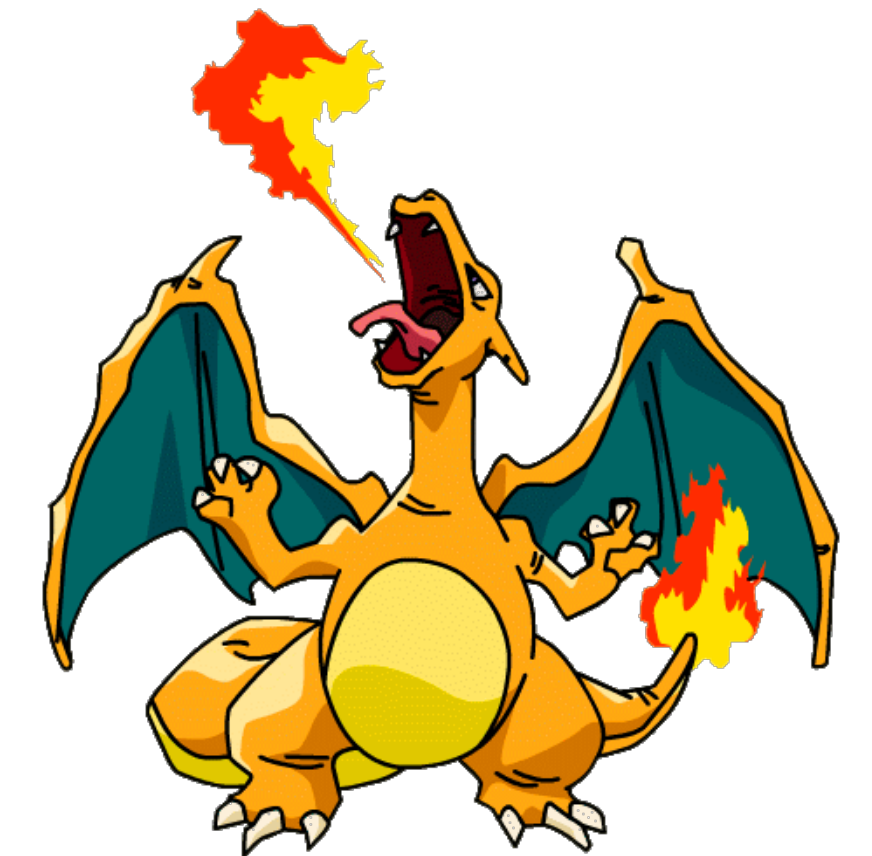
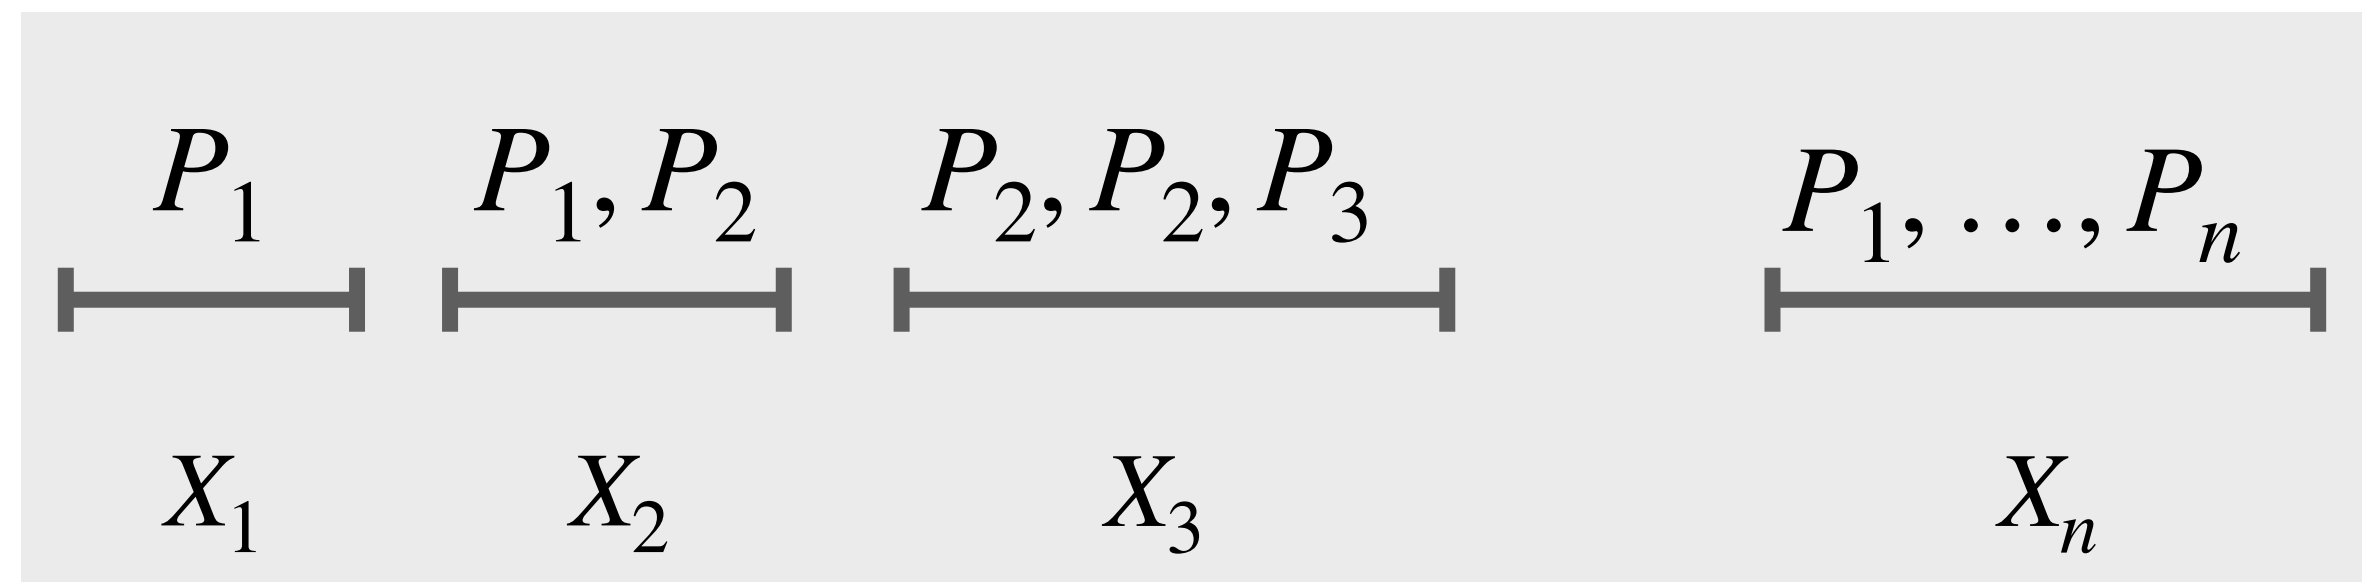


- We know, $E[X_i] = 1/p_i$ where p_i is the probability of success/ probability of seeing a heads during a coin flip in the i th phase
- Before the i th phase starts, we don't have $n - i + 1$ Pokemon
- Each of the n Pokemon are equally likely to be in a pack



Pokemon Collector Problem

- $E[X_i]$ is the expected number of coin flips until success
(expectation of a geometric r.v.)




- We know, $E[X_i] = 1/p_i$ where p_i is the probability of success/
probability of seeing a heads during a coin flip in the i th phase

- $$p_i = \frac{n - i + 1}{n}$$



Pokemon Collector Problem

- We know, $E[X_i] = 1/p_i$ where p_i is the probability of success/
probability of seeing a heads during a coin flip in the i th phase




- $E[X_i] = \text{Expected}[\text{number of flips until first heads}] = 1/p_i = \frac{n - i + 1}{n}$
- $E[X] = E\left[\sum_{i=1}^n X_i\right] = \sum_{i=1}^n E[X_i] = \sum_{i=1}^n \frac{n}{n - i + 1} = \sum_{i=1}^n \frac{n}{i} = nH_n = \Theta(n \log n)$

Random Walks and Recurrences


Random Walks

- A drunkard stumbles out of a bar. Each second, he either staggers 1 step to the left or staggers 1 step to the right, with equal probability. His home lies x steps to his left, and a canal lies y steps to his right.
- **Questions.** What is the probability that the drunkard arrives safely at home instead of falling into the canal? What is the expected duration of his journey, however it ends?
- The drunkard's meandering path is called a **random walk**
- Random walks are important as they model various phenomenon:
 - In Physics, random walks model gas diffusion
 - Google search engine uses random walks through the graph of web links to determine the relative importance of website
 - In finance theory, random walks can serve as a model for the fluctuation of market prices



Pass the Candy


- We have n students labelled $1, \dots, n$ and a professor labelled 0 sitting around in a circle
- Initially the professor has a candy bowl. She withdraws a piece of candy and then passes the bowl either to the left or right, with equal probability
- Each person who receives the bowl takes a piece of candy if they do not already have one; then passes it on randomly
- The ***last*** person to receive the candy wins the game
- Which player is most likely to win? Guess?
- Seems like 1 and n are almost always going to be eliminated right away. Seems like $n/2$ is most likely to win—but by how much?



Simpler Problem

- Suppose the players A, S_1, \dots, S_k, B are arranged in a line instead and S_1 initially has the candy
- As before, whenever a player gets the bowl they take a candy and pass it left or right with equal probability
- What is the probability that A gets the candy before B ?
- Let P_k be the probability that A gets the candy before B .
- **Base case.** Suppose $k = 1$, then $P_1 = 1/2$
- Suppose $k > 1$. In the first step there are two possibilities: the bowl either moves left to A or right to S_2

$$\begin{aligned} P_k &= \Pr(\text{first step is left}) \\ &\quad \cdot \Pr(A \text{ gets candy before } B \mid \text{first step is left}) \\ &+ \Pr(\text{first step is right}) \\ &\quad \cdot \Pr(A \text{ gets candy before } B \mid \text{first step is right}) \\ &= \frac{1}{2} \cdot 1 + \frac{1}{2} \cdot \Pr(A \text{ gets candy before } B \mid \text{first step is right}) \end{aligned}$$



Simpler Problem


- $P_k = \Pr(\text{first step is left})$
 - $\cdot \Pr(A \text{ gets candy before } B \mid \text{first step is left})$
 - $+ \Pr(\text{first step is right})$
 - $\cdot \Pr(A \text{ gets candy before } B \mid \text{first step is right})$
$$= \frac{1}{2} \cdot 1 + \frac{1}{2} \cdot \Pr(A \text{ gets candy before } B \mid \text{first step is right})$$

- Recurrence. $P_k = \frac{1}{2} \cdot 1 + \frac{1}{2} \cdot P_{k-1} \cdot P_k$ and $P_1 = \frac{1}{2}$ (Base case)

- Solve it using guess and check, and prove by induction

- $$P_2 = \frac{1}{2 - P_1} = \frac{2}{3}, \quad P_3 = \frac{1}{1 - P_2} = \frac{3}{4}$$

- $$P_k = \frac{k}{k+1} \quad (\text{Verify this is correct by induction})$$



New starting configuration


$$\Pr(S_1 \text{ gets candy before } B) = P_{k-1}$$

If S_1 gets candy, we are back in the initial configuration and A gets the candy before B with probability P_k .

Back to the Candy Game

- Consider player n on the right side of the professor. Only way n can win is the candy travels clockwise all the way around to player $n - 1$ before n ever touches it
- Thus, if we cut the circle and arrange on a line as shown, n wins if and only if $n - 1$ gets the bowl before n
- This fits our previous simpler problem model where we want to know the probability that B gets candy before A and $k = n - 1$
- $\Pr(n - 1 \text{ gets candy before } n) = 1 - \Pr(n \text{ gets candy before } n - 1)$

$$= 1 - \frac{n - 1}{(n - 1) + 1} = \frac{1}{n}$$
- Student n wins with probability $1/n$!
- We can extend this argument to show each student wins with probability $1/n$ --- **we would never have guessed this!**



Randomized Algorithm I

Min Cut

Randomized Min Cut

- **Global min-cut problem.** Given an undirected, unweighted graph $G = (V, E)$, find a cut (A, B) of minimum cardinality (that is, min # of edges crossing it).
- **Applications.** Network reliability, network design, circuit design, etc.
- **Poly-time network-flow solution** (by reduction to min s - t cut).
 - Replace every undirected edge (u, v) with $u \rightarrow v$ and $v \rightarrow u$, each of capacity 1
 - Fix any $s \in V$ and compute min s - t cut for every other node $t \in V - \{s\}$
 - $(n - 1)$ executions of min s - t cut
- Gives impression that finding global min cut is harder than finding a min s - t cut, which is not true
- Deceptively simple and efficient randomized algorithm [Karger 1992]

Karger's Min Cut

- **Edge contraction:** Contract edge e in G , denoted $G \leftarrow G/e$
 - Replace u and v by single new super-node w
 - Preserve edges, updating endpoints of u and v to w
 - Keep parallel edges, but delete self-loops
- An edge can be contracted in $O(n)$ time, assuming the graph is represented as an adjacency list



Karger's Min Cut

- Algorithm tries to *guess* the min cut by randomly contracting edges
- Running time:
 - $O(n)$ edge contraction, $O(n)$ iterations: $O(n^2)$
- **Correctness:** How often, if ever, does it return the min cut?


GUESSMINCUT(G):

for $i \leftarrow n$ downto 2

pick a random edge e in G

$G \leftarrow G/e$

return the only cut in G




Preserving Cuts

- **Observation.** Any cut in the contracted graph is a cut in the original graph
- Let $C = (S, V - S)$ be any cut, if algorithm never contracts an edge crossing this cut, then it will produce the cut C
- Let C be any arbitrary min cut of cardinality k
- If we pick an edge in G uniformly at random, what is the probability of picking an edge in C ?

- $\frac{k}{m}$


- We want to upper bound this probability
- Can we lower bound m in terms of k ?



Preserving Cuts

- The minimum cut C in G has cardinality k
- What can we say about degree of each vertex in G ?
 - Must be at least k
- G must have at least $nk/2$ edges
- If we pick an edge in G uniformly at random, what is the probability of picking an edge in C ?

$$\bullet \quad \frac{k}{m} \leq \frac{2k}{nk} = \frac{2}{n}$$



Karger's Analysis

- Let C be any arbitrary min cut of cardinality k
 - $\Pr(\text{picking an edge in } C) = \frac{k}{m} \leq \frac{k}{nk/2} = \frac{2}{n}$
- Probability we don't contract a cut edge in the 1st step $\geq 1 - \frac{2}{n}$
- After the first edge is contracted, the algorithm proceeds recursively (with independent random choices) on the $(n - 1)$ -vertex graph
- Let $P(n)$ denote the probability that the algorithm returns the correct min cut on an n -vertex graph, then

Karger's Analysis

- Let $P(n)$ denote the probability that the algorithm returns the correct min cut on an n -vertex graph, then

- $P(n) \geq \left(1 - \frac{2}{n}\right) \cdot P(n-1)$, with base case $P(2) = 1$

- Expanding the recurrence:

$$P(n) \geq \frac{n-2}{n} \cdot \frac{n-3}{n-1} \cdot \frac{n-4}{n-2} \cdot \frac{n-5}{n-3} \cdots \frac{4}{6} \cdot \frac{3}{5} \cdot \frac{2}{4} \cdot \frac{1}{3}$$

- Terms cancel out to get: $P(n) \geq \frac{2}{n(n-1)} = 1/\binom{n}{2}$

Amplifying Success Probability

- Thus, a single execution of Karger's min cut algorithm finds the min cut with probability at least $1/\binom{n}{2}$, which is low
 - But, we can amplify our success probability!
- Run the algorithm R times (using independent random choices) and pick the best min-cut among them
- What is probability we don't find the min cut after R repetitions?

$$\bullet \left(1 - 1/\binom{n}{2}\right)^R$$

Amplifying Success Probability

- If we execute $R = \binom{n}{2}$ times, the probability of failure is

- $\left(1 - 1/\binom{n}{2}\right)^{\binom{n}{2}}$: how can we simplify this?

- $\leq \frac{1}{e}$

- If we set $R = \binom{n}{2} c \ln n$, the failure probability becomes polynomially

small in n : $\left(\frac{1}{e}\right)^{c \ln n} = \frac{1}{n^c}$

Important Inequality:

$$(1 - x) \leq \left(\frac{1}{e}\right)^x \text{ for } x \geq 1$$


With High Probability

- If we run the algorithm $R = \binom{n}{2} c \ln n$ times, we can make the failure probability polynomially small in n : $\left(\frac{1}{e}\right)^{c \ln n} = \frac{1}{n^c}$
- Karger's algorithm finds the min-cut **with high probability (w.h.p.)**


An algorithm is correct **with high probability (w.h.p.)** with respect to input size n if it fails with probability at most $\frac{1}{n^c}$ for any constant $c > 1$.

Example Execution


trial 1




trial 2




trial 3




trial 4



trial 5
(finds min cut)



trial 6



...

Reference: Thore Husfeldt

Karger's Running Time

- Thus, Karger's algorithm finds the min-cut with high probability (w.h.p.)
- Running time: we perform $\Theta(n^2 \log n)$ iterations, each $O(n^2)$ time
 - $O(n^4 \log n)$ time
 - Faster than naive-flow-techniques, nothing to get excited about
- Improves to $O(n^2 \log^3 n)$ by guessing cleverly! [Karger-Stein 1996]
- **Idea:** Improve the guessing algorithm using the observation:
 - As the graph shrinks, the probability of contracting an edge in the minimum cut increases
 - At first the probability is very small: $2/n$ but by the time there are three nodes, we have a $2/3$ chance of screwing up!

Takeaways

- Karger's algorithm is an example of a "**Monte Carlo**" randomized algorithm
 - Find the correct answer most of the time
- You can increase the success rate of algorithms with one-sided errors by iterating it multiple times and taking the best solution
 - If the probability of success is $1/f(n)$, then running it $O(f(n)\log n)$ times gives a high probability of success
- If you're more intelligent about how you iterate the algorithm, you can often do much better than this
- Next, we'll see an example of a "**Las Vegas**" algorithm
 - Randomized selection and quick sort

Acknowledgments

- Some of the material in these slides are taken from
 - Kleinberg Tardos Slides by Kevin Wayne (<https://www.cs.princeton.edu/~wayne/kleinberg-tardos/pdf/04GreedyAlgorithmsI.pdf>)
 - Jeff Erickson's Algorithms Book (<http://jeffe.cs.illinois.edu/teaching/algorithms/book/Algorithms-JeffE.pdf>)