

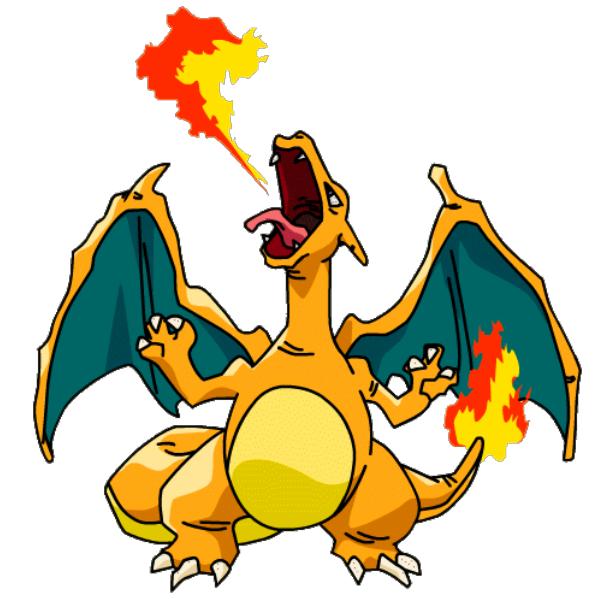
Probability and Recurrences

Admin

- **Assignment 8** is out due next Wed
 - Based on this week of lectures
 - **Lots of new definitions**--helpful to review notes and examples from class before/while attempting HW
- Assignment 9 will be the last one you turn in 
- Assignment 10 will consist of practice problems
 - Help prepare you for finals; review all topics
- **Health day** next week 

Recap: Last Lecture

- Expectation and law of total expectation
- Linearity of expectation
- Common tricks to compute expectation:
 - Break down into indicator random variables, or random variables we already know how to compute expectation of and use LoE
 - Condition random variable on an event and form a recurrence
- n Bernoulli trials when probability of success is p
 - Expectation of number of successes?
 - Expected number of trials until first success?
- n th Harmonic number?

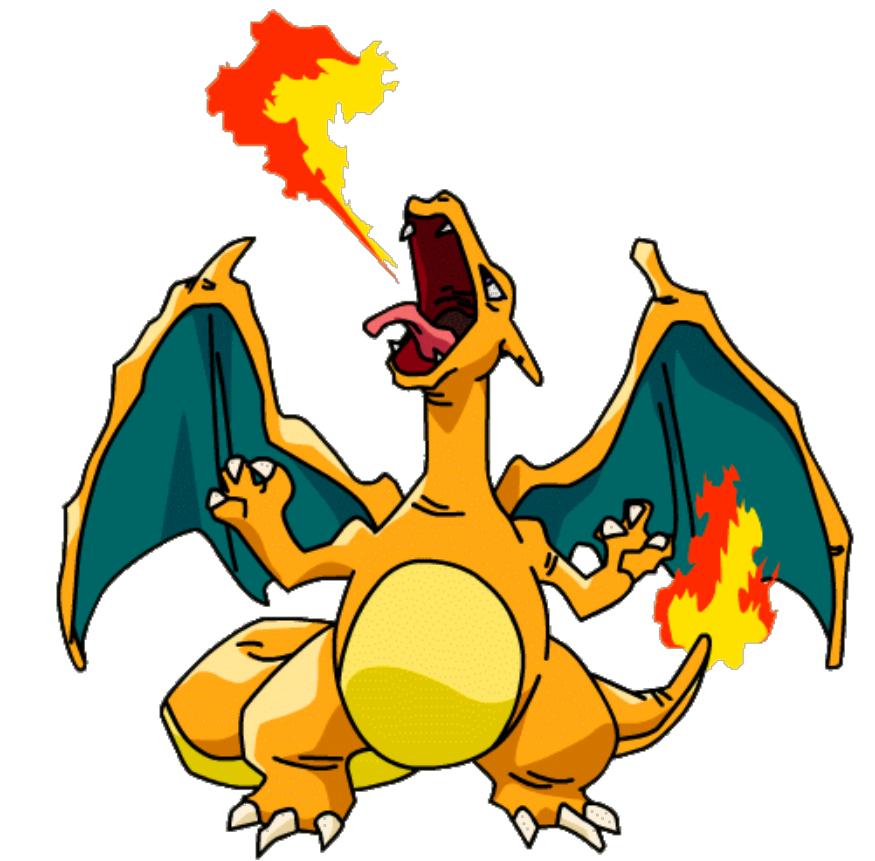


Coupon/Pokemon Collector Problem



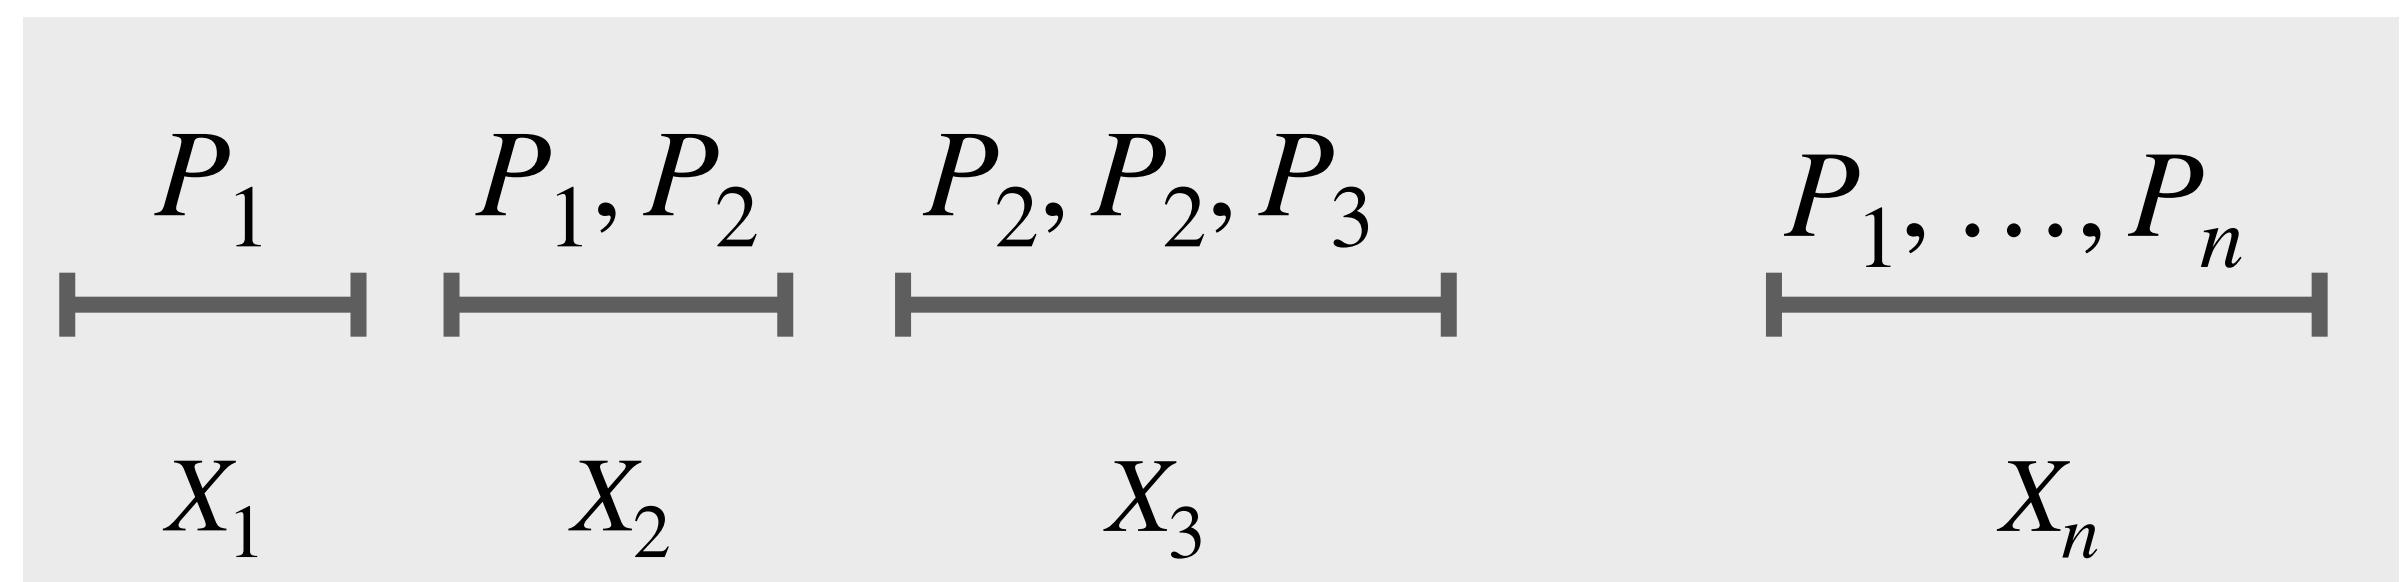
Gotta' Catch 'Em All

- Suppose there are n different types of Pokemon cards
- In each trial we purchase a pack that contains a Pokemon card
- We repeat until we have at least one of each type of card, how many packs does it take in expectation to collect all?
- Let X be the r.v. equal to the number of packs bought until you first have a card of each type. **Goal:** compute $E[X]$
- We break X into smaller random variables
- **Idea:** we make progress every time we get a card we don't already have

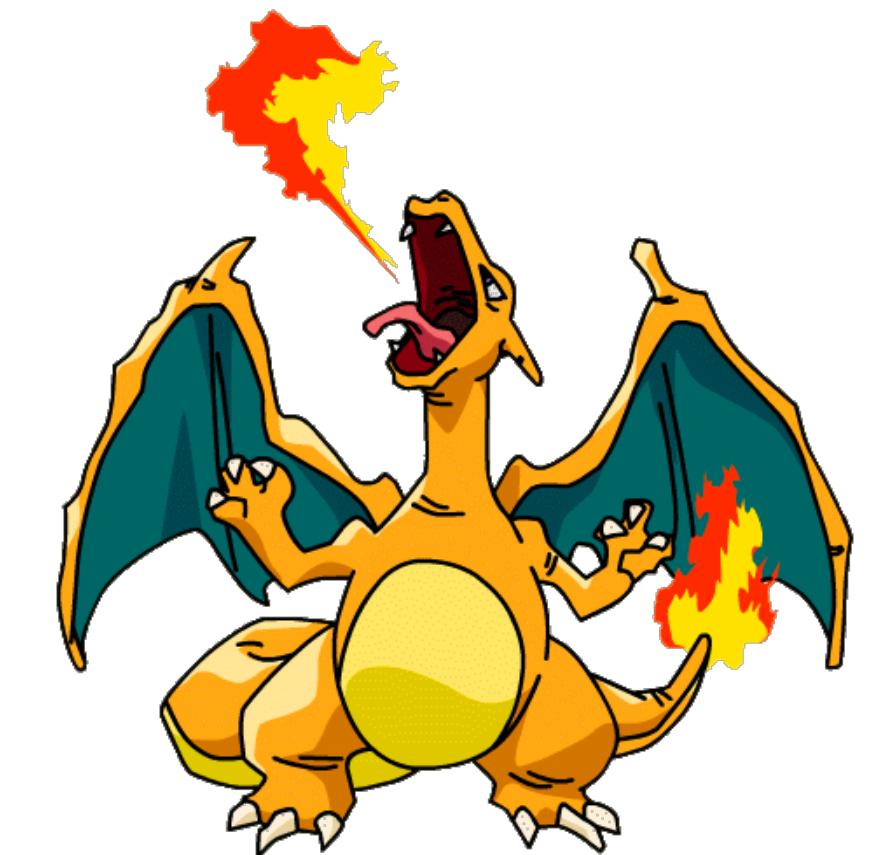


Pokemon Collector Problem

- Let X_i denote the "length of the i th phase", that is, the number of packs bought during the i th phase (i th phase ends as soon as we see the i th distinct card)
- Thus, $X = \sum_{i=1}^n X_i$

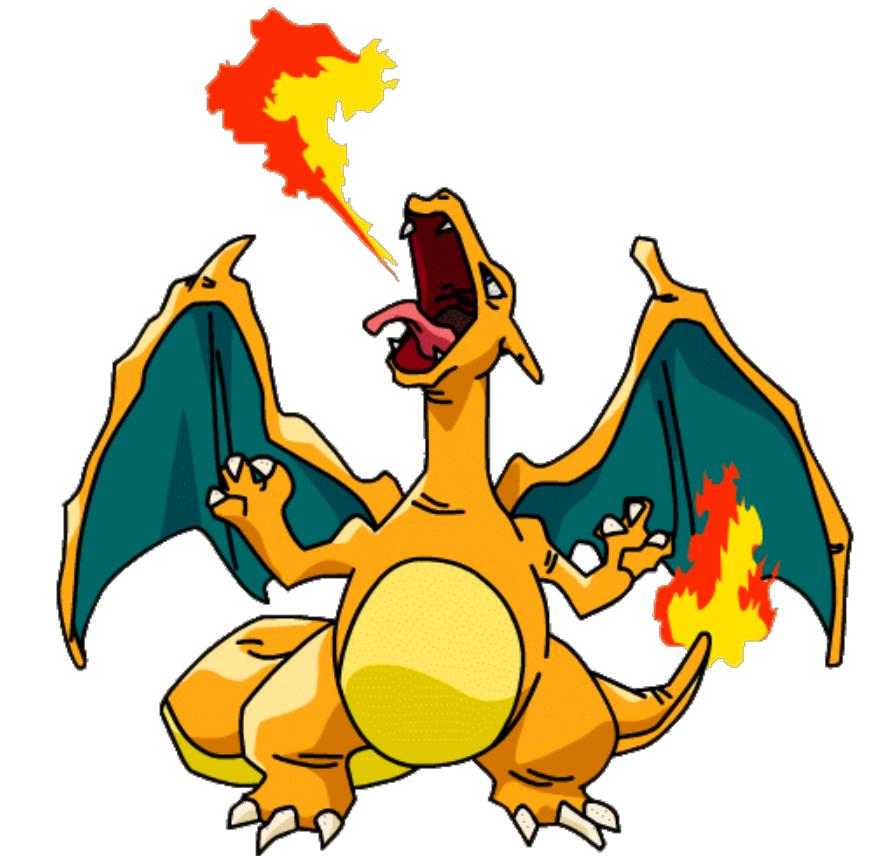
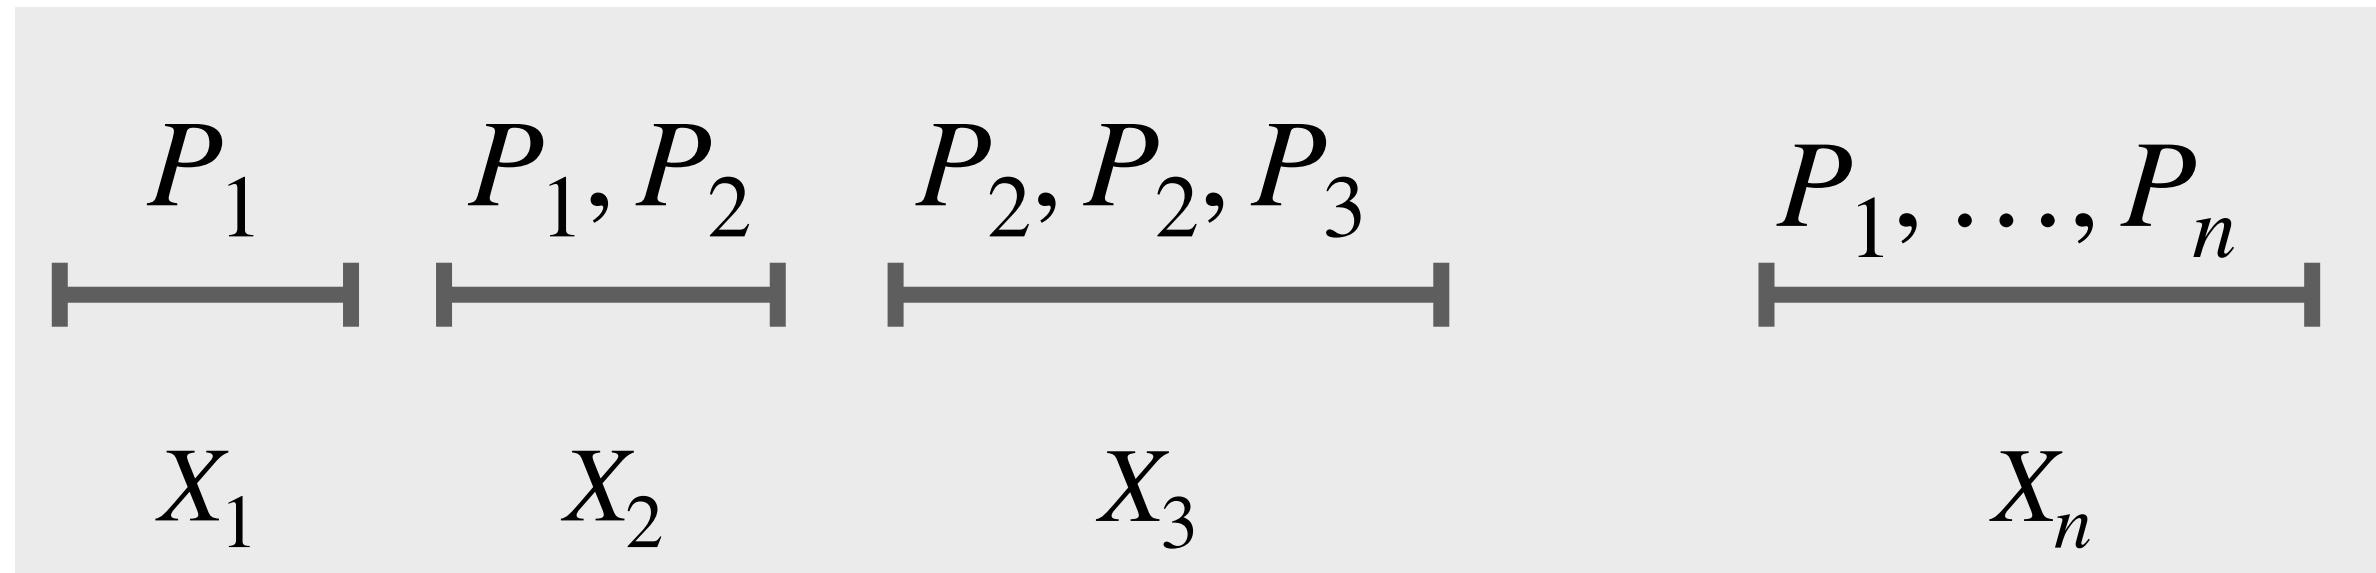


- Each phase can be thought of as flipping a biased coin until we see a head, where seeing a head = getting a new card



Pokemon Collector Problem

- $E[X_i]$ is the expected number of coin flips until success (expectation of a geometric r.v.)

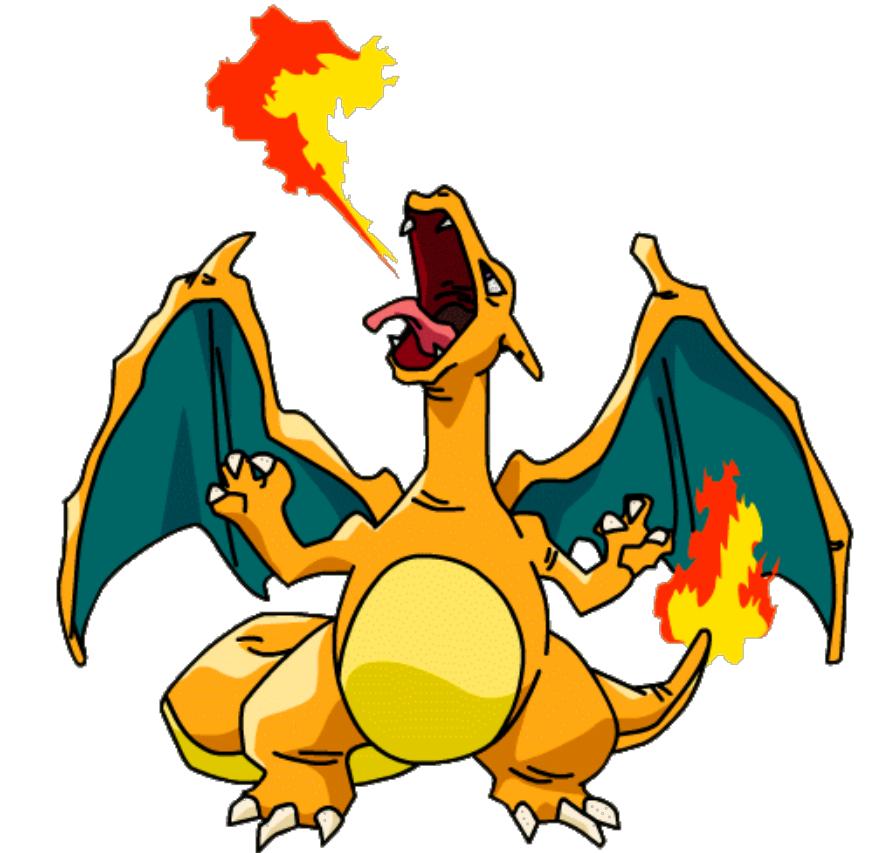
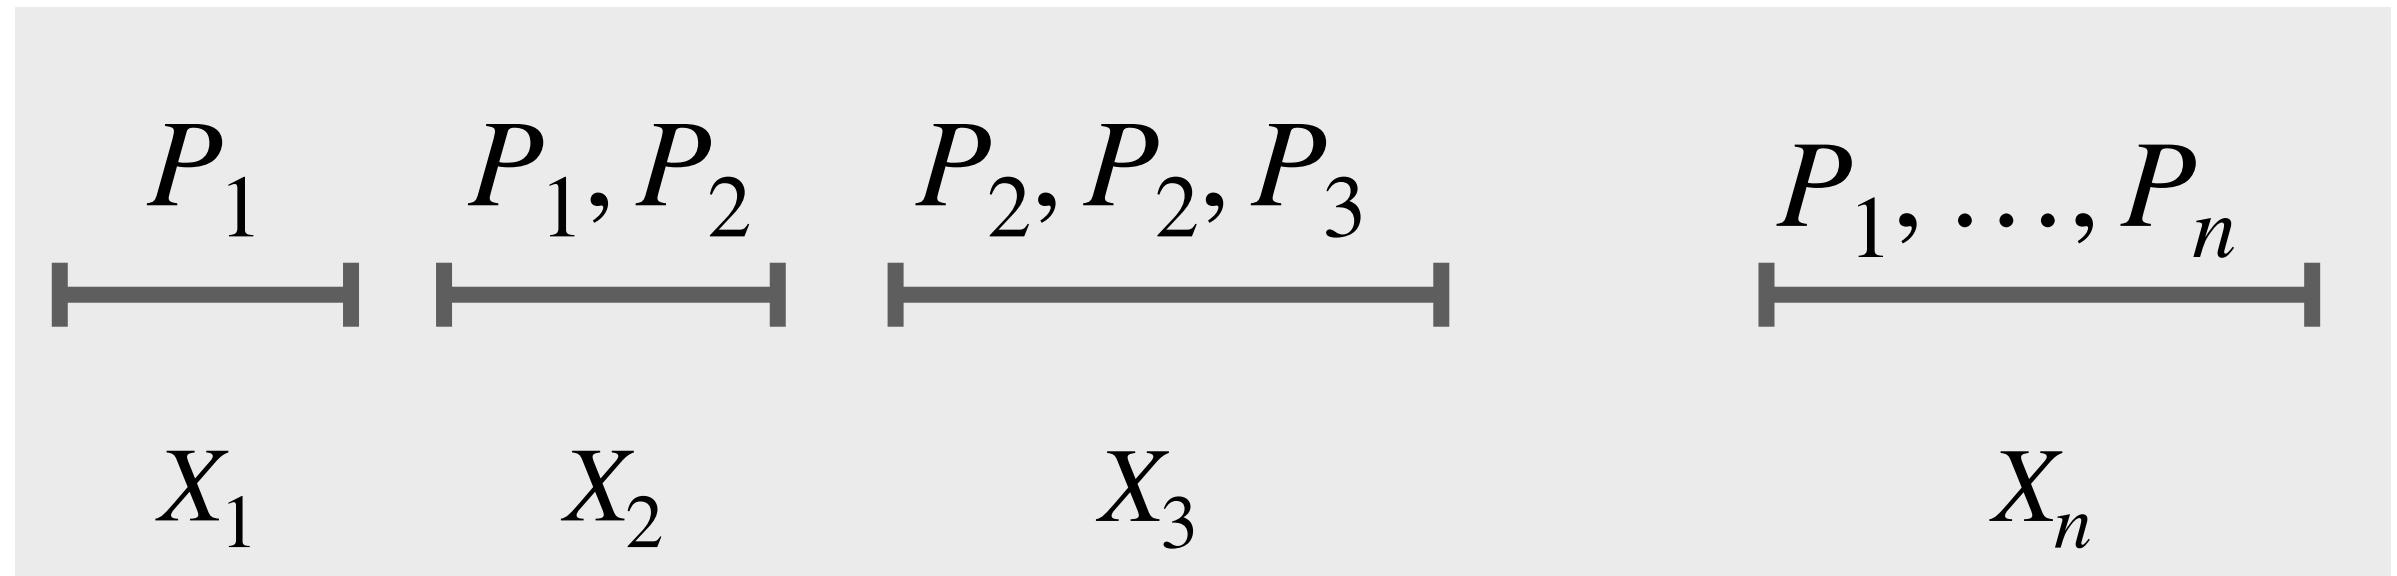


- We know, $E[X_i] = 1/p_i$ where p_i is the probability of success/ probability of seeing a heads during a coin flip in the i th phase
- Before the i th phase starts, we don't have $n - i + 1$ Pokemon
- Each of the n Pokemon are equally likely to be in a pack



Pokemon Collector Problem

- $E[X_i]$ is the expected number of coin flips until success
(expectation of a geometric r.v.)



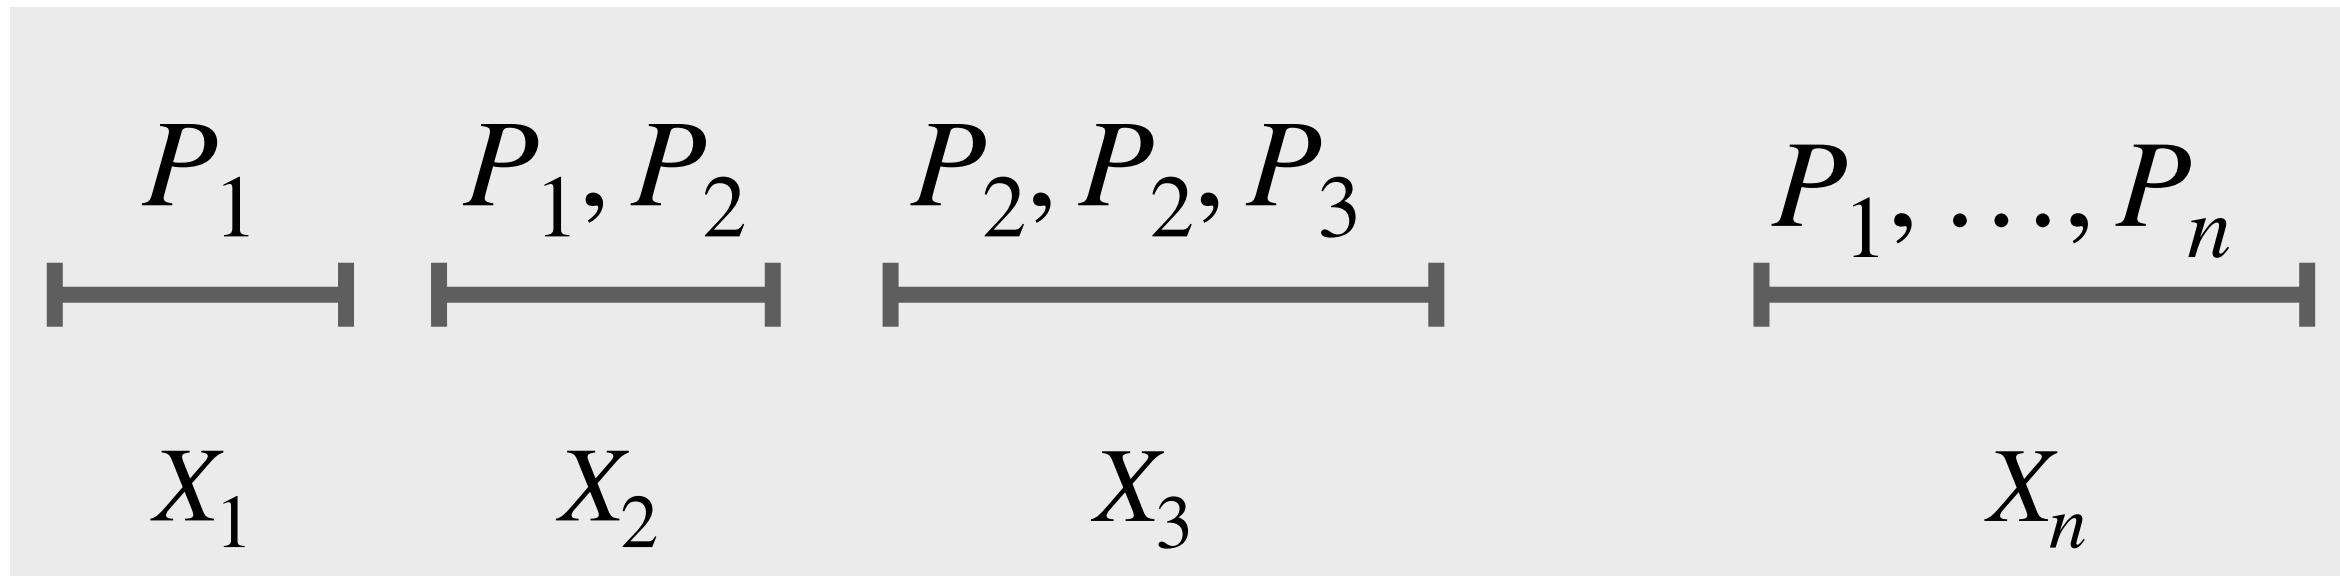
- We know, $E[X_i] = 1/p_i$ where p_i is the probability of success/probability of seeing a heads during a coin flip in the i th phase

$$\bullet p_i = \frac{n - i + 1}{n}$$



Pokemon Collector Problem

- We know, $E[X_i] = 1/p_i$ where p_i is the probability of success/probability of seeing a heads during a coin flip in the i th phase

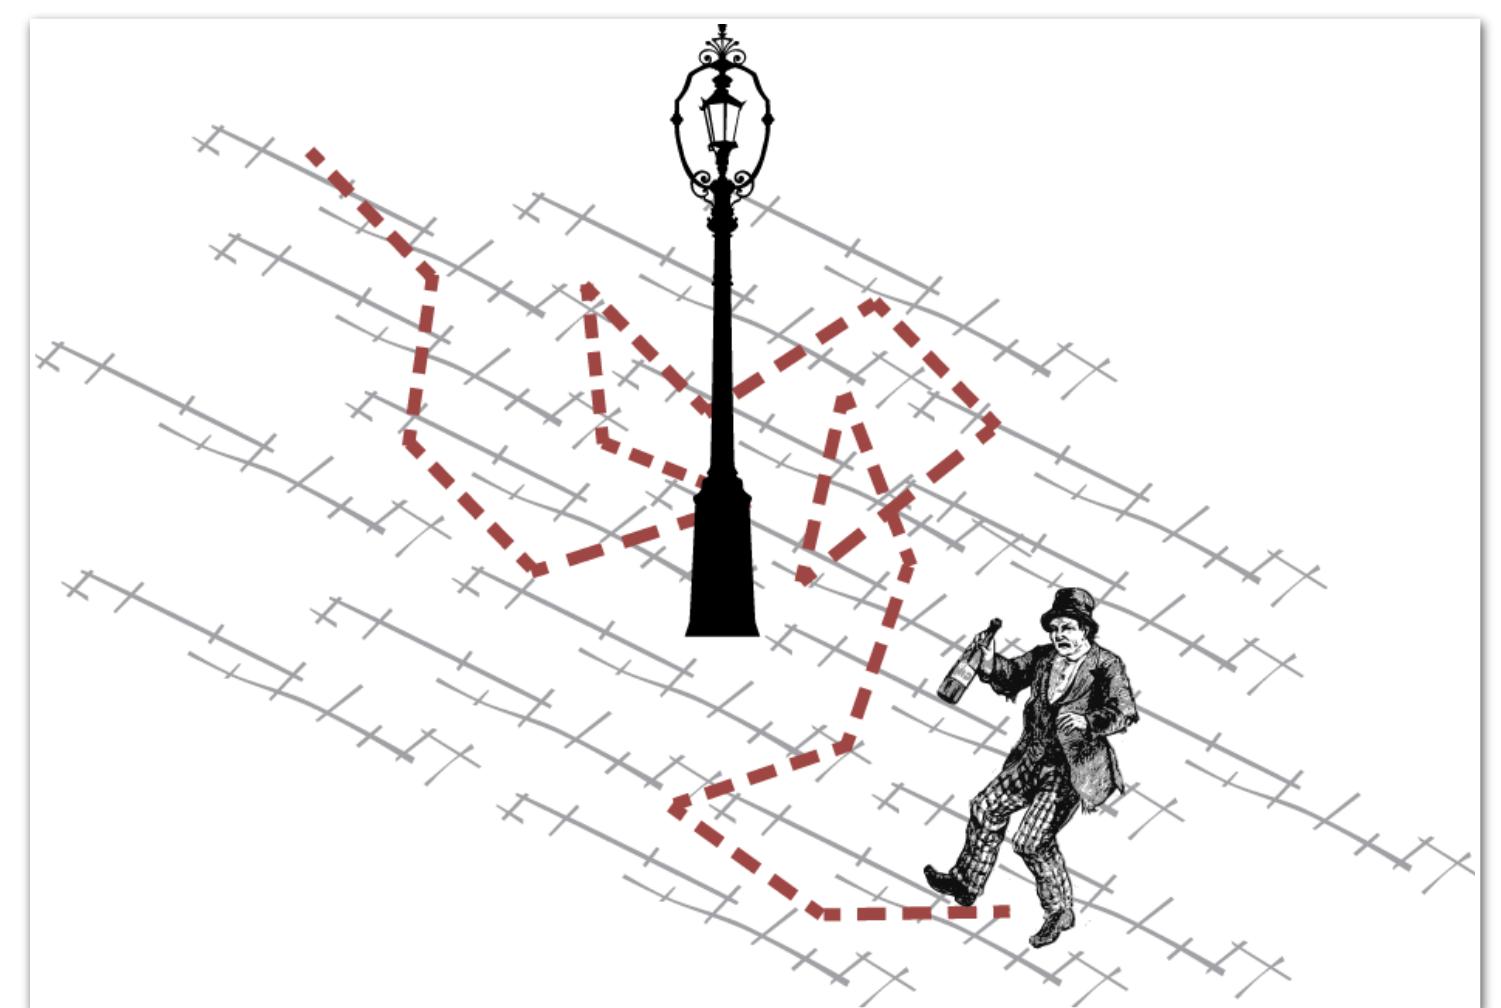


- $E[X_i] = \text{Expected}[\text{number of flips until first heads}] = 1/p_i = \frac{n - i + 1}{n}$
- $E[X] = E[\sum_{i=1}^n X_i] = \sum_{i=1}^n E[X_i] = \sum_{i=1}^n \frac{n}{n - i + 1} = \sum_{i=1}^n \frac{n}{i} = nH_n = \Theta(n \log n)$

Random Walks and Recurrences

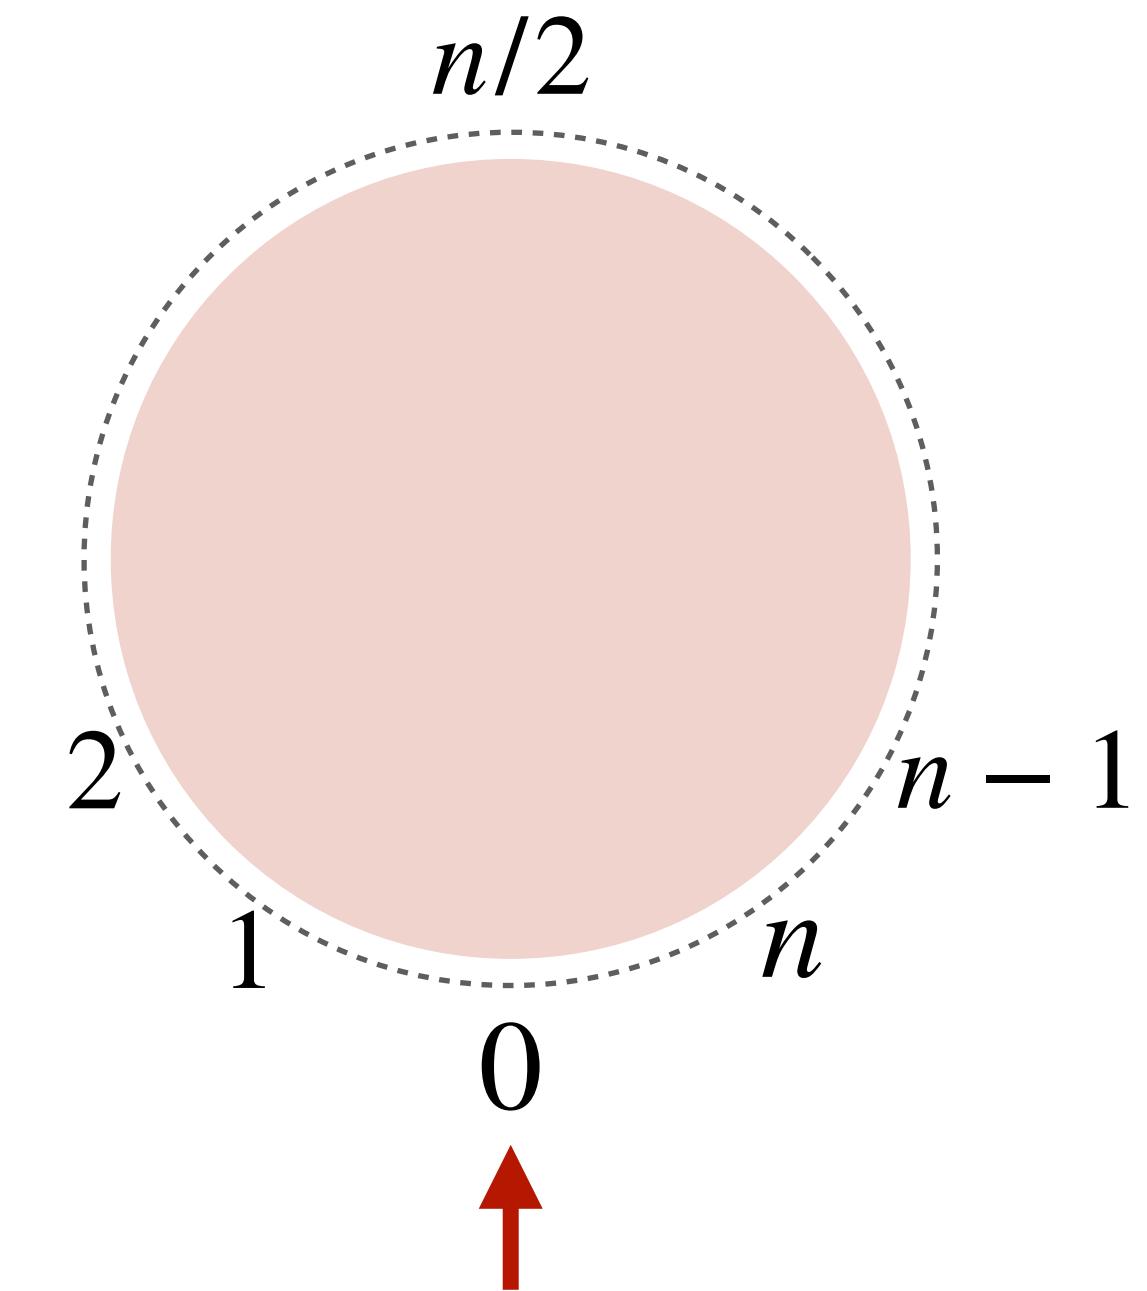
Random Walks

- A drunkard stumbles out of a bar. Each second, he either staggers 1 step to the left or staggers 1 step to the right, with equal probability. His home lies x steps to his left, and a canal lies y steps to his right.
- **Questions.** What is the probability that the drunkard arrives safely at home instead of falling into the canal? What is the expected duration of his journey, however it ends?
- The drunkard's meandering path is called a **random walk**
- Random walks are important as they model various phenomenon:
 - In Physics, random walks model gas diffusion
 - Google search engine uses random walks through the graph of web links to determine the relative importance of website
 - In finance theory, random walks can serve as a model for the fluctuation of market prices



Pass the Candy

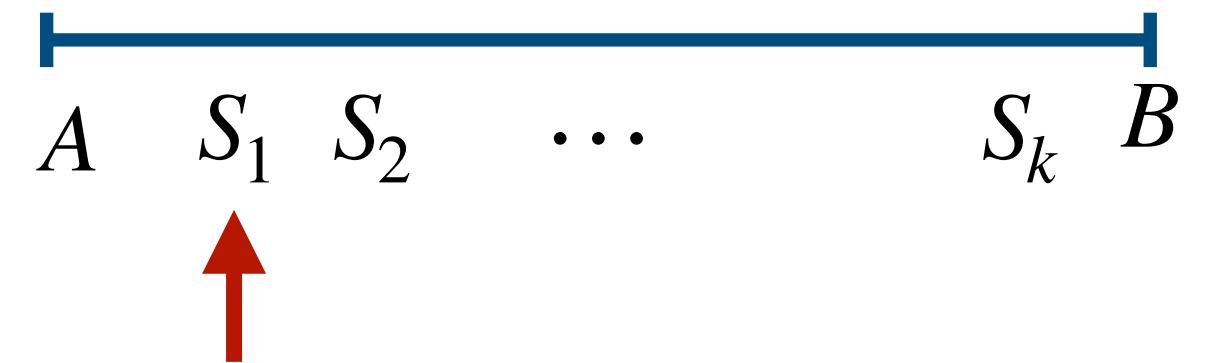
- We have n students labelled $1, \dots, n$ and a professor labelled 0 sitting around in a circle
- Initially the professor has a candy bowl. She withdraws a piece of candy and then passes the bowl either to the left or right, with equal probability
- Each person who receives the bowl takes a piece of candy if they do not already have one; then passes it on randomly
- The **last** person to receive the candy wins the game
- Which player is most likely to win? Guess?
- Seems like 1 and n are almost always going to be eliminated right away. Seems like $n/2$ is most likely to win—but by how much?



Simpler Problem

- Suppose the players A, S_1, \dots, S_k, B are arranged in a line instead and S_1 initially has a the candy
- As before, whenever a player gets the bowl they take a candy and pass it left or right with equal probability
- What is the probability that A gets the candy before B ?
- Let P_k be the probability that A gets the candy before B .
- **Base case.** Suppose $k = 1$, then $P_1 = 1/2$
- Suppose $k > 1$. In the first step there are two possibilities: the bowl either moves left to A or right to S_2

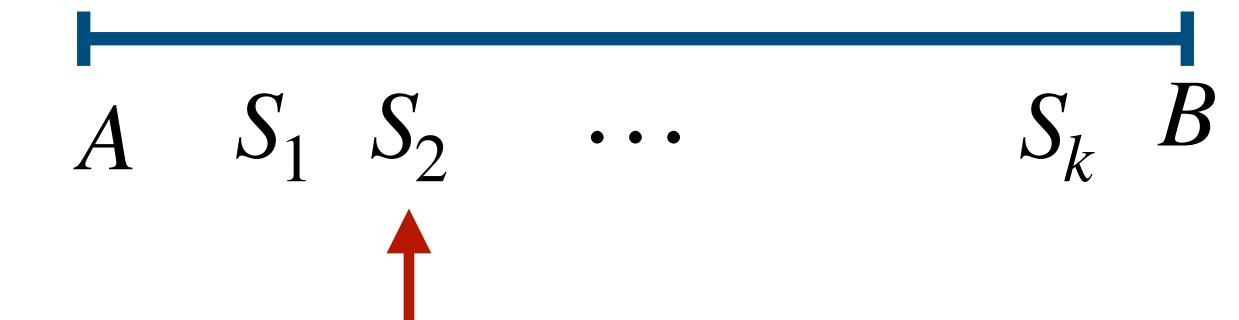
$$\begin{aligned} P_k &= \Pr(\text{first step is left}) \\ &\quad \cdot \Pr(A \text{ gets candy before } B \mid \text{first step is left}) \\ &\quad + \Pr(\text{first step is right}) \\ &\quad \cdot \Pr(A \text{ gets candy before } B \mid \text{first step is right}) \\ &= \frac{1}{2} \cdot 1 + \frac{1}{2} \cdot \Pr(A \text{ gets candy before } B \mid \text{first step is right}) \end{aligned}$$



Simpler Problem

- $P_k = \Pr(\text{first step is left})$
 - $\cdot \Pr(A \text{ gets candy before } B \mid \text{first step is left})$
 - $+ \Pr(\text{first step is right})$
 - $\cdot \Pr(A \text{ gets candy before } B \mid \text{first step is right})$
- $= \frac{1}{2} \cdot 1 + \frac{1}{2} \cdot \boxed{\Pr(A \text{ gets candy before } B \mid \text{first step is right})}$

- Recurrence. $P_k = \frac{1}{2} \cdot 1 + \frac{1}{2} \cdot P_{k-1} \cdot P_k$ and $P_1 = \frac{1}{2}$ (Base case)
- Solve it using guess and check, and prove by induction
- $P_2 = \frac{1}{2 - P_1} = \frac{2}{3}$, $P_3 = \frac{1}{1 - P_2} = \frac{3}{4}$
- $P_k = \frac{k}{k+1}$ (Verify this is correct by induction)

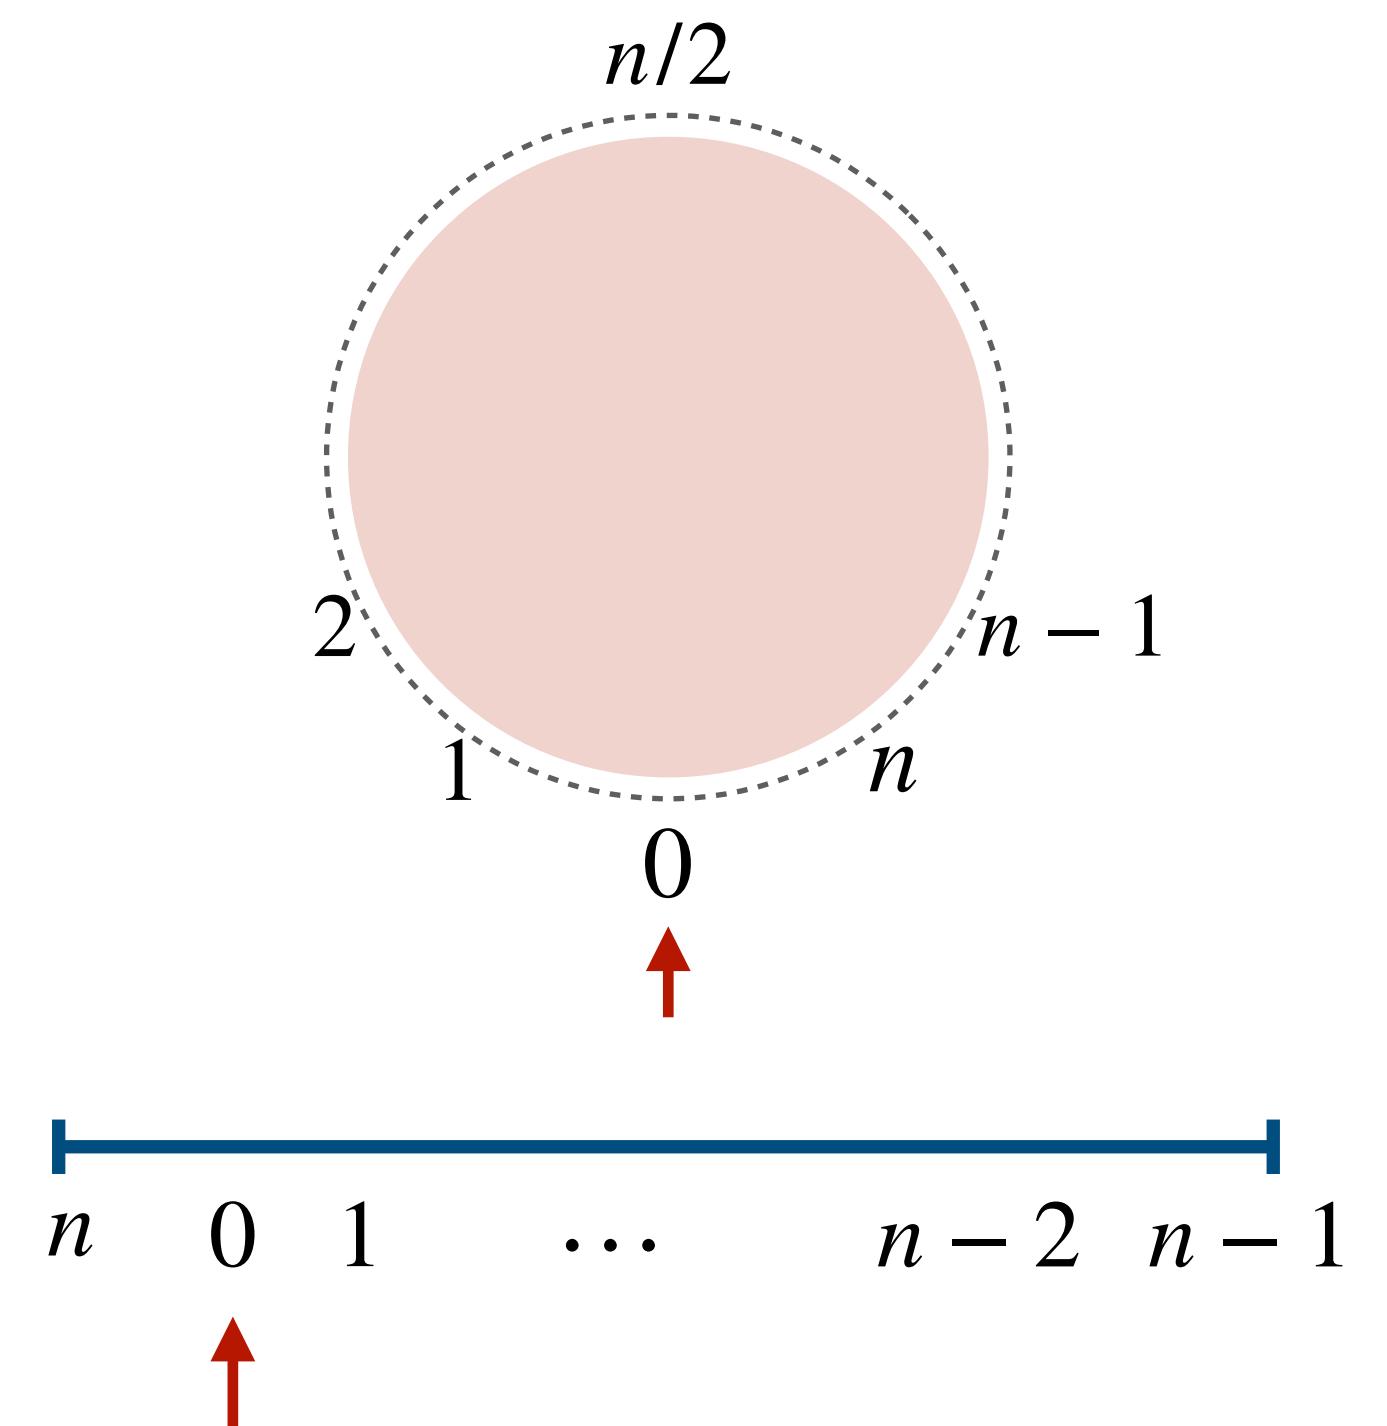


$\boxed{\Pr(S_1 \text{ gets candy before } B) = P_{k-1}}$

$\boxed{\text{If } S_1 \text{ gets candy, we are back in the initial configuration and } A \text{ gets the candy before } B \text{ with probability } P_k.}$

Back to the Candy Game

- Consider player n on the right side of the professor. Only way n can win is the candy travels clockwise all the way around to player $n - 1$ before n ever touches it
 - Thus, if we cut the circle and arrange on a line as shown, n wins if and only if $n - 1$ gets the bowl before n
 - This fits our previous simpler problem model where we want to know the probability that B gets candy before A and $k = n - 2$
 - $\Pr(n - 1 \text{ gets candy before } n) = 1 - \Pr(n \text{ gets candy before } n - 1)$
$$= 1 - \frac{n - 2}{(n - 2) + 1} = \frac{1}{n}$$
 - Student n wins with probability $1/n$!
 - We can extend this argument to show each student wins with probability $1/n$ --- we would never have guessed this!



n wins only if $n - 1$ gets candy before n

Randomized Algorithm I

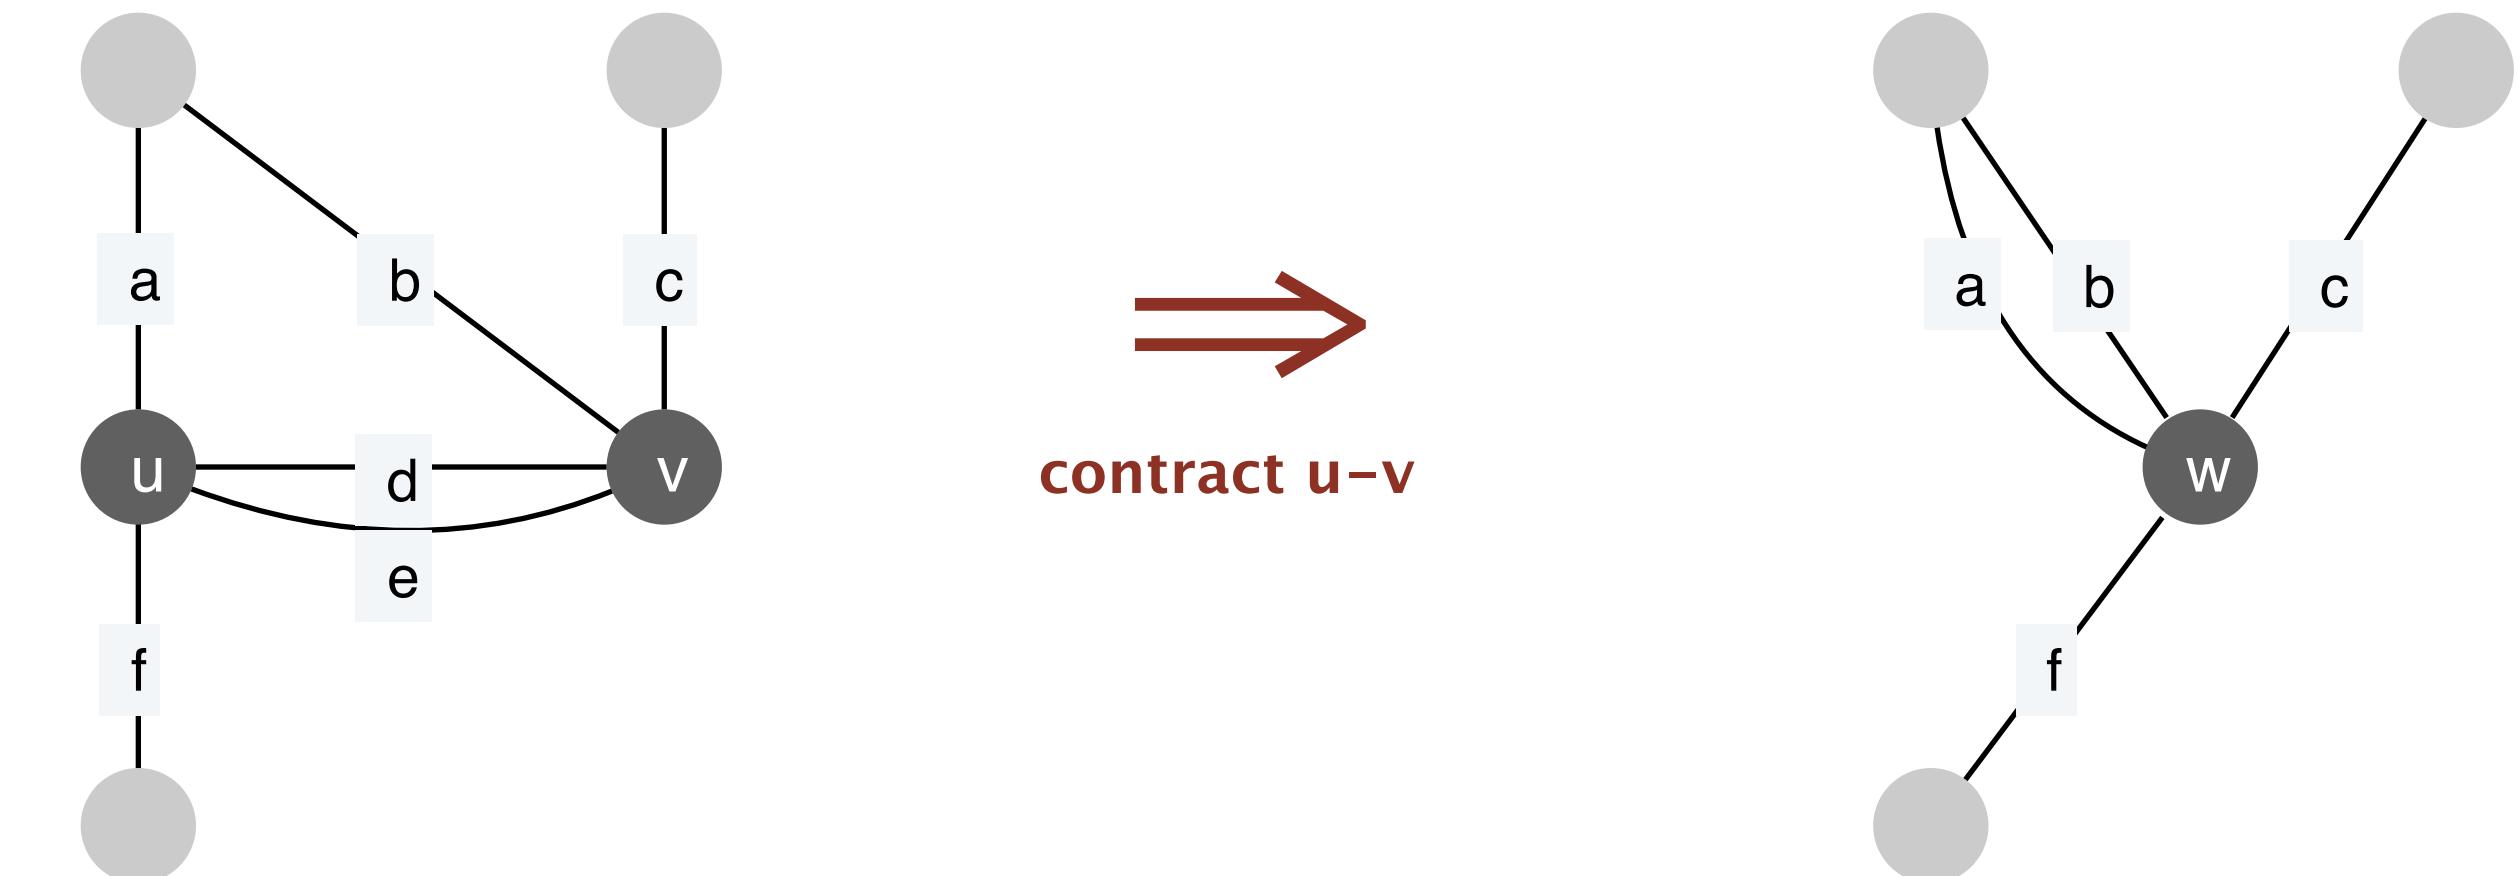
Min Cut

Randomized Min Cut

- **Global min-cut problem.** Given an undirected, unweighted graph $G = (V, E)$, find a cut (A, B) of minimum cardinality (that is, min # of edges crossing it).
- **Applications.** Network reliability, network design, circuit design, etc.
- **Poly-time network-flow solution** (by reduction to min s - t cut).
 - Replace every undirected edge (u, v) with $u \rightarrow v$ and $v \rightarrow u$, each of capacity 1
 - Fix any $s \in V$ and compute min s - t cut for every other node $t \in V - \{s\}$
 - $(n - 1)$ executions of min s - t cut
- Gives impression that finding global min cut is harder than finding a min s - t cut, which is not true
- Deceptively simple and efficient randomized algorithm [Karger 1992]

Karger's Min Cut

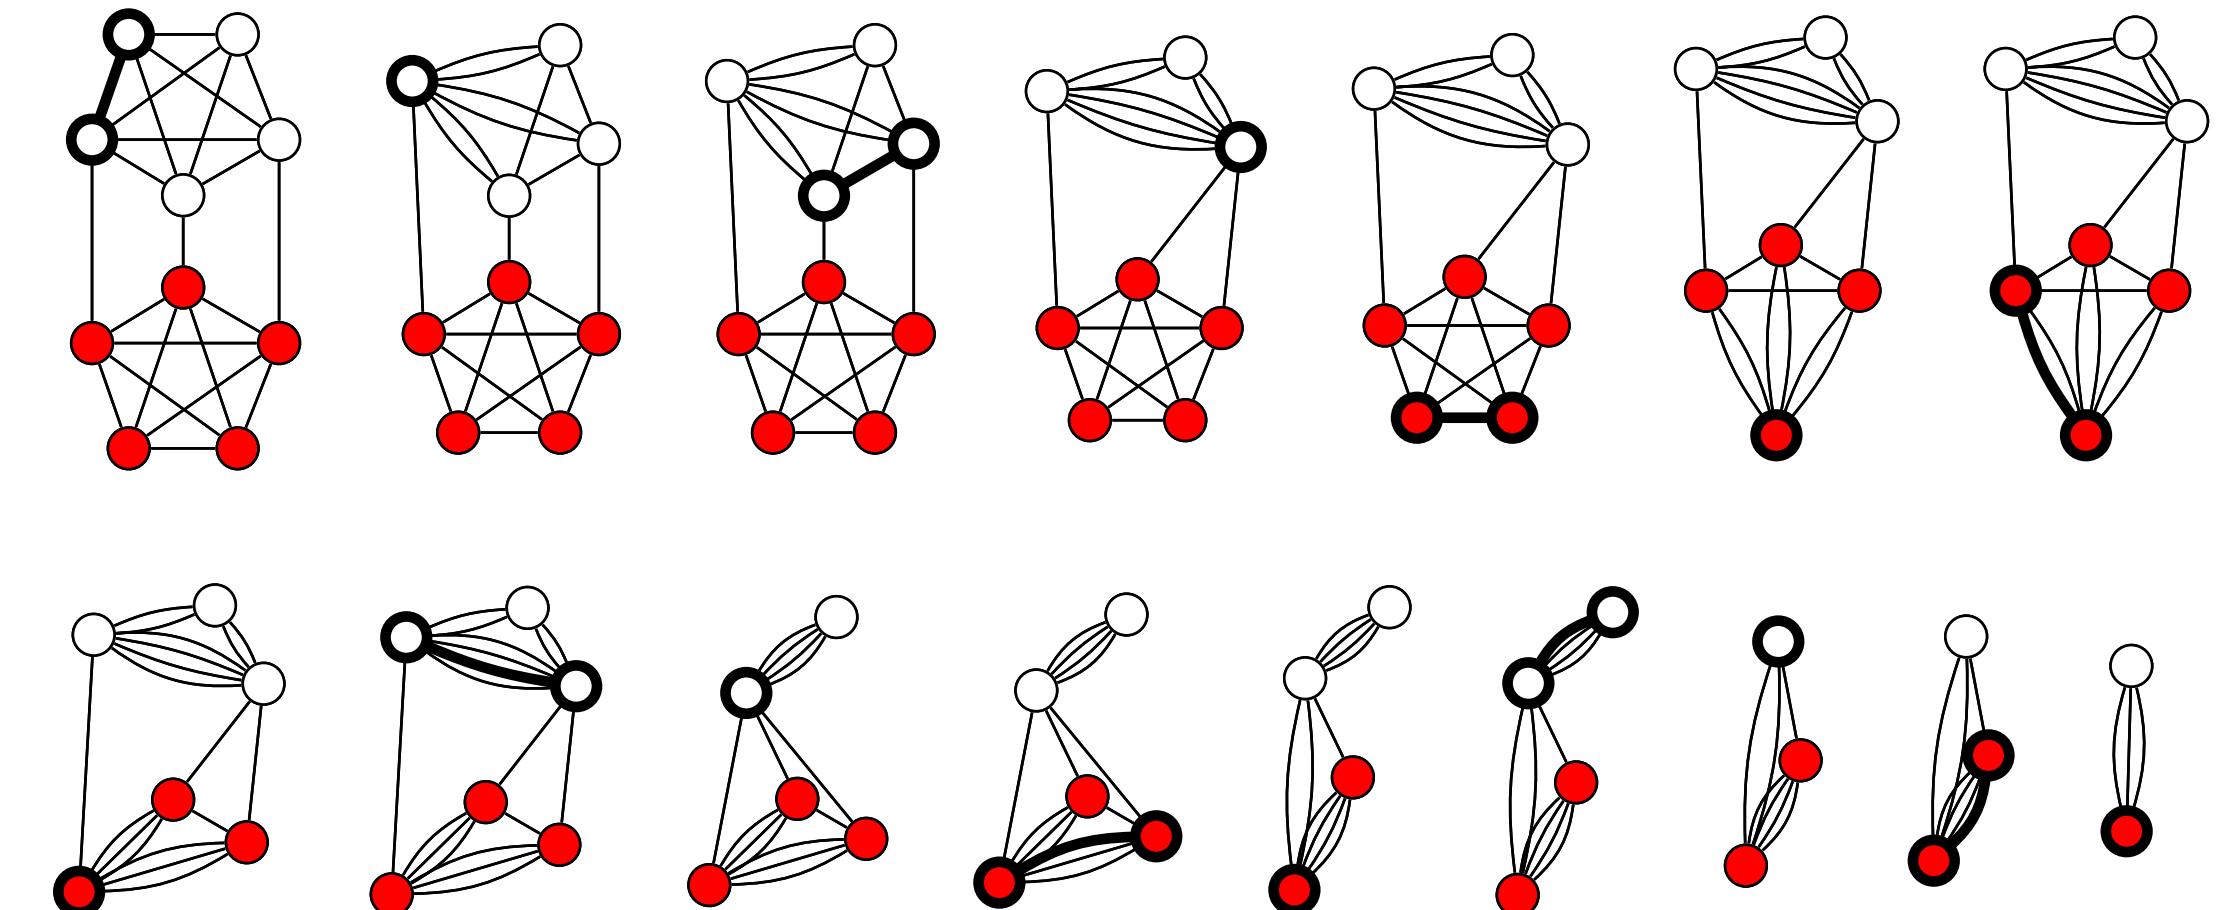
- **Edge contraction:** Contract edge e in G , denoted $G \leftarrow G/e$
 - Replace u and v by single new super-node w
 - Preserve edges, updating endpoints of u and v to w
 - Keep parallel edges, but delete self-loops
- An edge can be contracted in $O(n)$ time, assuming the graph is represented as an adjacency list



Karger's Min Cut

- Algorithm tries to guess the min cut by randomly contracting edges
- Running time:
 - $O(n)$ edge contraction, $O(n)$ iterations: $O(n^2)$
- **Correctness:** How often, if ever, does it return the min cut?

```
GUESSMINCUT( $G$ ):
    for  $i \leftarrow n$  downto 2
        pick a random edge  $e$  in  $G$ 
         $G \leftarrow G/e$ 
    return the only cut in  $G$ 
```

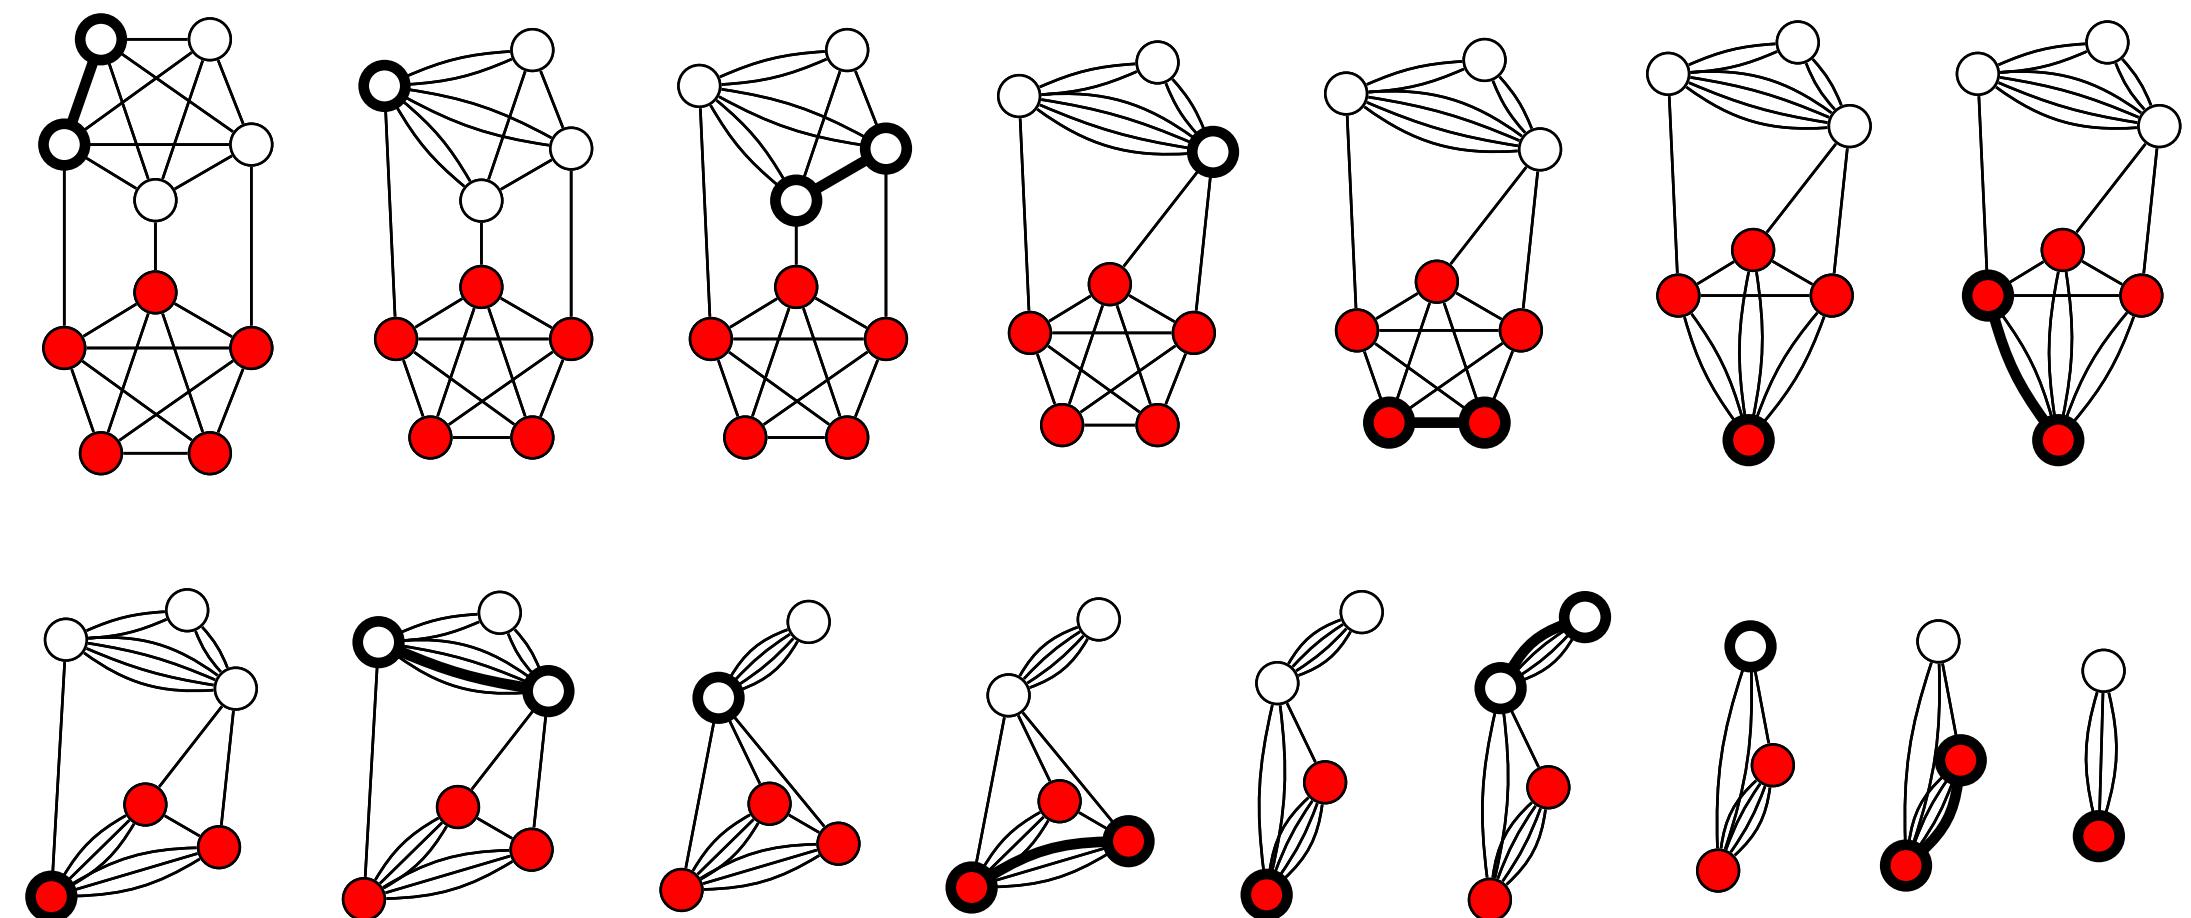


Preserving Cuts

- **Observation.** Any cut in the contracted graph is a cut in the original graph
- Let $C = (S, V - S)$ be any cut, if algorithm never contracts an edge crossing this cut, then it will produce the cut C
- Let C be any arbitrary min cut of cardinality k
- If we pick an edge in G uniformly at random, what is the probability of picking an edge in C ?

$$\bullet \frac{k}{m}$$

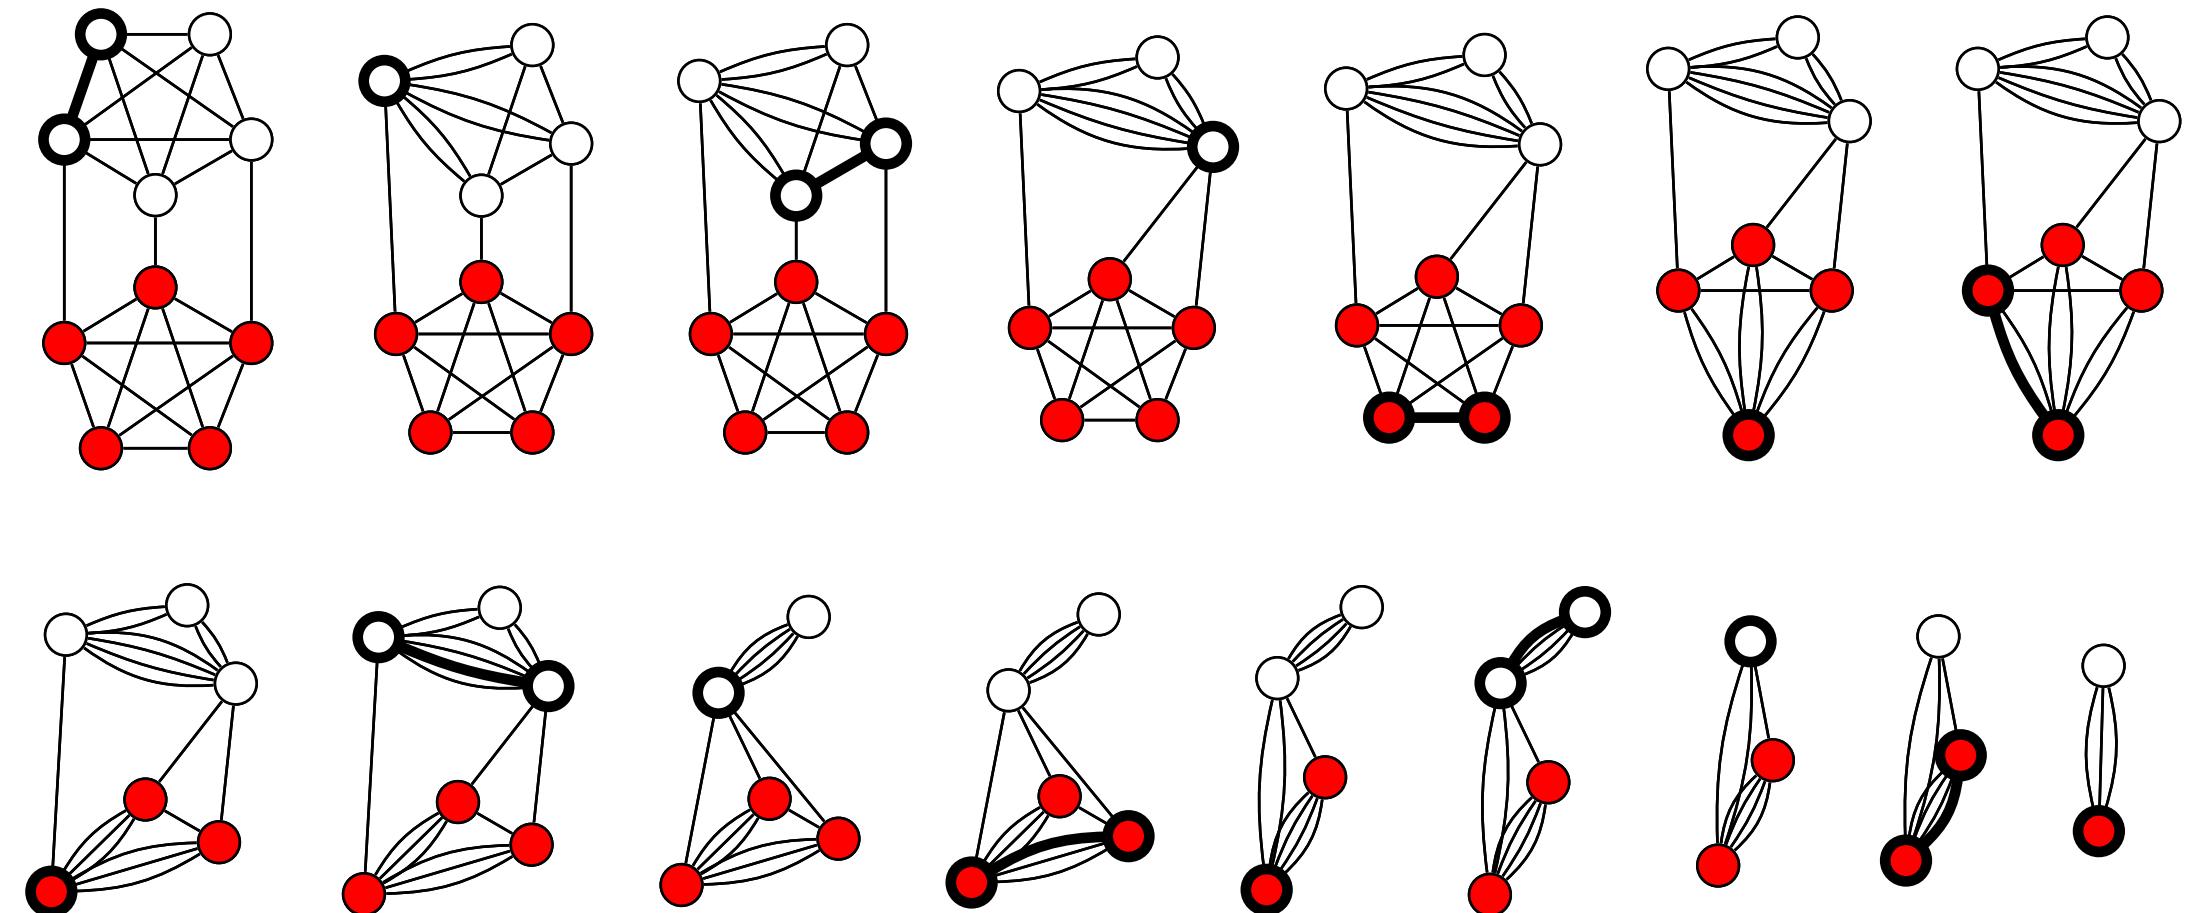
- We want to upper bound this probability
- Can we lower bound m in terms of k ?



Preserving Cuts

- The minimum cut C in G has cardinality k
- What can we say about degree of each vertex in G ?
 - Must be at least k
- G must have at least $nk/2$ edges
- If we pick an edge in G uniformly at random, what is the probability of picking an edge in C ?

$$\bullet \frac{k}{m} \leq \frac{2k}{nk} = \frac{2}{n}$$



Karger's Analysis

- Let C be any arbitrary min cut of cardinality k
 - $\Pr(\text{picking an edge in } C) = \frac{k}{m} \leq \frac{k}{nk/2} = \frac{2}{n}$
 - Probability we don't contract a cut edge in the 1st step $\geq 1 - \frac{2}{n}$
 - After the first edge is contracted, the algorithm proceeds recursively (with independent random choices) on the $(n - 1)$ -vertex graph
 - Let $P(n)$ denote the probability that the algorithm returns the correct min cut on an n -vertex graph, then

Karger's Analysis

- Let $P(n)$ denote the probability that the algorithm returns the correct min cut on an n -vertex graph, then

- $$P(n) \geq \left(1 - \frac{2}{n}\right) \cdot P(n-1), \text{ with base case } P(2) = 1$$

- Expanding the recurrence:

$$P(n) \geq \frac{n-2}{n} \cdot \frac{n-3}{n-1} \cdot \frac{n-4}{n-2} \cdot \frac{n-5}{n-3} \cdots \cdots \frac{4}{6} \cdot \frac{3}{5} \cdot \frac{2}{4} \cdot \frac{1}{3}$$

- Terms cancel out to get: $P(n) \geq \frac{2}{n(n-1)} = 1/\binom{n}{2}$

Amplifying Success Probability

- Thus, a single execution of Karger's min cut algorithm finds the min cut with probability at least $1/\binom{n}{2}$, which is low
 - But, we can amplify our success probability!
- Run the algorithm R times (using independent random choices) and pick the best min-cut among them
- What is probability we don't find the min cut after R repetitions?

$$\cdot \left(1 - 1/\binom{n}{2}\right)^R$$

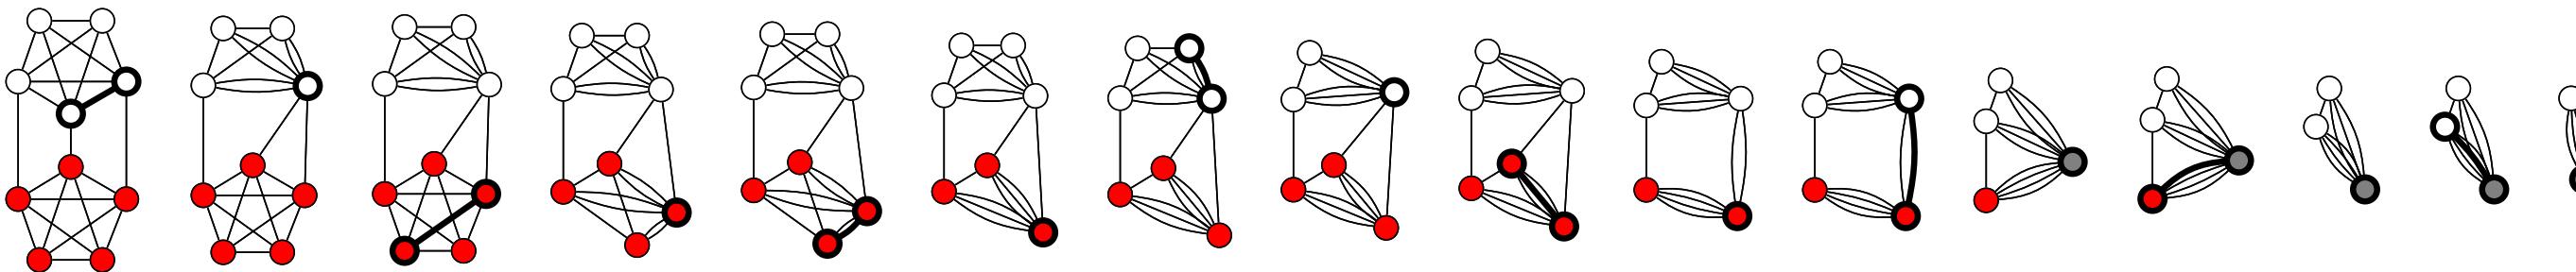
With High Probability

- If we run the algorithm $R = \binom{n}{2} c \ln n$ times, we can make the failure probability polynomially small in n : $\left(\frac{1}{e}\right)^{c \ln n} = \frac{1}{n^c}$
- Karger's algorithm finds the min-cut **with high probability (w.h.p.)**

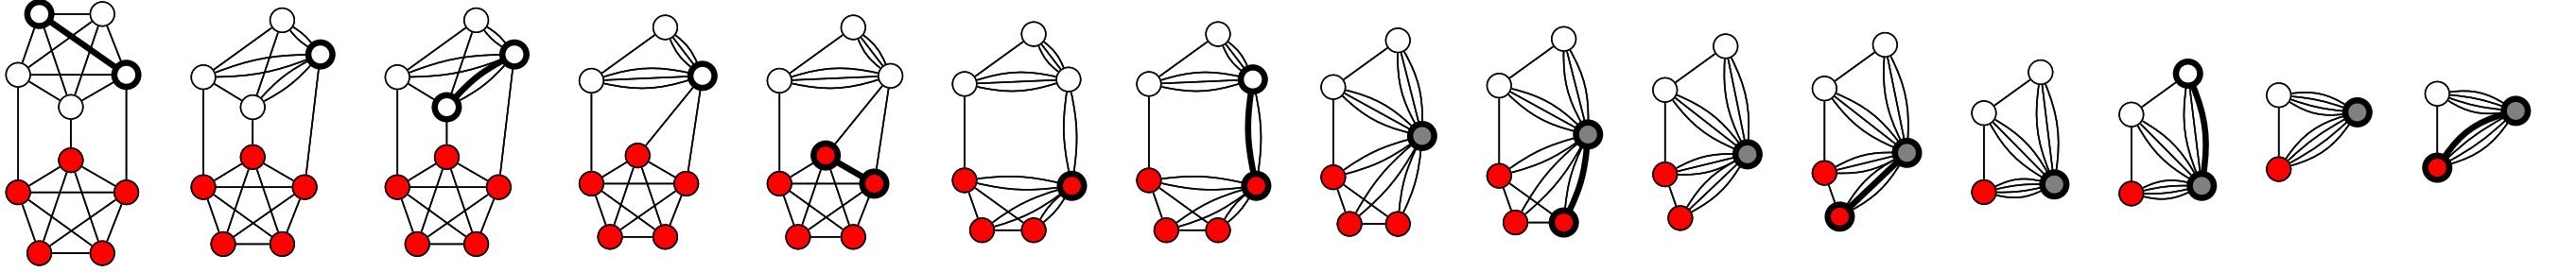
An algorithm is correct **with high probability (w.h.p.)** with respect to input size n if it fails with probability at most $\frac{1}{n^c}$ for any constant $c > 1$.

Example Execution

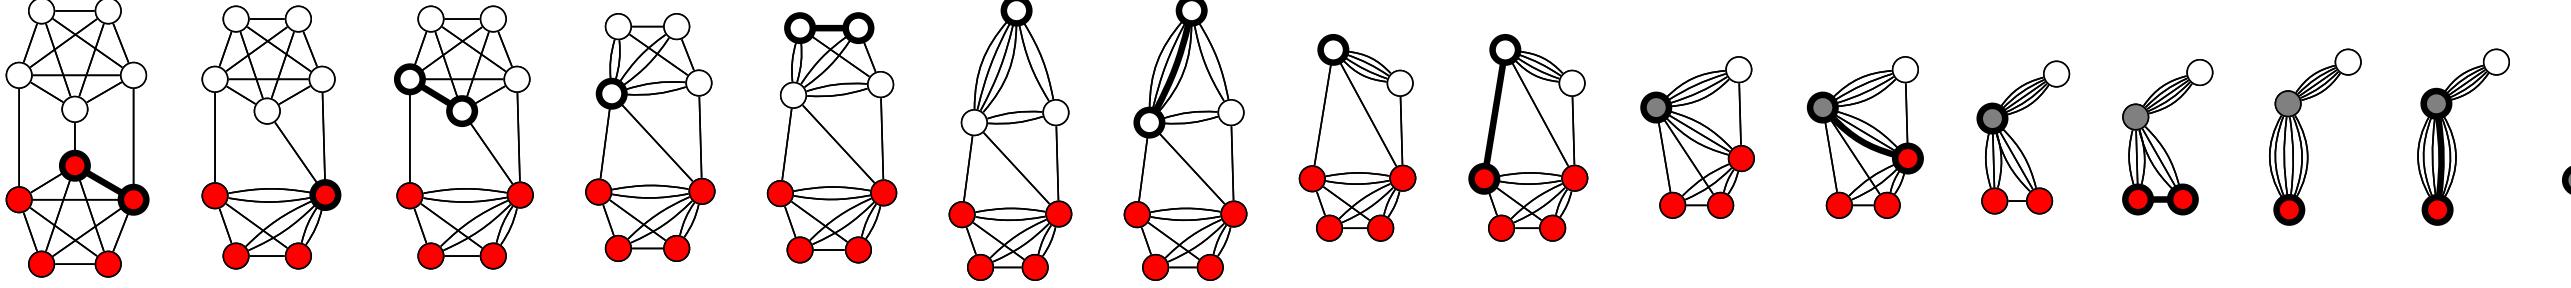
trial 1



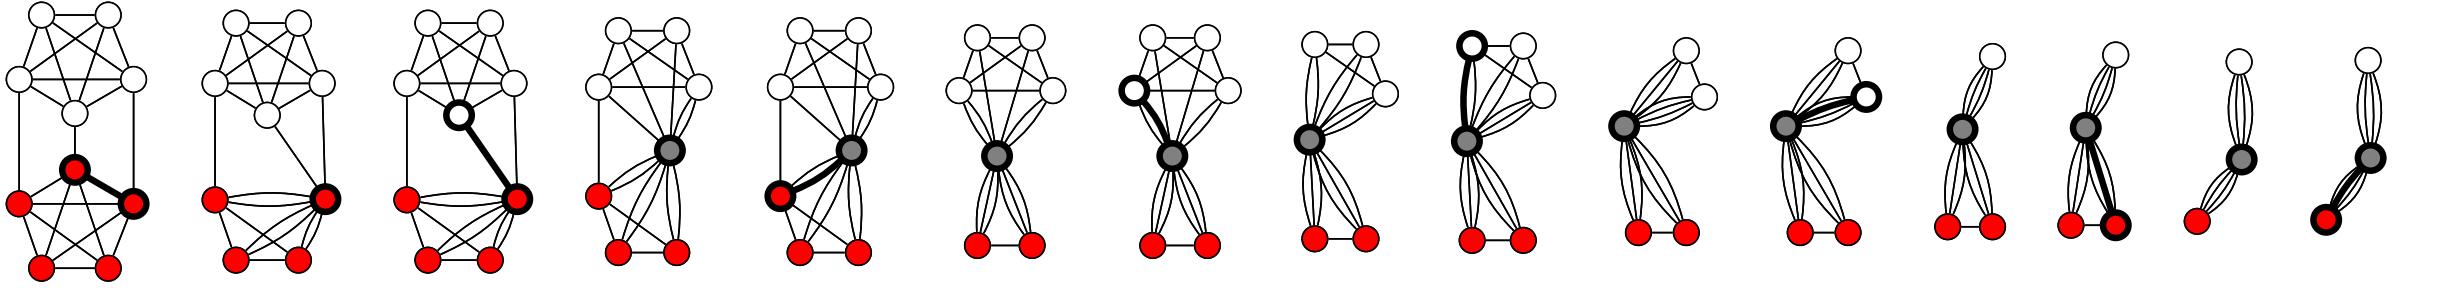
trial 2



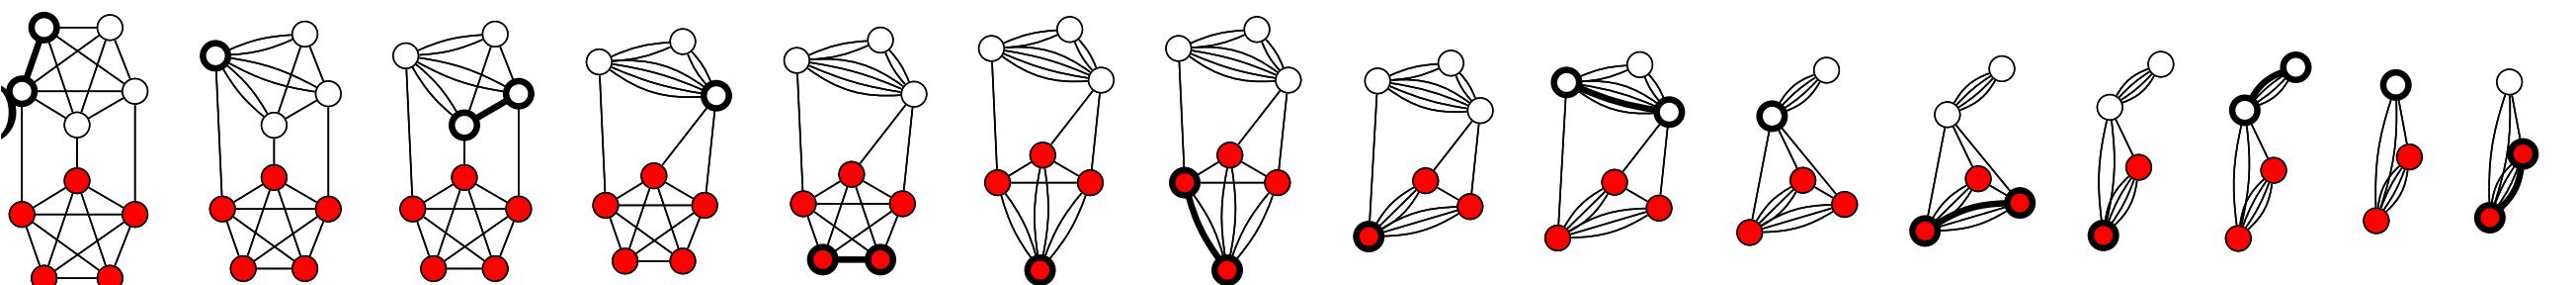
trial 3



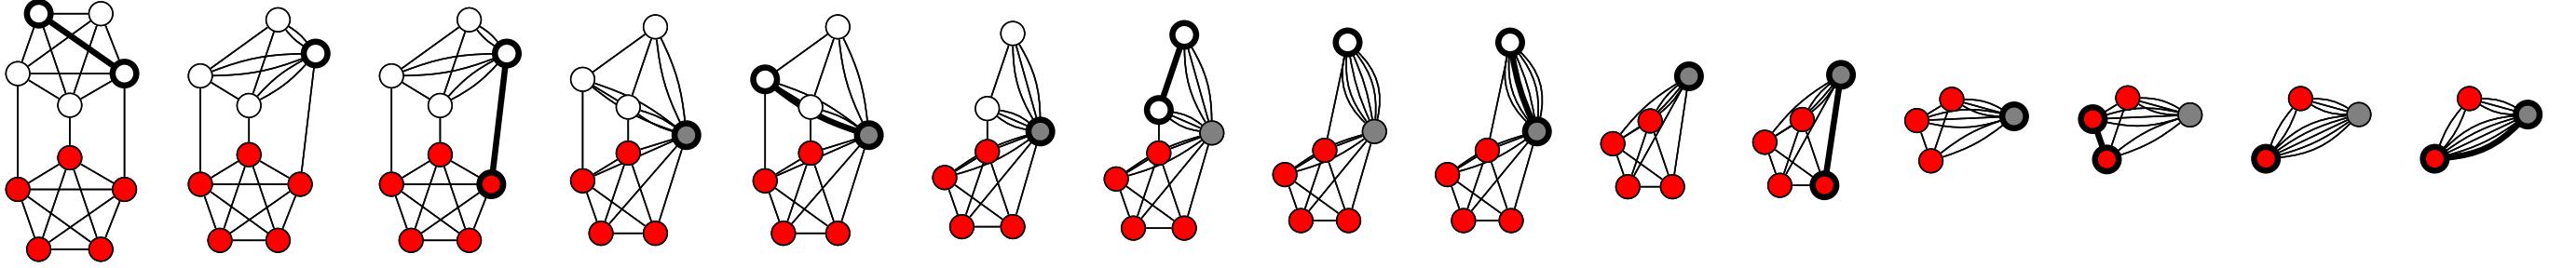
trial 4



trial 5
(finds min cut)



trial 6



...

Reference: Thore Husfeldt

Karger's Running Time

- Thus, Karger's algorithm finds the min-cut [with high probability \(w.h.p.\)](#)
- Running time: we perform $\Theta(n^2 \log n)$ iterations, each $O(n^2)$ time
 - $O(n^4 \log n)$ time
 - Faster than naive-flow-techniques, nothing to get excited about
- Improves to $O(n^2 \log^3 n)$ by guessing cleverly! [Karger-Stein 1996]
- **Idea:** Improve the guessing algorithm using the observation:
 - As the graph shrinks, the probability of contracting an edge in the minimum cut increases
 - At first the probability is very small: $2/n$ but by the time there are three nodes, we have a $2/3$ chance of screwing up!

Takeaways

- Karger's algorithm is an example of a "**Monte Carlo**" randomized algorithm
 - Find the correct answer most of the time
- You can increase the success rate of algorithms with one-sided errors by iterating it multiple times and taking the best solution
 - If the probability of success is $1/f(n)$, then running it $O(f(n)\log n)$ times gives a high probability of success
- If you're more intelligent about how you iterate the algorithm, you can often do much better than this
- Next, we'll see an example of a "**Las Vegas**" algorithm
 - Randomized selection and quick sort

Randomized Algorithms & Data Structures

- *Monte-Carlo algorithms*
 - Find the correct answer most of the time
 - Can usually amplify probability of success with repetitions
 - Example, Karger's min cut
- *Las-Vegas algorithms*
 - Always find the correct answer, e.g. RandQuick sort
 - But the running time guarantees are not worst (but hold in expectation or with high probability depending on the randomness)
- *Randomized data structures*: hashing, search trees, filters, etc.



Randomized Algorithm II

Randomized Selection

Randomized Selection

- **Problem.** Find the k th smallest/largest element in an unsorted array
- Recall our selection algorithm

Select (A, k) :

If $|A| = 1$: return $A[1]$

Else:

Choose a pivot $p \leftarrow A[1, \dots, n]$; let r be the rank of p

$r, A_{<p}, A_{>p} \leftarrow \text{Partition}((A, p))$

If $k == r$, return p

Else if $k < r$: Select $(A_{<p}, k)$

Else: Select $(A_{>p}, k - r)$

Selection with a Good Pivot

- Recall: pivot is “good” if it reduced the array size by at least a constant
 - Gives a recurrence $T(n) \leq T(\alpha n) + O(n)$ for some constant $\alpha < 1$
 - Expands to a decreasing geometric series $T(n) = O(n)$
- In the deterministic algorithm, how did we find a good pivot?
 - Split array into groups of 5
 - And computed the median of group medians
 - The pivot guaranteed that $n \rightarrow 7n/10$
- **Here is a silly idea:** What if we pick the pivot uniformly at random?
 - Seems like the pivot is “usually” around the midpoint
 - What is the expected running time?

Randomized Selection

- **Problem.** Find the k th smallest/largest element in an unsorted array
- Recall our selection algorithm

Select (A, k) :

If $|A| = 1$: return $A[1]$

Else:

Choose a pivot $p \leftarrow A[1, \dots, n]$ uniformly at random; let r be the rank of p

$r, A_{<p}, A_{>p} \leftarrow \text{Partition}((A, p))$

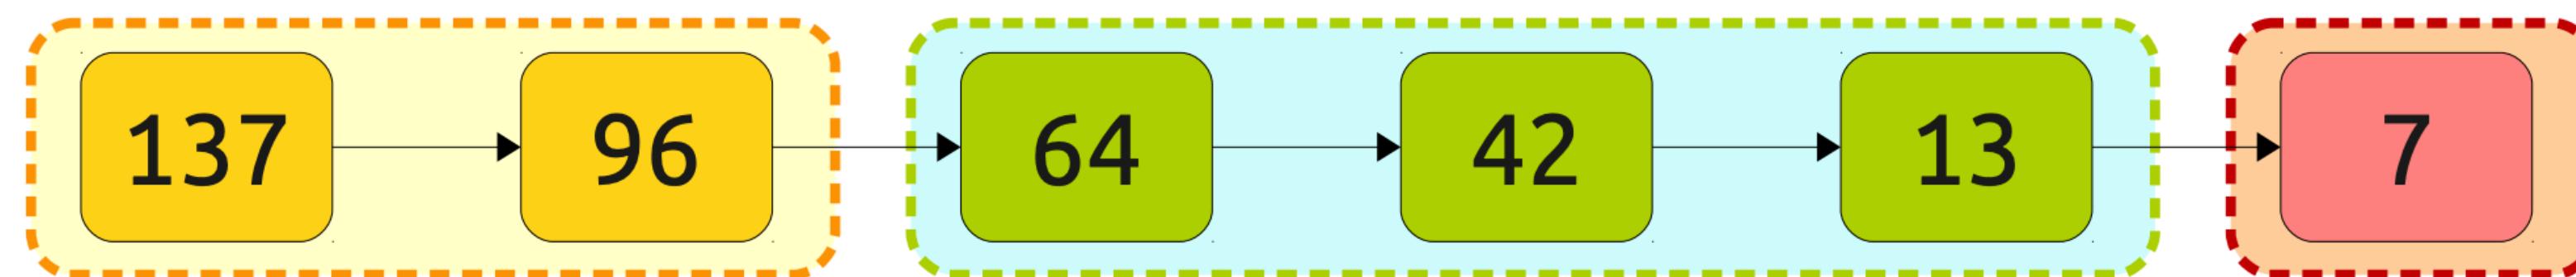
If $k == r$, return p

Else if $k < r$: Select $(A_{<p}, k)$

Else: Select $(A_{>p}, k - r)$

Analyzing Randomized Selection

- Normally, we'd write a recurrence relation for a recursive function
- A bit complicated now--- input size of later recursive call depends on the random choice of pivots in earlier calls
- We will use a different accounting trick for running time
- Randomized selection makes at most one recursive call each time:
 - Group multiple recursive call in “phases”
 - Sum of work done by all calls is equal to the sum of the work done in all the phases



Analyzing in Phases

- **Idea:** let a “phase” of the algorithm be the time it takes for the array size to drop by a constant factor (say $n \rightarrow (3/4) \cdot n$)
- If array shrinks by a constant factor in each phase and linear work done in each phase, what would be the running time?
- $T(n) = c(n + 3n/4 + (3/4)^2n + \dots + 1) = O(n)$
- If we want a $1/4$ th, $3/4$ th split, what range should our pivot be in?
 - Middle half of the array (if n size array, then pivot in $[n/4, 3n/4]$)
 - What is the probability of picking such a pivot?
 - $1/2$
 - Phase ends as soon as we pick a pivot in the middle half
 - Expected # of recursive calls until phase ends? 2

Expected Running Time

- Let the algorithm be in phase j when the size of the array is

- At least $n \left(\frac{3}{4}\right)^j$ but not greater than $n \left(\frac{3}{4}\right)^{j+1}$

- Expected number of iterations within a phase: 2
- Let X_j be the expected number of steps spent in phase j
- $X = X_0 + X_1 + X_2 \dots$ be the total number of steps taken by the algorithm
- Within a phase, the algorithm does work linear in the size of the array in one iterations and thus,

$$E[X_j] \leq 2cn \left(\frac{3}{4}\right)^j$$

- Expected running time: $E[X] = \sum_j E[X_j] \leq \sum_j 2cn \left(\frac{3}{4}\right)^j = 2cn \sum_j \left(\frac{3}{4}\right)^j \leq 8cn = O(n)$

Pivot Selection

- Deterministic and random both take $O(n)$ time
 - What's the advantage of the deterministic algorithm?
 - Worst-case guarantee—the random algorithm could be very slow sometimes
 - What's the advantage of the random algorithm?
 - Much much simpler and better constants hidden in $O()$
- Which should you use?
 - Pretty much always random
 - Question to ask yourself: how often is the randomized algorithm going to be much worse than $O(n)$?

Acknowledgments

- Some of the material in these slides are taken from
 - Kleinberg Tardos Slides by Kevin Wayne (<https://www.cs.princeton.edu/~wayne/kleinberg-tardos/pdf/04GreedyAlgorithmsI.pdf>)
 - Jeff Erickson's Algorithms Book (<http://jeffe.cs.illinois.edu/teaching/algorithms/book/Algorithms-JeffE.pdf>)