


Randomized Quicksort

Admin

- Assignment 8 is due this Wed
- Grading feedback of HW 7 in a couple of days
 - HW 7 Solutions posted on GLOW in the meantime
- Health day: no Lecture on Friday 

Randomized Algorithm I

Min Cut (Wrap Up)

Karger's Min Cut

- Algorithm tries to *guess* the min cut by randomly contracting edges
- Running time $O(n^2)$ (why?)
- Correctness:
How often, if ever, does it return the min cut?


GUESSMINCUT(G):

 for $i \leftarrow n$ downto 2

 pick a random edge e in G

$G \leftarrow G/e$

 return the only cut in G



Amplifying Success Probability

- If we execute $R = \binom{n}{2}$ times, the probability of failure is

- $\left(1 - 1/\binom{n}{2}\right)^{\binom{n}{2}}$: how can we simplify this?

- $\leq \frac{1}{e}$

- If we set $R = \binom{n}{2} c \ln n$, the failure probability becomes polynomially

small in n : $\left(\frac{1}{e}\right)^{c \ln n} = \frac{1}{n^c}$

Important Inequality:

$$(1 - x) \leq \left(\frac{1}{e}\right)^x \text{ for } x \geq 1$$


With High Probability

- If we run the algorithm $R = \binom{n}{2} c \ln n$ times, we can make the failure probability polynomially small in n : $\left(\frac{1}{e}\right)^{c \ln n} = \frac{1}{n^c}$
- Karger's algorithm finds the min-cut **with high probability (w.h.p.)**


An algorithm is correct **with high probability (w.h.p.)** with respect to input size n if it fails with probability at most $\frac{1}{n^c}$ for any constant $c > 1$.

Example Execution


trial 1




trial 2




trial 3




trial 4



trial 5
(finds min cut)



trial 6



...

Reference: Thore Husfeldt

Karger's Running Time

- Thus, Karger's algorithm finds the min-cut with high probability (w.h.p.)
- Running time: we perform $\Theta(n^2 \log n)$ iterations, each $O(n^2)$ time
 - $O(n^4 \log n)$ time
 - Faster than naive-flow-techniques, nothing to get excited about
- **Improves to $O(n^2 \log^3 n)$** by guessing cleverly! [Karger-Stein 1996]
- **Idea:** Improve the guessing algorithm using the observation:
 - As the graph shrinks, the probability of contracting an edge in the minimum cut increases
 - At first the probability is very small: $2/n$ but by the time there are three nodes, we have a $2/3$ chance of screwing up!

Takeaways

- Karger's algorithm is an example of a "**Monte Carlo**" randomized algorithm
 - Find the correct answer most of the time
- You can increase the success rate of algorithms with one-sided errors by iterating it multiple times and taking the best solution
 - If the probability of success is $1/f(n)$, then running it $O(f(n)\log n)$ times gives a high probability of success
- If you're more intelligent about how you iterate the algorithm, you can often do much better than this
- Next, we'll see an example of a "**Las Vegas**" algorithm
 - Randomized selection and quick sort

Randomized Algorithms & Data Structures

- *Monte-Carlo algorithms*
 - Find the correct answer most of the time
 - Can usually amplify probability of success with repetitions
 - Example, Karger's min cut
- *Las-Vegas algorithms*
 - Always find the correct answer, e.g. RandQuick sort
 - But the running time guarantees are not worst (but hold in expectation or with high probability depending on the randomness)
- *Randomized data structures*: hashing, search trees, filters, etc.



Randomized Algorithm II

Randomized Selection

Randomized Selection

- **Problem.** Find the k th smallest/largest element in an unsorted array
- Recall our selection algorithm

Select (A, k):

If $|A| = 1$: return $A[1]$

Else:

Choose a pivot $p \leftarrow A[1, \dots, n]$; let r be the rank of p

$r, A_{<p}, A_{>p} \leftarrow \text{Partition}((A, p))$

If $k = r$, return p

Else if $k < r$: Select ($A_{<p}, k$)

Else: Select ($A_{>p}, k - r$)

Selection with a Good Pivot

- Recall: pivot is “good” if it reduced the array size by at least a constant
 - Gives a recurrence $T(n) \leq T(\alpha n) + O(n)$ for some constant $\alpha < 1$
 - Expands to a decreasing geometric series $T(n) = O(n)$
- In the deterministic algorithm, how did we find a good pivot?
 - Split array into groups of 5
 - And computed the median of group medians
 - The pivot guaranteed that $n \rightarrow 7n/10$
- **Here is a silly idea:** What if we pick the pivot uniformly at random?
 - Seems like the pivot is “usually” around the midpoint
 - What is the expected running time?

Randomized Selection

- **Problem.** Find the k th smallest/largest element in an unsorted array
- Recall our selection algorithm

Select (A, k):

If $|A| = 1$: return $A[1]$

Else:

Choose a pivot $p \leftarrow A[1, \dots, n]$ uniformly at random; let r be the rank of p

$r, A_{<p}, A_{>p} \leftarrow \text{Partition}((A, p))$


If $k = r$, return p

Else if $k < r$: Select ($A_{<p}, k$)

Else: Select ($A_{>p}, k - r$)

Analyzing Randomized Selection

- Normally, we'd write a recurrence relation for a recursive function
- A bit complicated now--- input size of later recursive call depends on the random choice of pivots in earlier calls
- We will use a different accounting trick for running time
- Randomized selection makes at most one recursive call each time:
 - Group multiple recursive call in “phases”
 - Sum of work done by all calls is equal to the sum of the work done in all the phases



Analyzing in Phases

- **Idea:** let a “phase” of the algorithm be the time it takes for the array size to drop by a constant factor (say $n \rightarrow (3/4) \cdot n$)
- If array shrinks by a constant factor in each phase and linear work done in each phase, what would be the running time?
- $T(n) = c(n + 3n/4 + (3/4)^2n + \dots + 1) = O(n)$
- If we want a 1/4th, 3/4th split, what range should our pivot be in?
 - Middle half of the array (if n size array, then pivot in $[n/4, 3n/4]$)
 - What is the probability of picking such a pivot?
 - 1/2
- Phase ends as soon as we pick a pivot in the middle half
 - Expected # of recursive calls until phase ends? 2

Expected Running Time

- Let the algorithm be in phase j when the size of the array is
 - At least $n \left(\frac{3}{4}\right)^j$ but not greater than $n \left(\frac{3}{4}\right)^{j+1}$
- Expected number of iterations within a phase: 2
- Let X_j be the expected number of steps spent in phase j
- $X = X_0 + X_1 + X_2 \dots$ be the total number of steps taken by the algorithm
- $E(X_j) = E(\# \text{ of iterations until } j\text{th phase ends} \cdot \# \text{ steps in phase } j)$
- $E(X_j) \leq n(3/4)^j \cdot E(\# \text{ iterations until } j\text{th phase ends}) = n(3/4)^j$

Expected Running Time

- Let X_j be the expected number of steps spent in phase j
- $X = X_0 + X_1 + X_2 \dots$ be the total number of steps taken by the algorithm
- $E(X_j) = E(\# \text{ of iterations until } j\text{th phase ends} \cdot \# \text{ steps in phase } j)$
- $E(X_j) \leq n(3/4)^j \cdot E(\# \text{ iterations until } j\text{th phase ends}) = n(3/4)^j$
- Now we can apply linearity of expectation:

$$\begin{aligned} E[X] &= \sum_j E[X_j] \leq \sum_j 2cn \left(\frac{3}{4}\right)^j = 2cn \sum_j \left(\frac{3}{4}\right)^j \\ &\leq 8cn = \Theta(n) \end{aligned}$$

Pivot Selection

- Deterministic and random both take $O(n)$ time
 - What's the advantage of the deterministic algorithm?
 - Worst-case guarantee—the random algorithm could be very slow sometimes
 - What's the advantage of the random algorithm?
 - Much much simpler and better constants hidden in $O()$
- Which should you use?
 - Pretty much always random
 - Question to ask yourself:
 - how often is the randomized algorithm going to be much worse than $O(n)$?

Randomized Algorithm III

Randomized QuickSort

Randomized Quicksort

- Recall deterministic Quicksort
- Depending on the choice pivot, could be $O(n^2)$
- What if we pick the pivot uniformly at random?
 - We saw in that in randomized selection this lead to good pivots half the time

Quicksort(A):

If $|A| < 3$: Sort(A) directly

Else: choose a pivot element $p \leftarrow A$

$A_{<p}, A_{>p} \leftarrow$ Partition around p

Quicksort($A_{<p}$)

Quicksort($A_{>p}$)

Randomized Quicksort

- Intuitively half the pivots will be good, half bad
- We analyze quick sort using another accounting trick
- Total work done can be split into to types:
 - Work done making recursive calls (lower order term, turns out)
 - Work partitioning the elements
- How many recursive calls in the worst case?
 - Each time at least element in the smaller partition
 - $O(n)$

Randomized Quicksort

- We thus need to bound the work partitioning elements
- Partitioning an array of size n around a pivot p takes exactly $n - 1$ comparisons
- We won't look at partitions made in each recursive calls, which depend on the choice of random pivot
- **Idea:** Account for the total work done by the partition step by summing up the total number of comparisons made
- Two ways to count total comparisons:
 - Look at the size of arrays across recursive calls and sum
 - Look at all pairs of elements and count total # of times they are compared (easier to do in this case)

Aside: Randomized Analysis

- Often multiple ways to determine a randomized algorithm's cost
- We can split into phases, or count the cost directly. We can calculate each probability, or use linearity of expectation
- Intrinsically some “cleverness” involved in choosing the way that gets you a clean answer
- In this class I'm going to try to ask you problems where there's a clear path to finding the solution (either it follows directly from the question, or I'll ask about problems you've seen before)
- That said, here's a very clever way to calculate Quicksort's running time

Counting Total Comparisons

- Just for analysis, let B denote the sorted version of input array A , that is, $B[i]$ is the i th smallest element in A
- Define random variable X_{ij} as the number of times Quicksort compares $B[i]$ and $B[j]$
- Observation: $X_{ij} = 0$ or $X_{ij} = 1$, why?
 - $B[i], B[j]$ only compared when one of them is the current pivot; pivots are excluded from future recursive calls
- Let $T = \sum_{i=1}^n \sum_{j=i+1}^n X_{ij}$ be the total number of comparisons made by randomized Quicksort



Expected Running Time

- **Goal:** $E[T] = E \left[\sum_{i=1}^n \sum_{j=i+1}^n X_{ij} \right] = \sum_{i=1}^n \sum_{j=i+1}^n E[X_{ij}]$
- $E[X_{ij}] = \Pr[X_{ij} = 1]$
- When is $X_{ij} = 1$? That is, when are $B[i]$ and $B[j]$ compared?
- Consider a particular recursive call. Let rank of pivot p be r .
 - Let's think about where $B[i], B[j]$ lie with respect to p

Expected Running Time

- Goal: $E[T] = E \left[\sum_{i=1}^n \sum_{j=i+1}^n X_{ij} \right] = \sum_{i=1}^n \sum_{j=i+1}^n E[X_{ij}]$
- $E[X_{ij}] = \Pr[X_{ij} = 1]$
- When is $X_{ij} = 1$? That is, when are $B[i]$ and $B[j]$ compared?
- Consider a particular recursive call. Let rank of pivot p be r .
 - Case 1. One of them is the pivot: $r = i$ or $r = j$
 - Case 2. Pivot is between them: $r > i$ and $r < j$
 - Case 3. Both less than the pivot: $r > i, j$
 - Case 4. Both greater than the pivot: $r < i, j$

Comparisons for Each Case

- **Case 1.** $r = i$ or $r = j$
 - $B[i]$ and $B[j]$ are compared once and one of them is excluded from all future calls
- **Case 2.** $r > i$ and $r < j$
 - $B[i]$ and $B[j]$ are both compared to the pivot but not to each other, after which they are in different recursive calls: will never be compared again
- **Case 3.** $r > i, j$ and **Case 4.** $r < i, j$
 - $B[i]$ and $B[j]$ are not compared to each other, they are both in the same subarray and may be compared in the future
- **Takeaway:** $B[i]$, $B[j]$ are compared for the 1st time when one of them is chosen as pivot from $B[i], B[i + 1], \dots, B[j]$ & never again

Expected Running Time

- $\Pr[X_{ij} = 1] = \Pr(\text{one of them is picked as pivot from } B[i], B[i + 1], \dots, B[j])$

- $\Pr[X_{ij} = 1] = \frac{2}{j - i + 1}$

- $E[T] = \sum_{i=1}^n \sum_{j=i+1}^n E[X_{ij}] = 2 \sum_{i=1}^n \sum_{j=i+1}^n \frac{1}{j - i + 1}$

Expected Running Time

- $B[i]$ and $B[j]$ are compared iff one of them is the first pivot chosen from the range $B[i], B[i + 1], \dots, B[j]$

- $\Pr[X_{ij} = 1] = \frac{2}{j - i + 1}$

- $E[T] = \sum_{i=1}^n \sum_{j=i+1}^n E[X_{ij}] = 2 \sum_{i=1}^n \sum_{j=i+1}^n \frac{1}{j - i + 1}$

- For fixed i , inner sum is $\frac{1}{2} + \frac{1}{3} + \frac{1}{4} + \dots + \frac{1}{n - i + 1} \leq \sum_{\ell=2}^n \frac{1}{\ell} = O(\log n)$

- Thus, expected number of comparisons is:
 $E[T] = O(n \log n + n) = O(n \log n)$

Quick Sort Summary

- Las Vegas algorithms like Quicksort and Selection are always correct but their running time guarantees hold in expectation
- We can actually prove that the number of comparisons made by Quicksort is $O(n \log n)$ **with high probability**
 - This means the the probability that the running time of quicksort is more than a constant factor away from its expectation is very small (polynomially small: less than $1/n^c$ for constant $c \geq 1$)
 - Whp bounds are called **concentration bounds**

Randomized Algorithms & Data Structures

- *Monte-Carlo algorithms*
 - Find the correct answer most of the time
 - Can usually amplify probability of success with repetitions
 - Example, Karger's min cut
- *Las-Vegas algorithms*
 - Always find the correct answer, e.g. RandQuick sort
 - But the running time guarantees are not worst (but hold in expectation or with high probability depending on the randomness)
- *Randomized data structures*: hashing, search trees, filters, etc.



Acknowledgments

- Some of the material in these slides are taken from
 - Kleinberg Tardos Slides by Kevin Wayne (<https://www.cs.princeton.edu/~wayne/kleinberg-tardos/pdf/04GreedyAlgorithmsI.pdf>)
 - Jeff Erickson's Algorithms Book (<http://jeffe.cs.illinois.edu/teaching/algorithms/book/Algorithms-JeffE.pdf>)