

Wrapping Up CS 256

We Did It!

Admin: Last Lecture

- **Course evaluations** in second half of class today
- Will post HW 10 (solution sketch) on GLOW later today
- Help hours before the final:
 - Office hours today 1.30- 3pm, TA hours rest of the day
- **Final.** 24-hour final, will be available on Gradescope from **Thursday May 20 8.30 am** and must be submitted by **May 28, 8.30 pm**
- **Honor Code (final).**
 - Can only refer to course materials (notes/book/GLOW), honor code violation to use external resources, google terms, or discuss exam with others



Thesis Presentations: Tomorrow

10:00 am: David Lee, "A Practical Adaptive Quotient Filter"

10:45 am: Markus Feng, "Shard: Shell for Ad Hoc, Reliable, Distributed Programs"

---Break---

1:00 pm: Andrew Thai, "A Filtering Approach for NFA Processing on GPUs"

1:45 pm: Peter Zhao, "Time-Space Tradeoffs in Indexing Problems via Reductions to Function Inversion

Algorithmic Design Paradigms

- Graph Traversals
 - Can do a lot with just BFS and DFS!
- Greedy Algorithms
 - Simple solution/ hard to show why they work
- Divide & Conquer and Recurrences
 - Using recursion as a powerful design technique
- Dynamic Programming
 - Exponential-seeming problems with polynomial time recursive solutions
- Network Flows
 - Many complicated looking optimization problems in disguise!
- Randomized Algorithms
 - Simple, elegant algorithms that work really well!
- Approximation Algorithms
 - Worst-case guarantees for intractable problems

Problem Solving Strategies

- How to make an abstract problem concrete
- How different parts of the problem commute/compose together
- **Problem reduction:** recognizing one problem in another
 - Reduce an instance of Problem A to Problem B
 - Use algorithm for Problem B to solve it
 - Why it works (necessary and sufficient conditions)

Biggest Takeaways:

- Learning to think algorithmically
- How to formalize your intuition

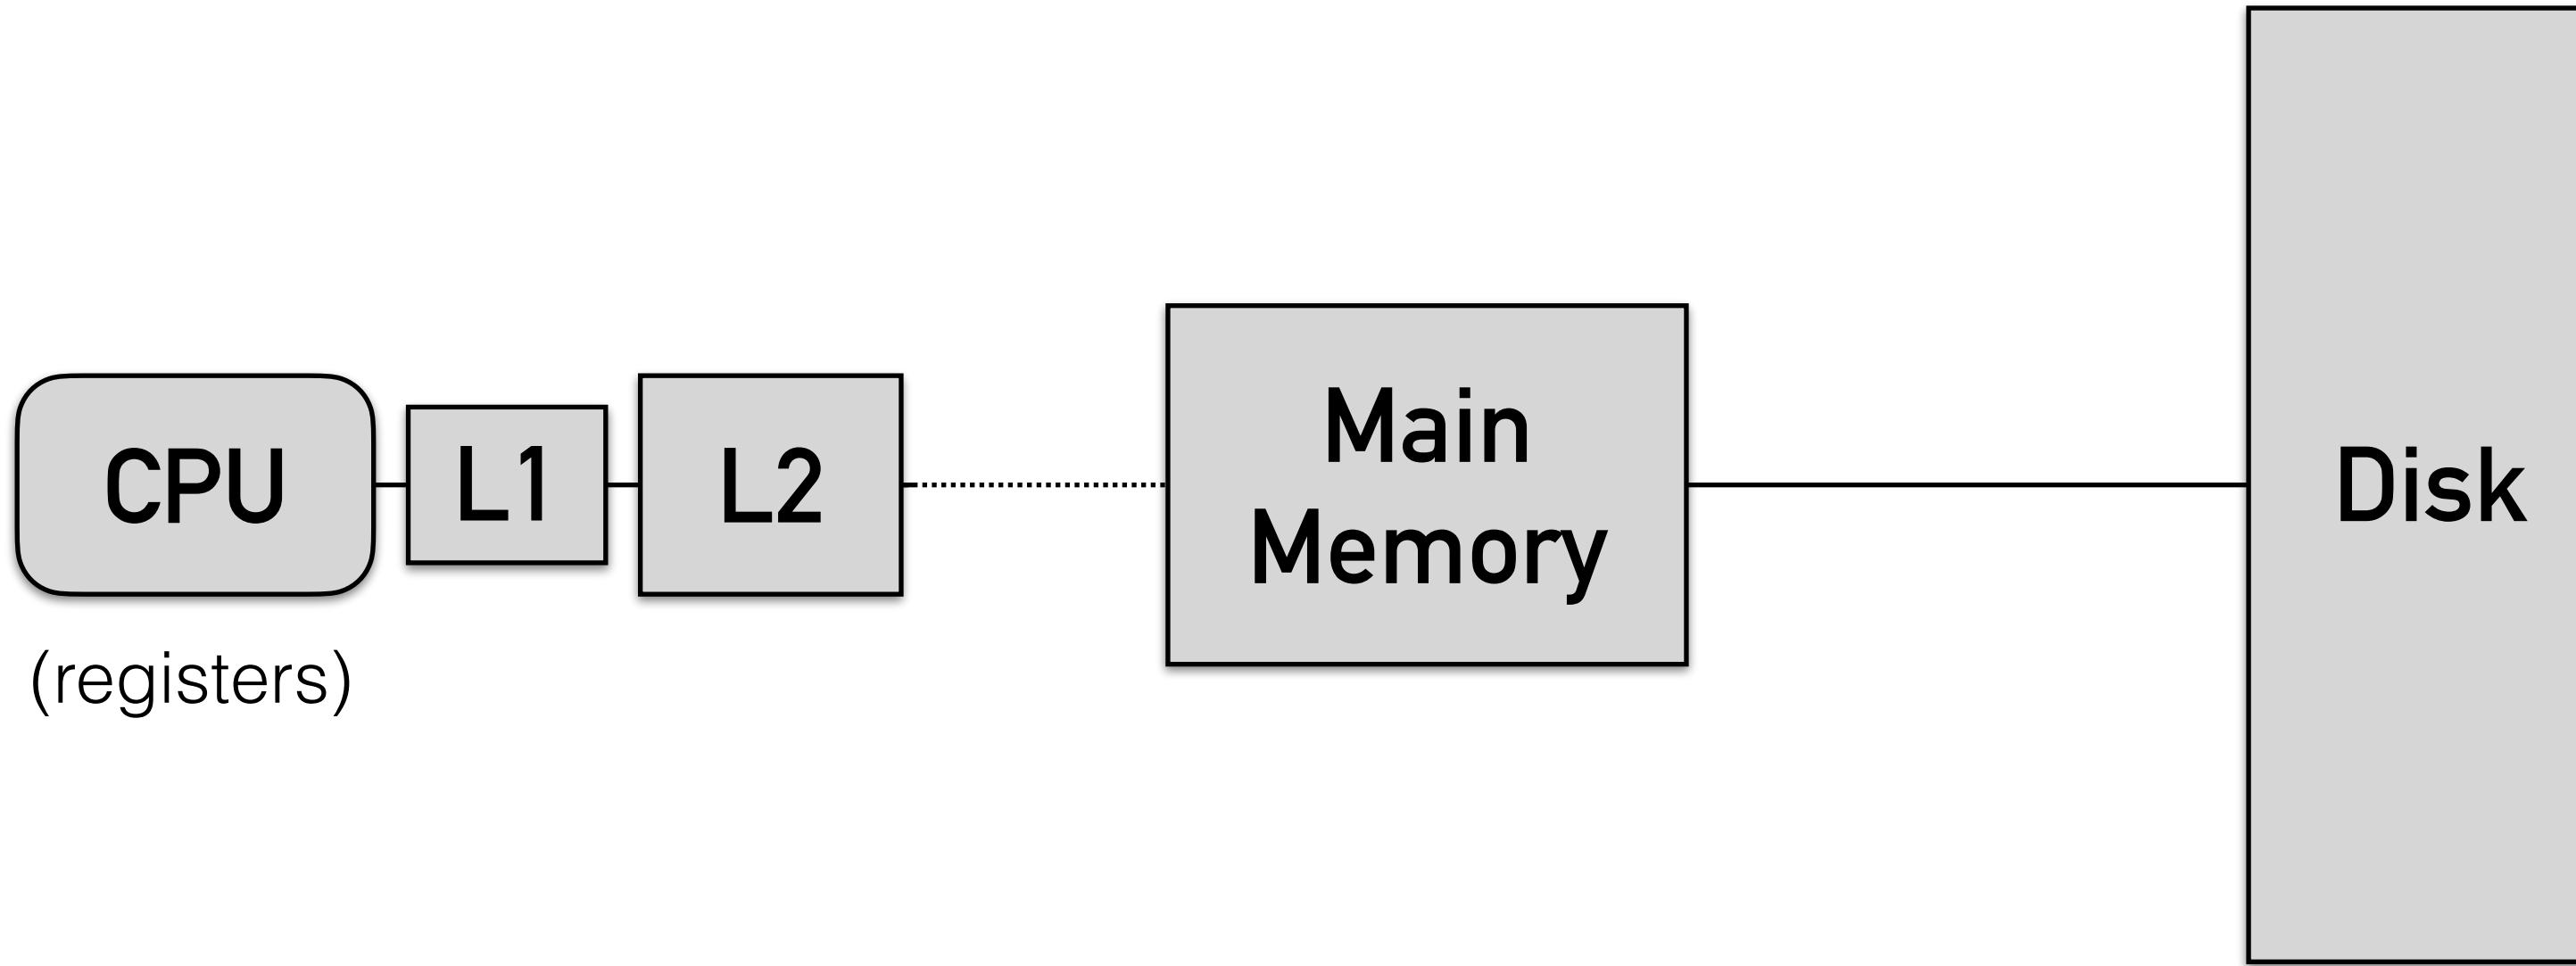
Flip Side of Design: Intractability

- How do we draw boundaries between what is computable and what is impossible
- No matter how much we'd want it, drawing these boundaries isn't easy
- Brings up the famous **P** versus **NP** question
- When we can't prove something is unconditionally hard, we prove it conditioned on some conjecture we all believe is not true
 - SAT/ 3SAT, Independent Set, Vertex cover
 - Set cover, Hamiltonian Cycle/Path
 - Subset Sum, Graph Coloring, Clique
- Shows the power of mathematical characterization
- *If you are CS major, you will see these again in CS 361*

Other Models of Computation

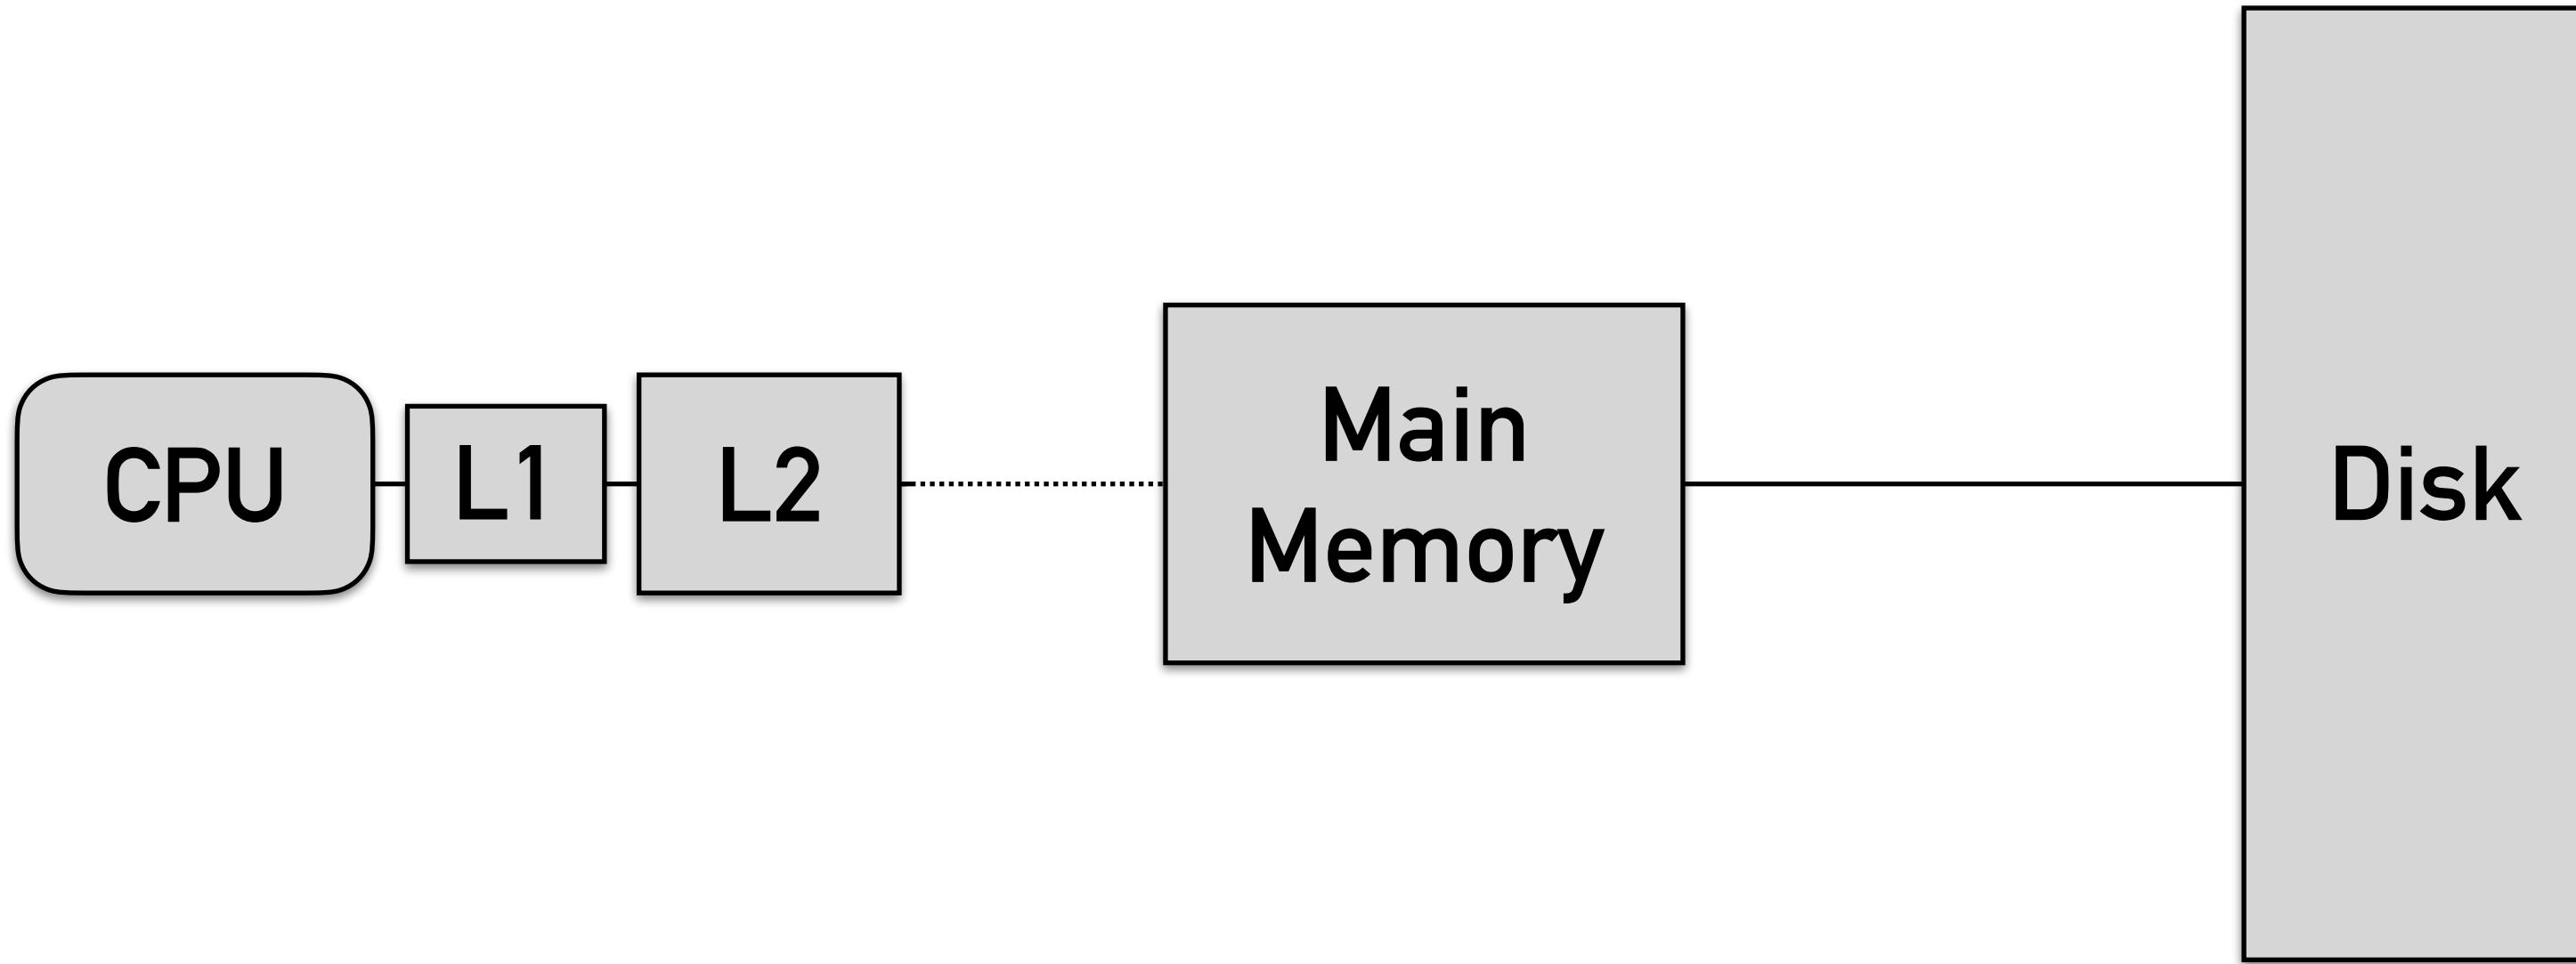
Aka: Applying what
we've learned in
practice

Modern Memory Hierarchy



- $\approx 10K \rightarrow 100K \rightarrow \text{MBs} \rightarrow \text{GBs} \rightarrow \text{TBs}$
- $\approx \text{ns} \rightarrow 10\text{ns} \rightarrow 10\mu\text{s} \rightarrow \mu\text{s} \rightarrow 10\text{ms}$
- (Left to right). Bigger but slower latency (distance travelled and physical seek time on disk)
- Bandwidth (amount of data that can transferred) usually matched

Modern Memory Hierarchy

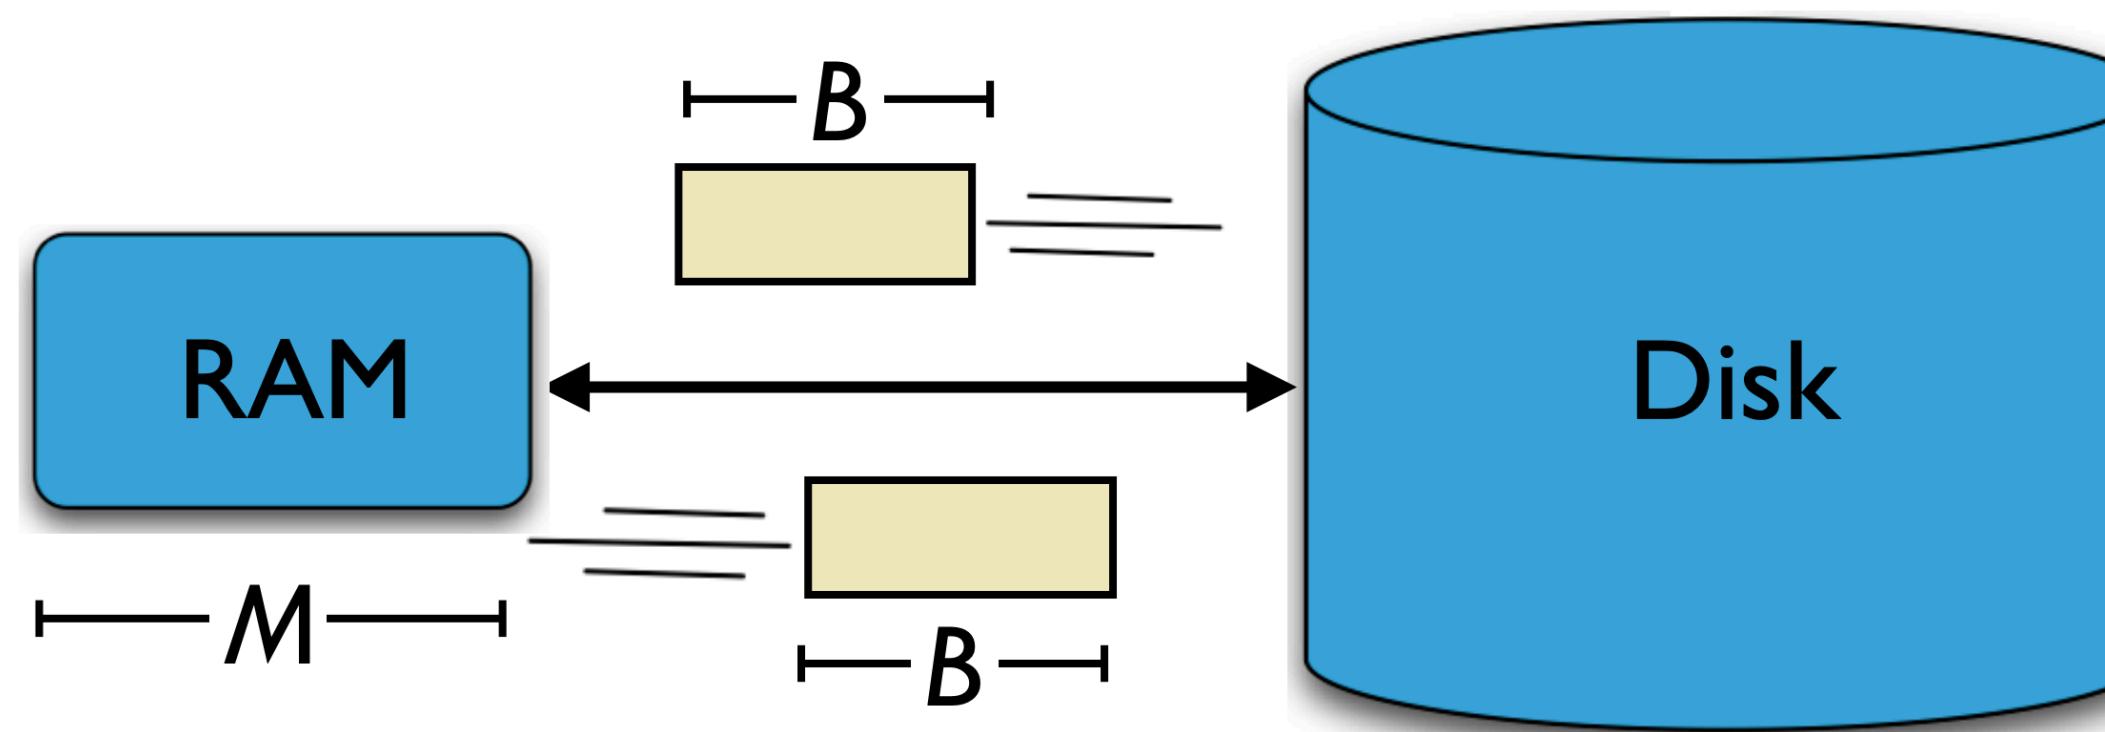


- **Blocking.** Idea to mitigate latency.
 - When fetching data, get an entire block containing it
 - Cache-efficient algorithms have good locality
 - Use items in same block or nearby blocks (spatial locality)
 - Reuse blocks already in cache (temporal locality)

External-Memory Model

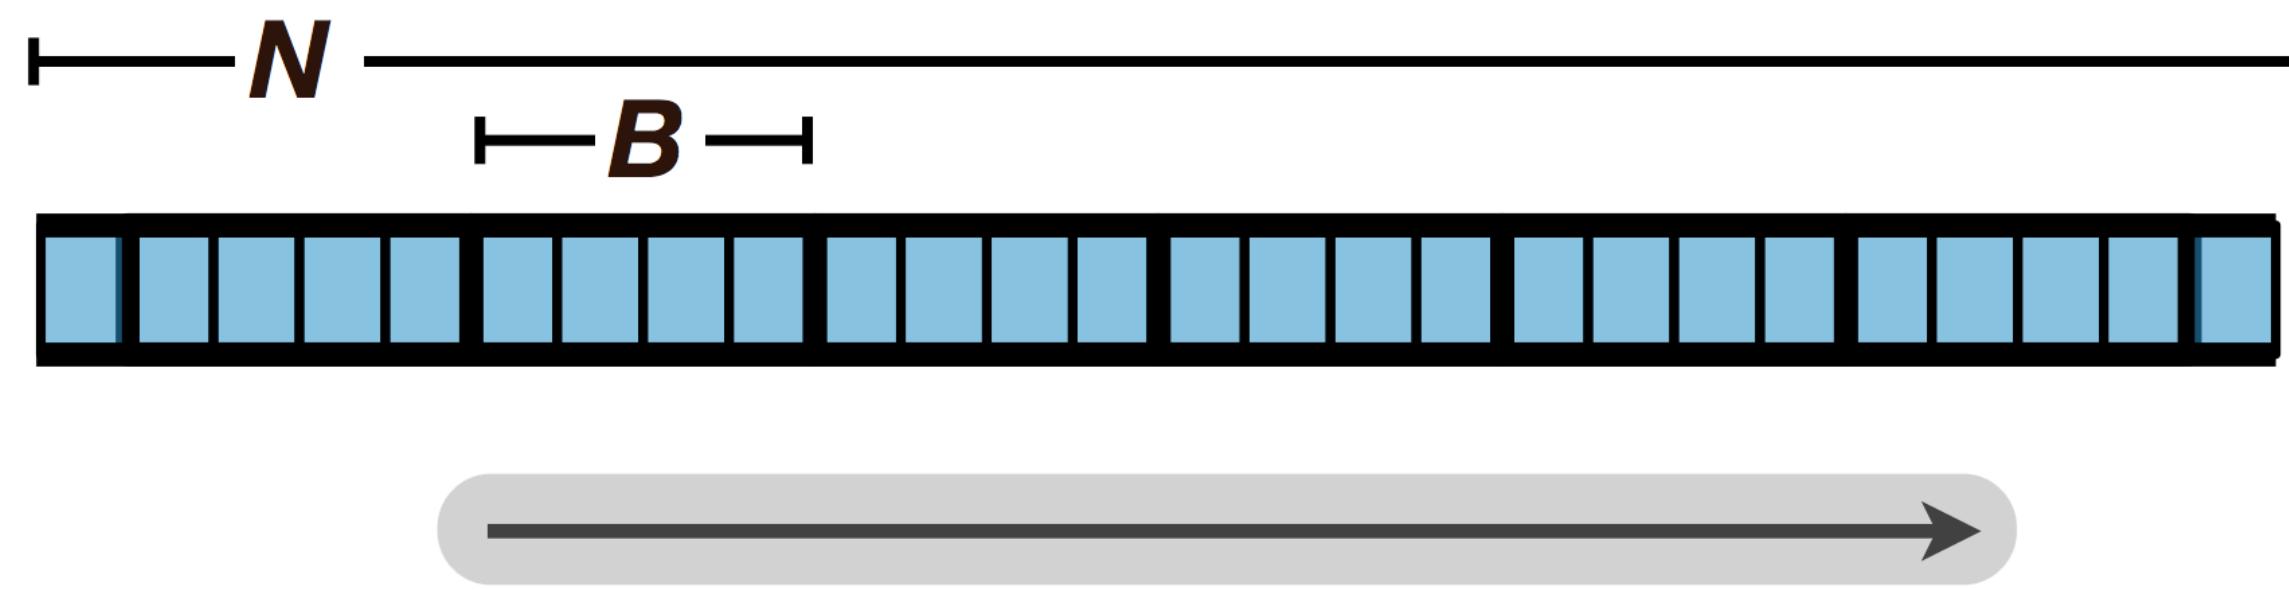
[Aggarwal & Vitter 1988]

- Models any two levels of memory hierarchy
- Memory of size M , disk of unlimited size
- Data is transferred in blocks of size B
- Algorithmic performance metric: number of block transfers (I/Os)
- Clean and useful model to predict cache-efficiency of algorithms
- Usually algorithms don't need to know M, B (cache oblivious)



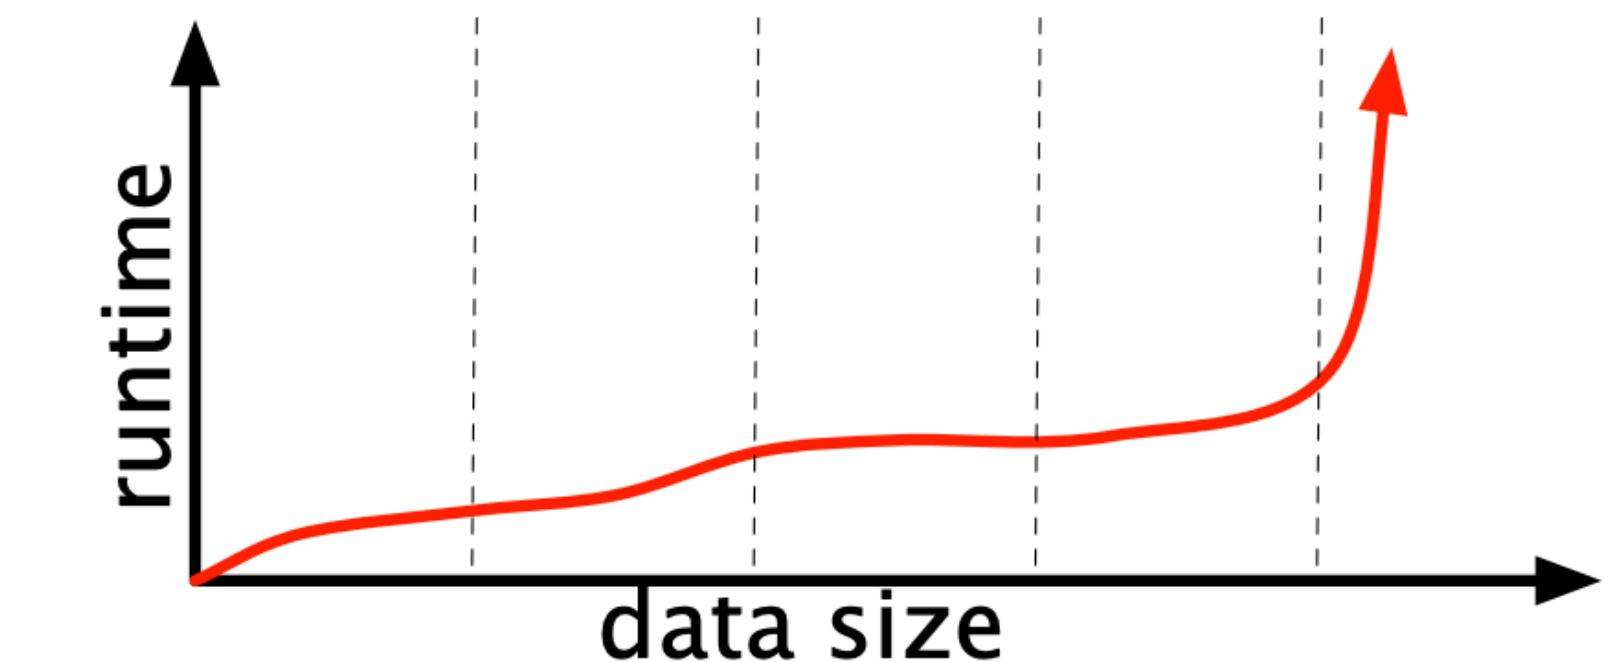
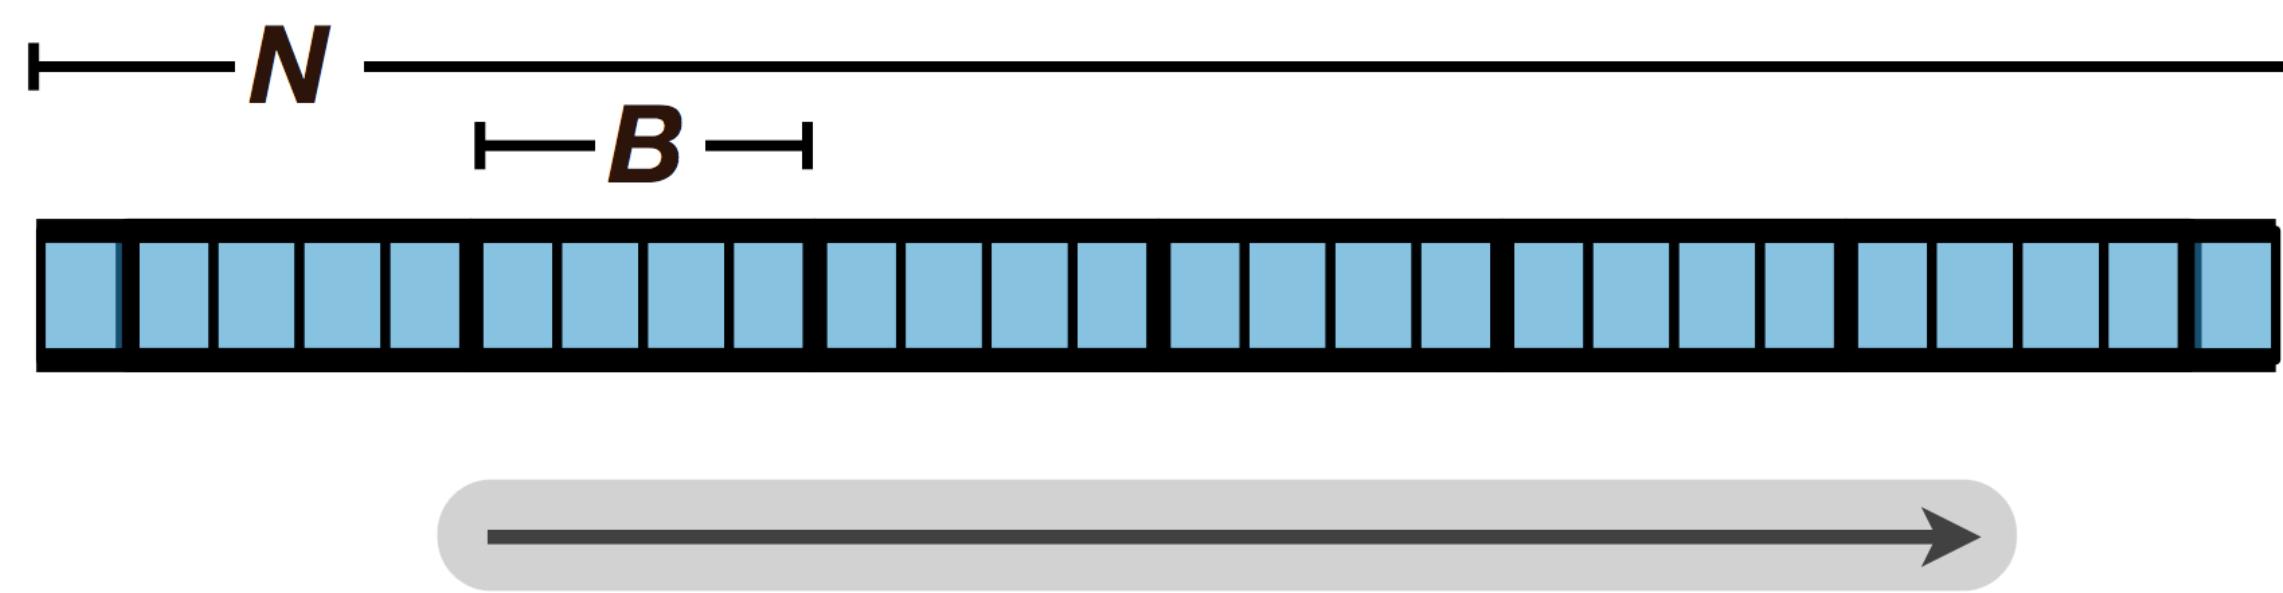
Sequential Versus Random Access

- Disk access usually 10^6 times more expensive than RAM
- Sequential access patterns are better than random
 - Due to cache misses regardless of where in the memory hierarchy
- Consider traveling a linked list
- Naive: $O(N)$ I/Os (each hop to a new block)
- Smart: $O(N/B)$ I/Os (if sequential nodes stored in a single block)



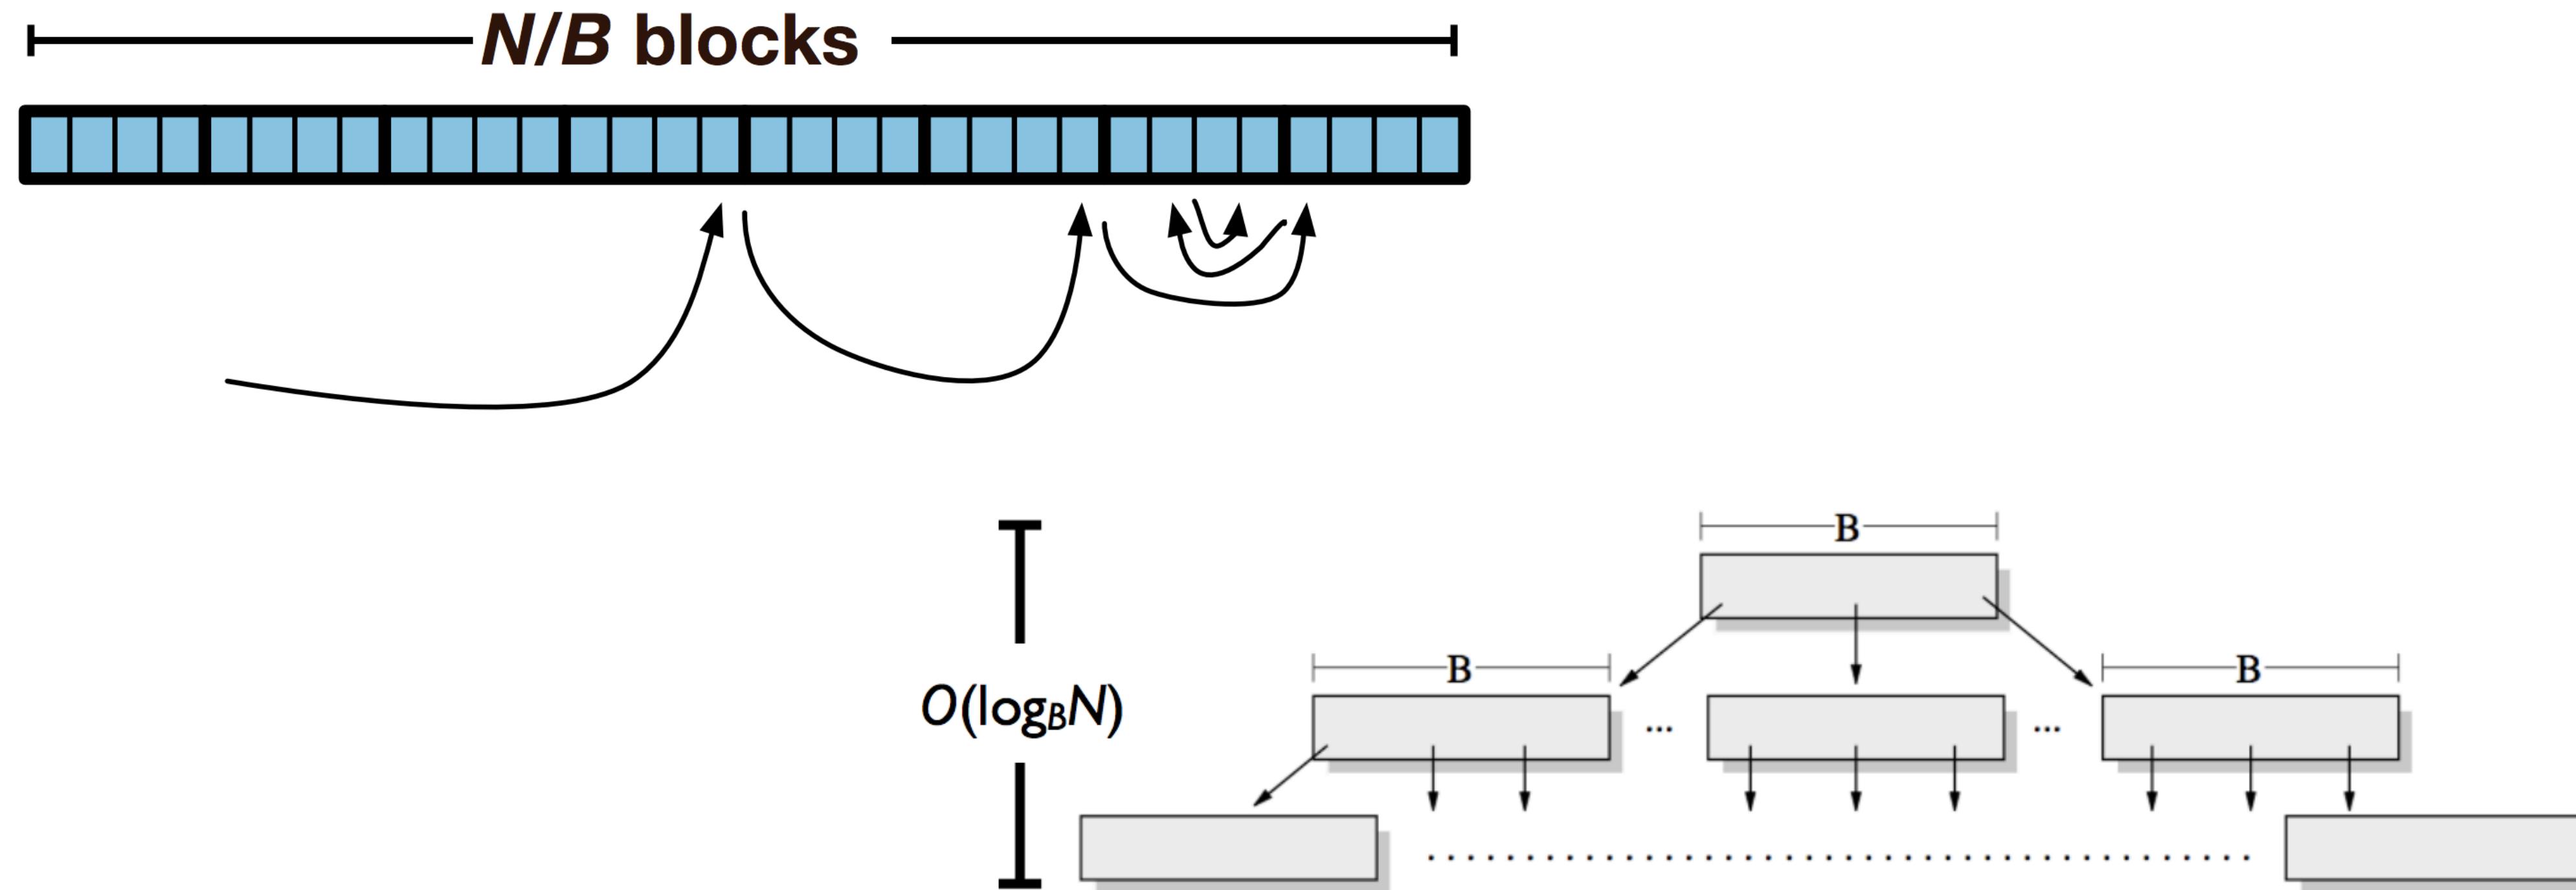
Difference Between N and N/B

- Consider $N = 256 \times 10^6$, $B = 8000$
 - With disks: 1ms disk access time, SSDs: ~25 microseconds
- N I/Os takes 256×10^3 sec = **71 hours** disks, ~**4 mins** for SSDs
- N/B I/Os takes $256/8$ sec = **32 sec**
- But, most problems developed in RAM model; why don't they thrash?
 - Under the hood OS stuff paging and prefetching
- Massive data and random access pattern can still hurt a lot



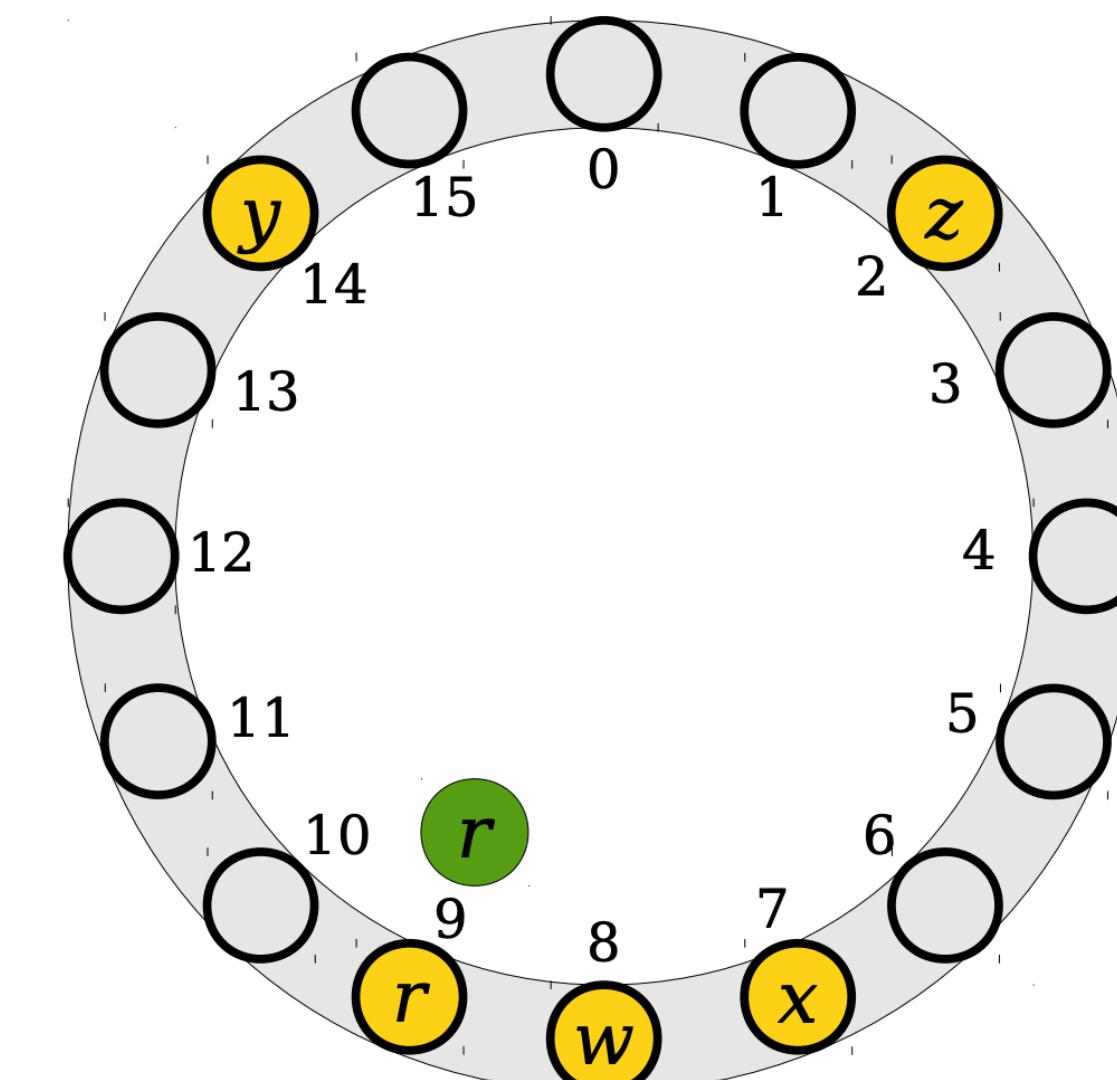
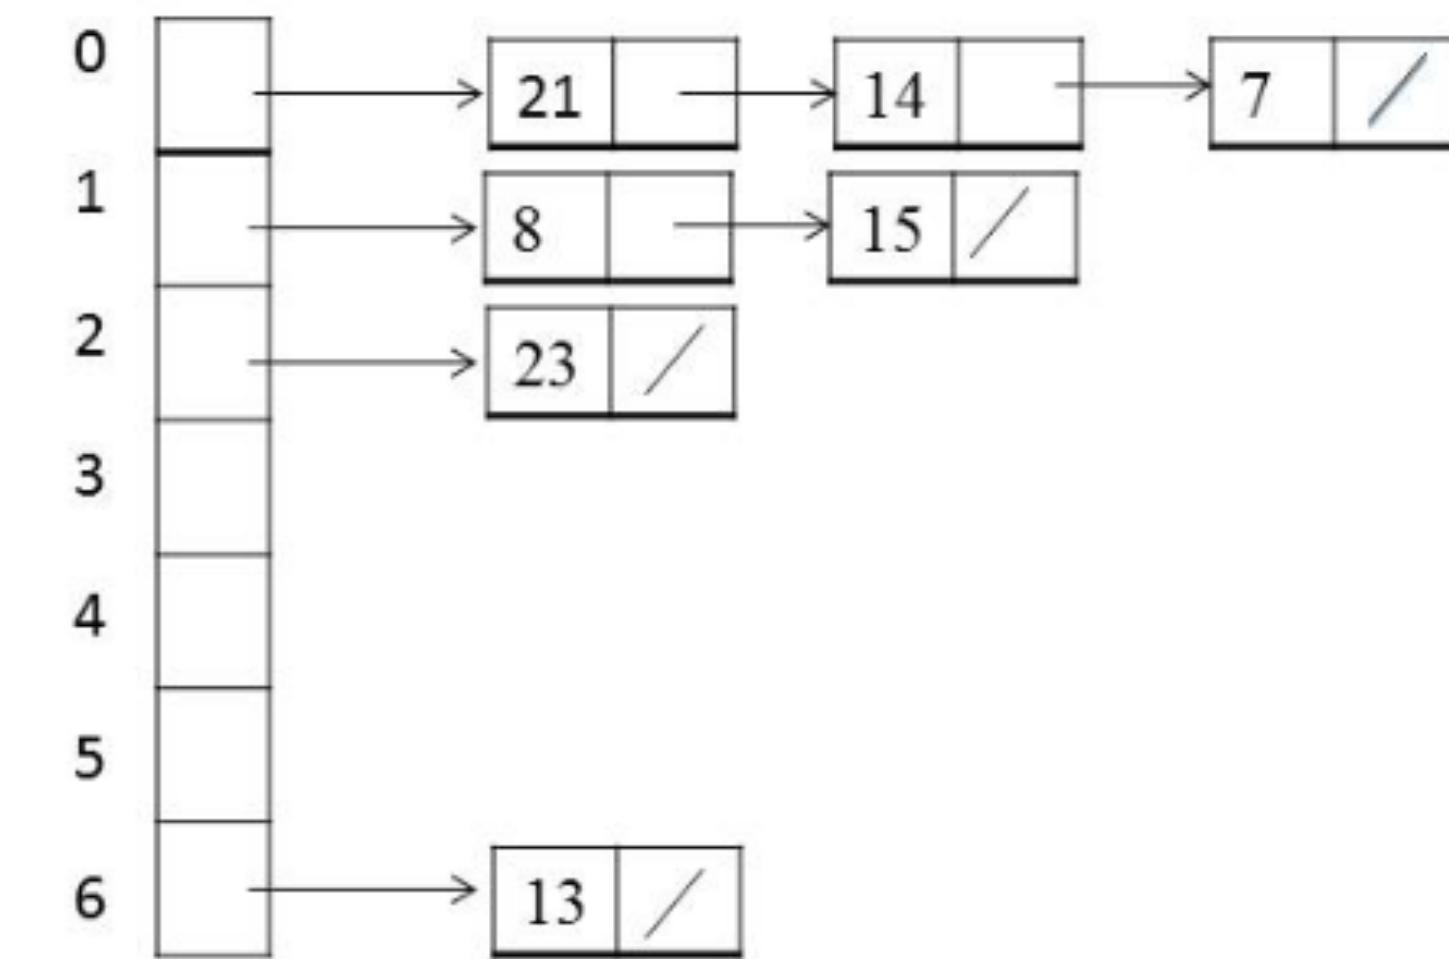
Examples

- **Standard binary search.** Cache-efficient or not?
 - Number of I/Os: $O(\log_2(N/B)) \approx O(\log_2 N)$
 - Much more I/O efficient to search using a B -tree (a B -ary) tree



Examples

- Hashing with **Chaining** vs **Linear Probing**
- Linear probing is much more cache-efficient
- In practice, linear probing is one of the fastest general-purpose hashing strategy
 - Low overhead
 - Excellent locality!
 - Great cache performance



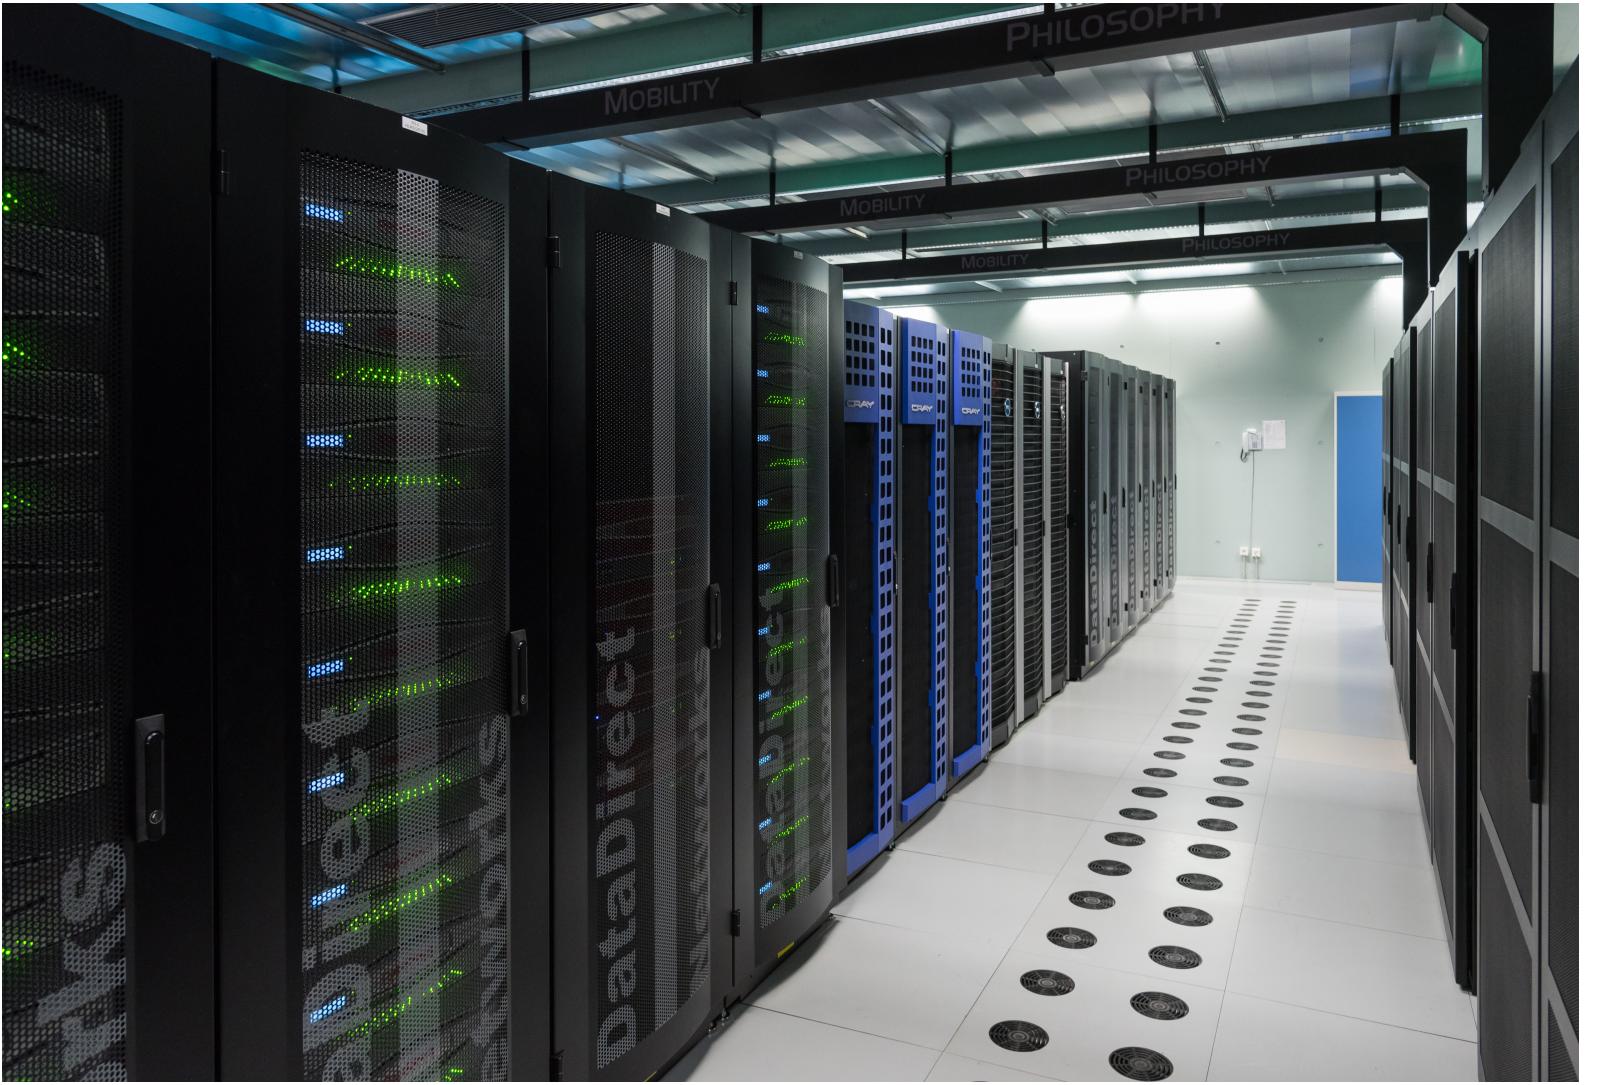
Operations Aren't the Same

- Addition and subtraction are fast, multiplication is fast
- Division and modulo are slow
- Integers are faster than floats (?)
- Can we model this in theory?
- Not really. Asymptotics ignore this intentionally
 - Going from $O(2^n)$ to $O(n)$, or even $O(n^2)$ to $O(n \log n)$, is more important
- Occasionally something like: $O(n \log n)$ additions, $O(n)$ divisions



Parallelism

- Modern computers almost always have multiple compute cores
- If we have p identical “processors” can we speed up our algorithms?
 - Maybe by a factor of p ?
 - Many models for algorithm analysis
 - PRAM is the classical model
 - MapReduce model: massive number of cores, want to minimize communication rounds
 - Many others



Take Aways

- Big Oh analysis we did is the first strong predictor for performance
- But there are other factors like cache efficiency, that determine how well an algorithm will perform in practice
- When applying these to practice eye towards:
 - Other costs (like cache efficiency)
 - Abilities of modern hardware (parallelism, GPUs)

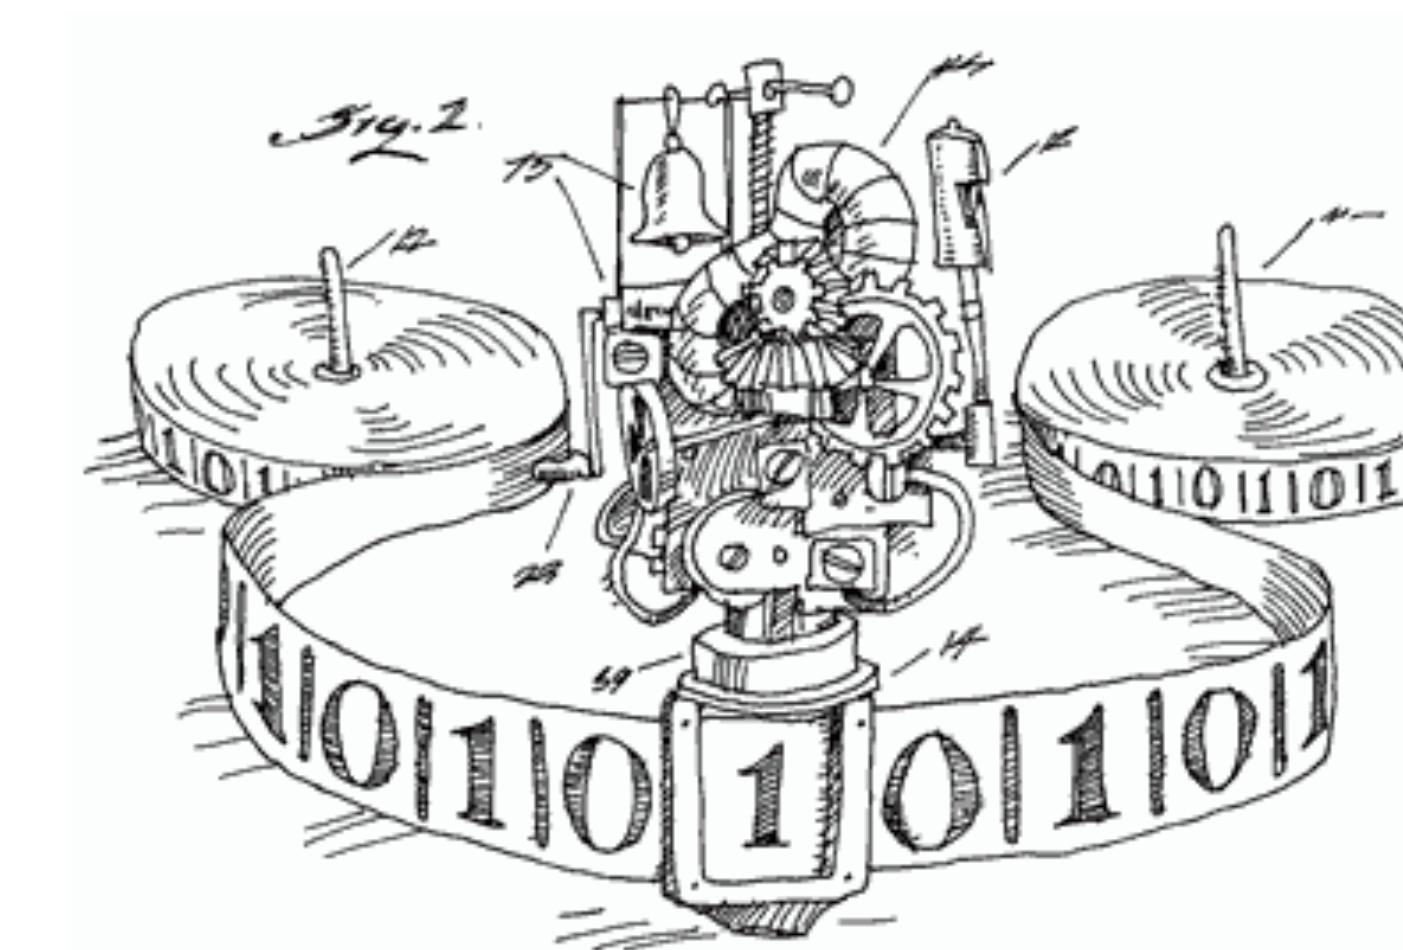
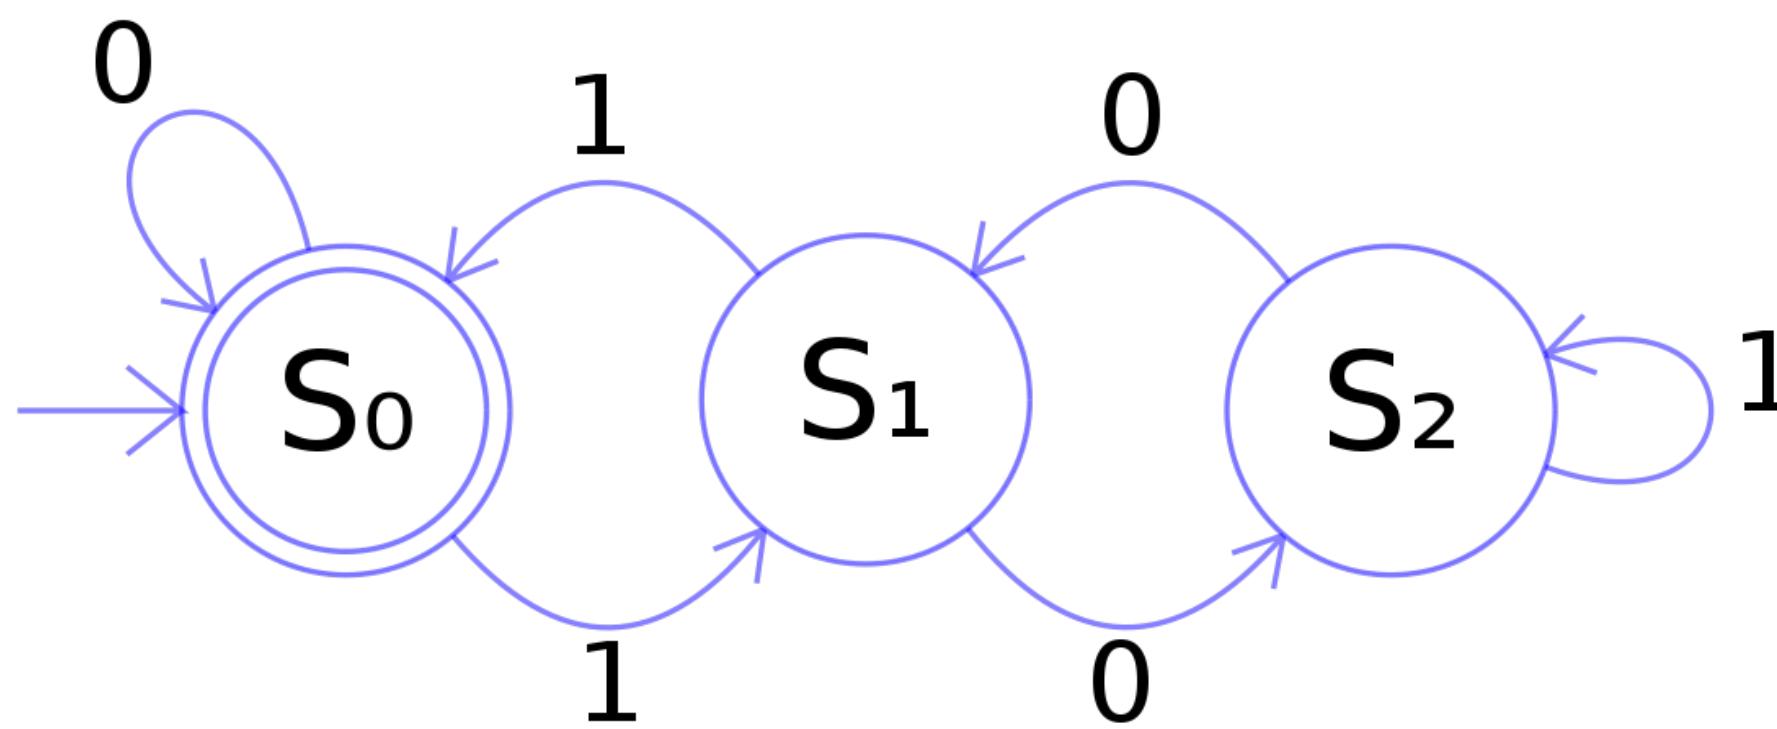
Beyond CSCI 256: Similar Courses

Beyond CS 256: CSCI 361: Theory of Computation

Required course if you are CS major

- Formal framework for investigating both the computability and complexity of problems
- Several models of computation including finite automata, regular languages, context-free grammars, and Turing machines
- These models provide a mathematical basis for the study of computability theory

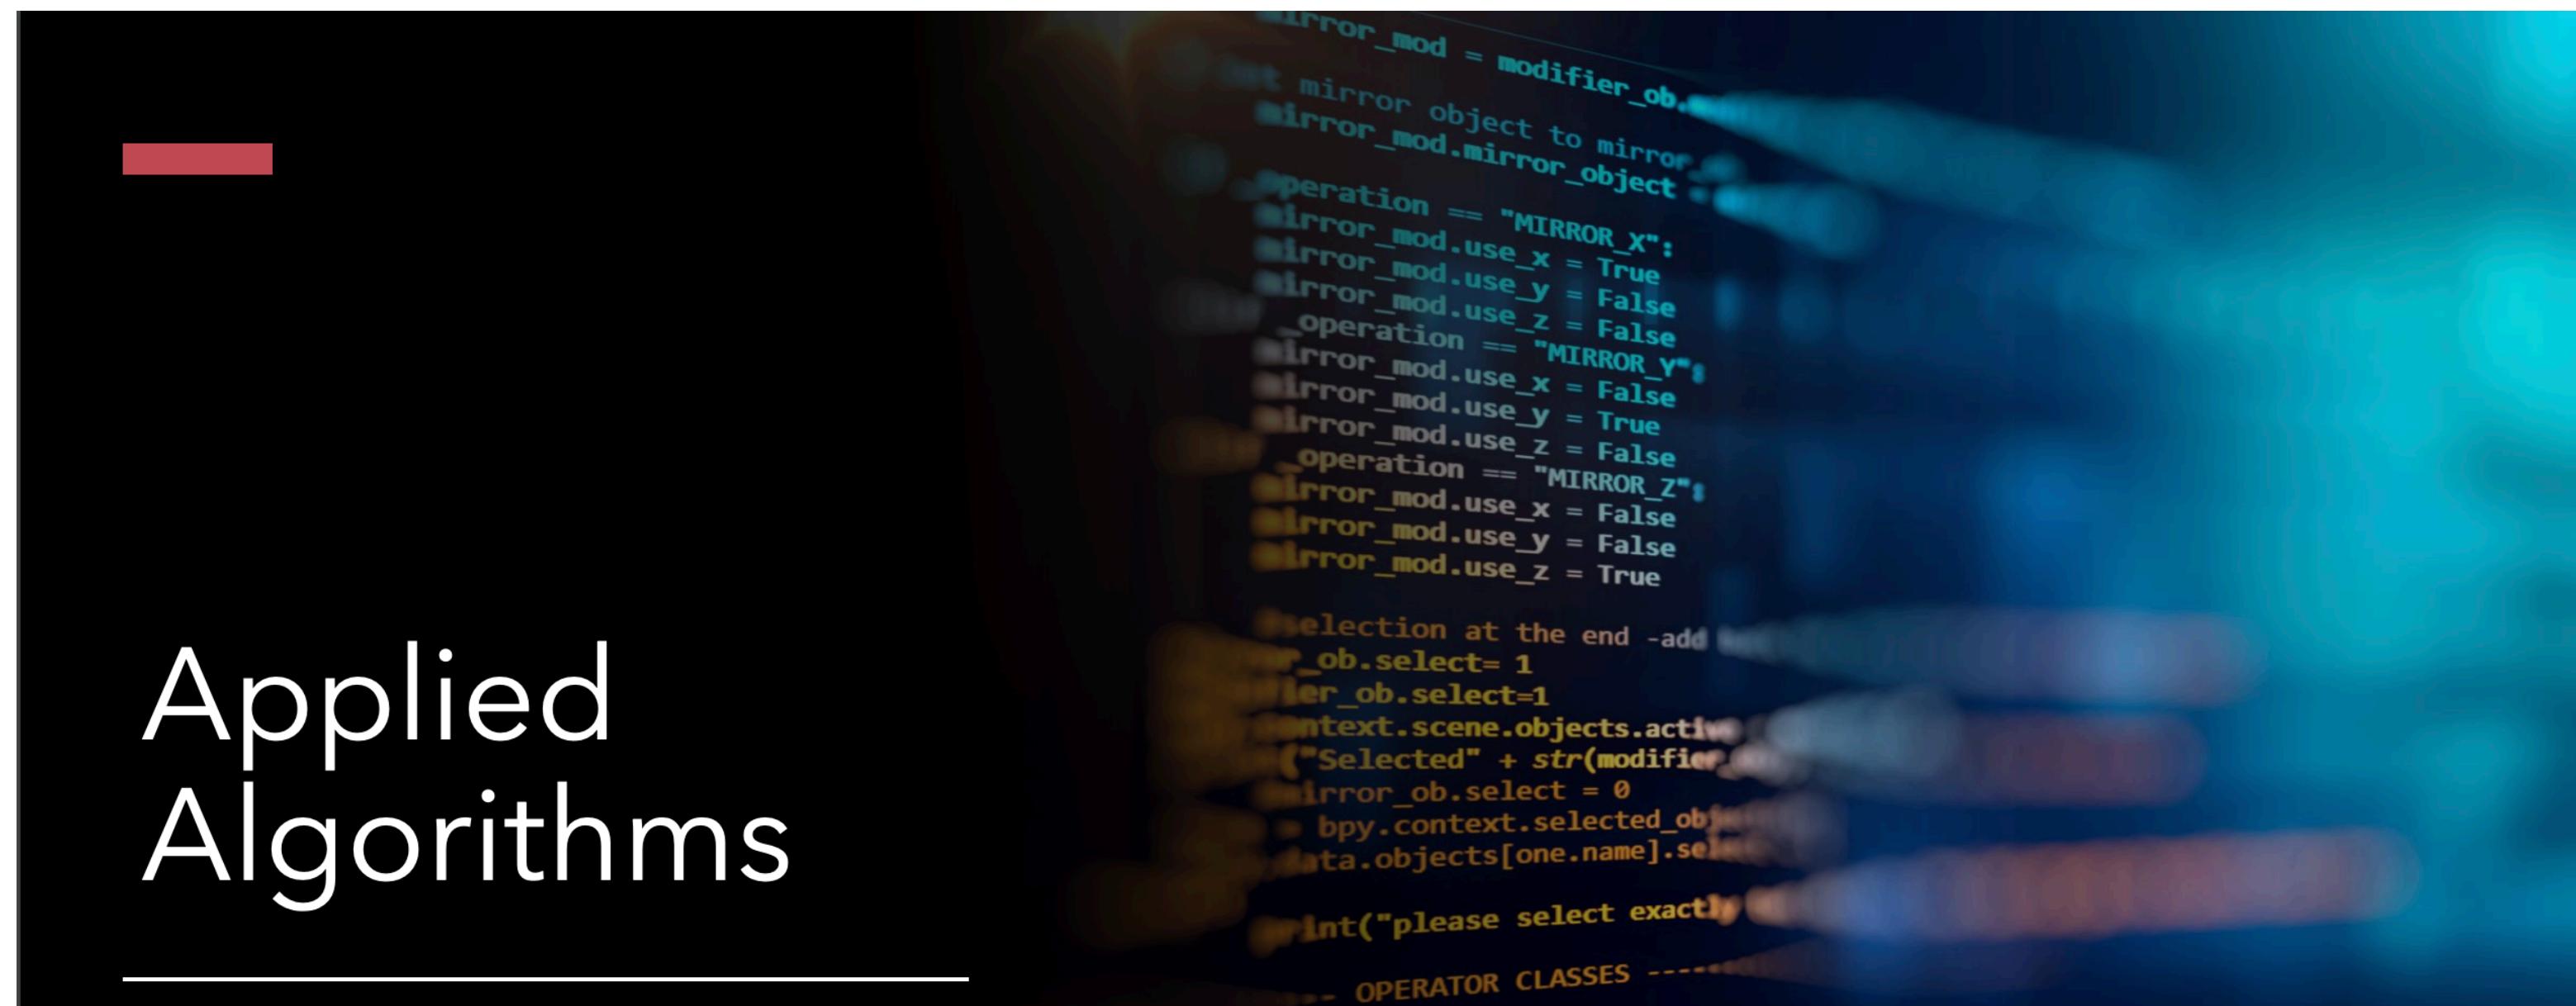
Prerequisite:
CS 256



Beyond CS 256: CSCI 358: Applied Algorithms

- Bridging the gap between theoretical running time and writing fast code in practice (taking advantage of modern hardware)
- Learn about the most useful tools in a coder's toolkits
- How to implement algorithms in an practically-efficient way
- Learn new algorithmic techniques not covered in 256

Prerequisites:
CS 256 and CS237

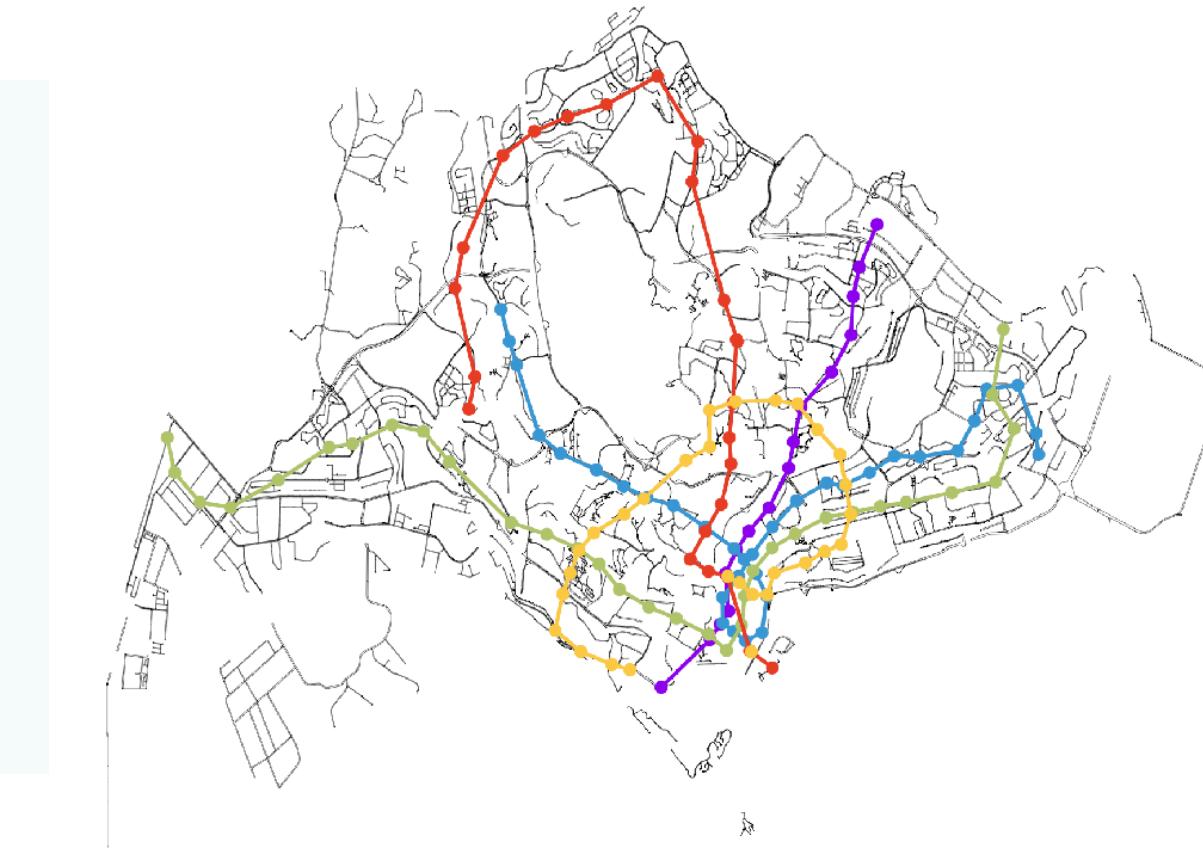


Beyond CS 256: CSCI 357: Algorithmic Game Theory

Topics in game theory/mechanism design from a computational perspective

- how to design algorithms that incentivize truthful behavior, that is, where the participants have no incentive to cheat?
- Overarching goal is to understand and analyze selfish behavior and whether it can or should influence system design

Prerequisite:
CS 256

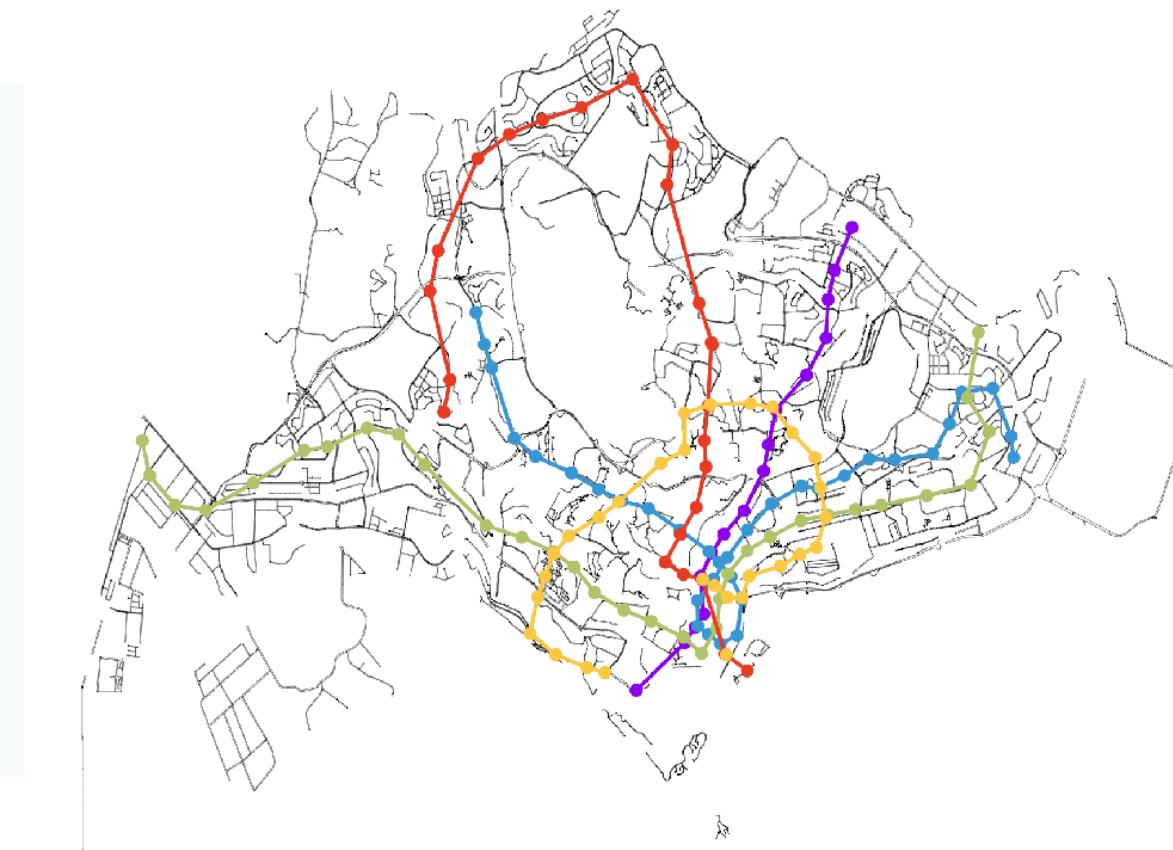


Source: <https://courses.csail.mit.edu/6.851/spring12/>

Beyond CS 256: CSCI 357: Algorithmic Game Theory

"The next time we bemoan people exploiting loopholes to subvert the intent of the rule makers, instead of asking 'What's wrong with these people?' let's instead ask, 'What's wrong with the rules?' and then adopt a scientifically principled approach to fixing them."

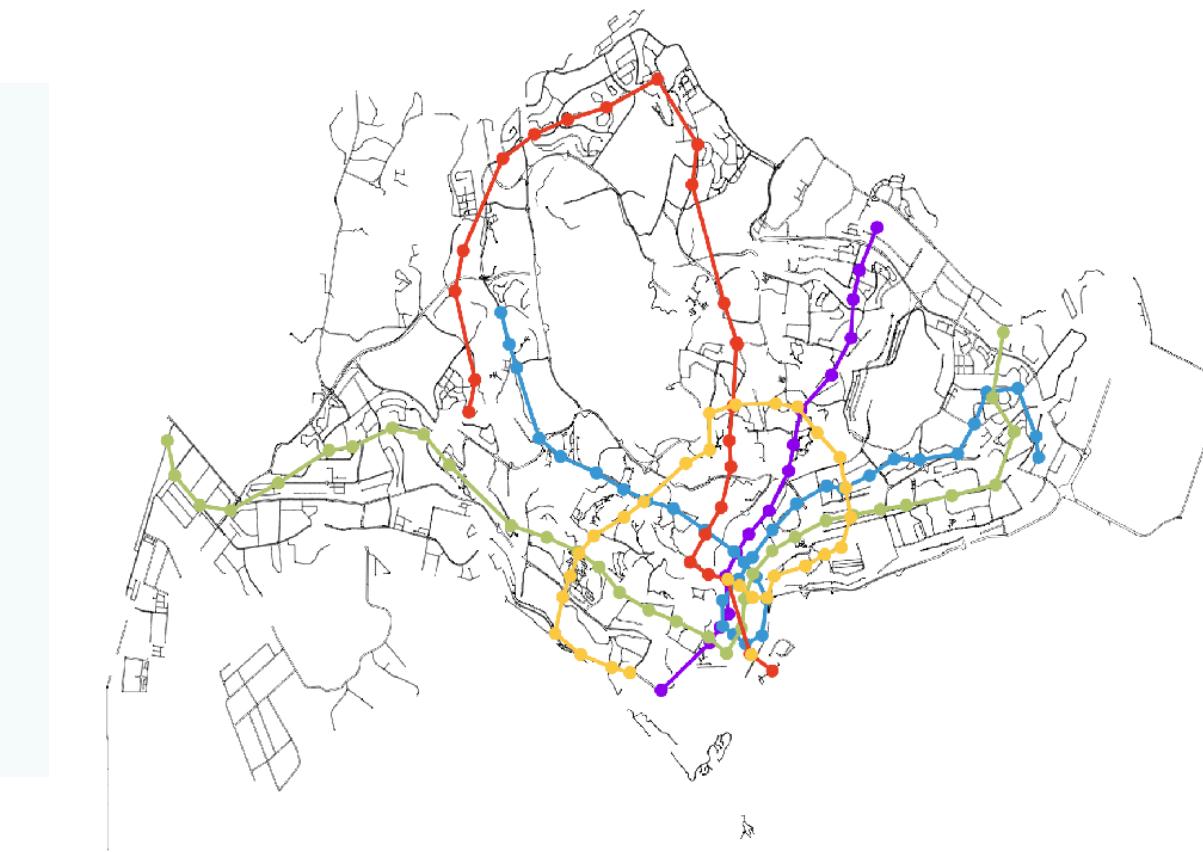
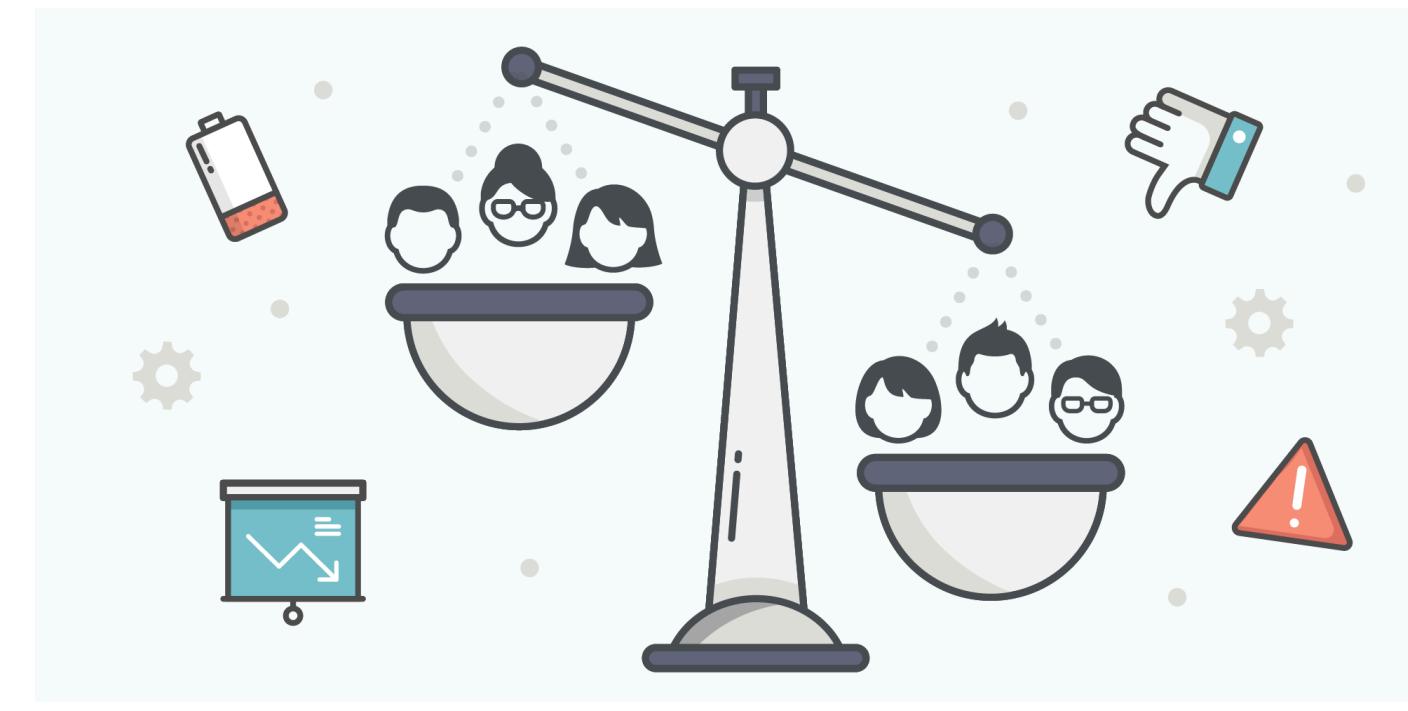
— Hartline and Kleinberg



Beyond CS 256: CSCI 357: Algorithmic Game Theory

Examples of topics:

- auction design, efficiency of equilibria, network games
- two-sided markets, crowdsourcing markets
- incentives in computing applications such as file sharing, cryptocurrencies, etc
- computational social choice: selfish voting, resource allocation, etc.

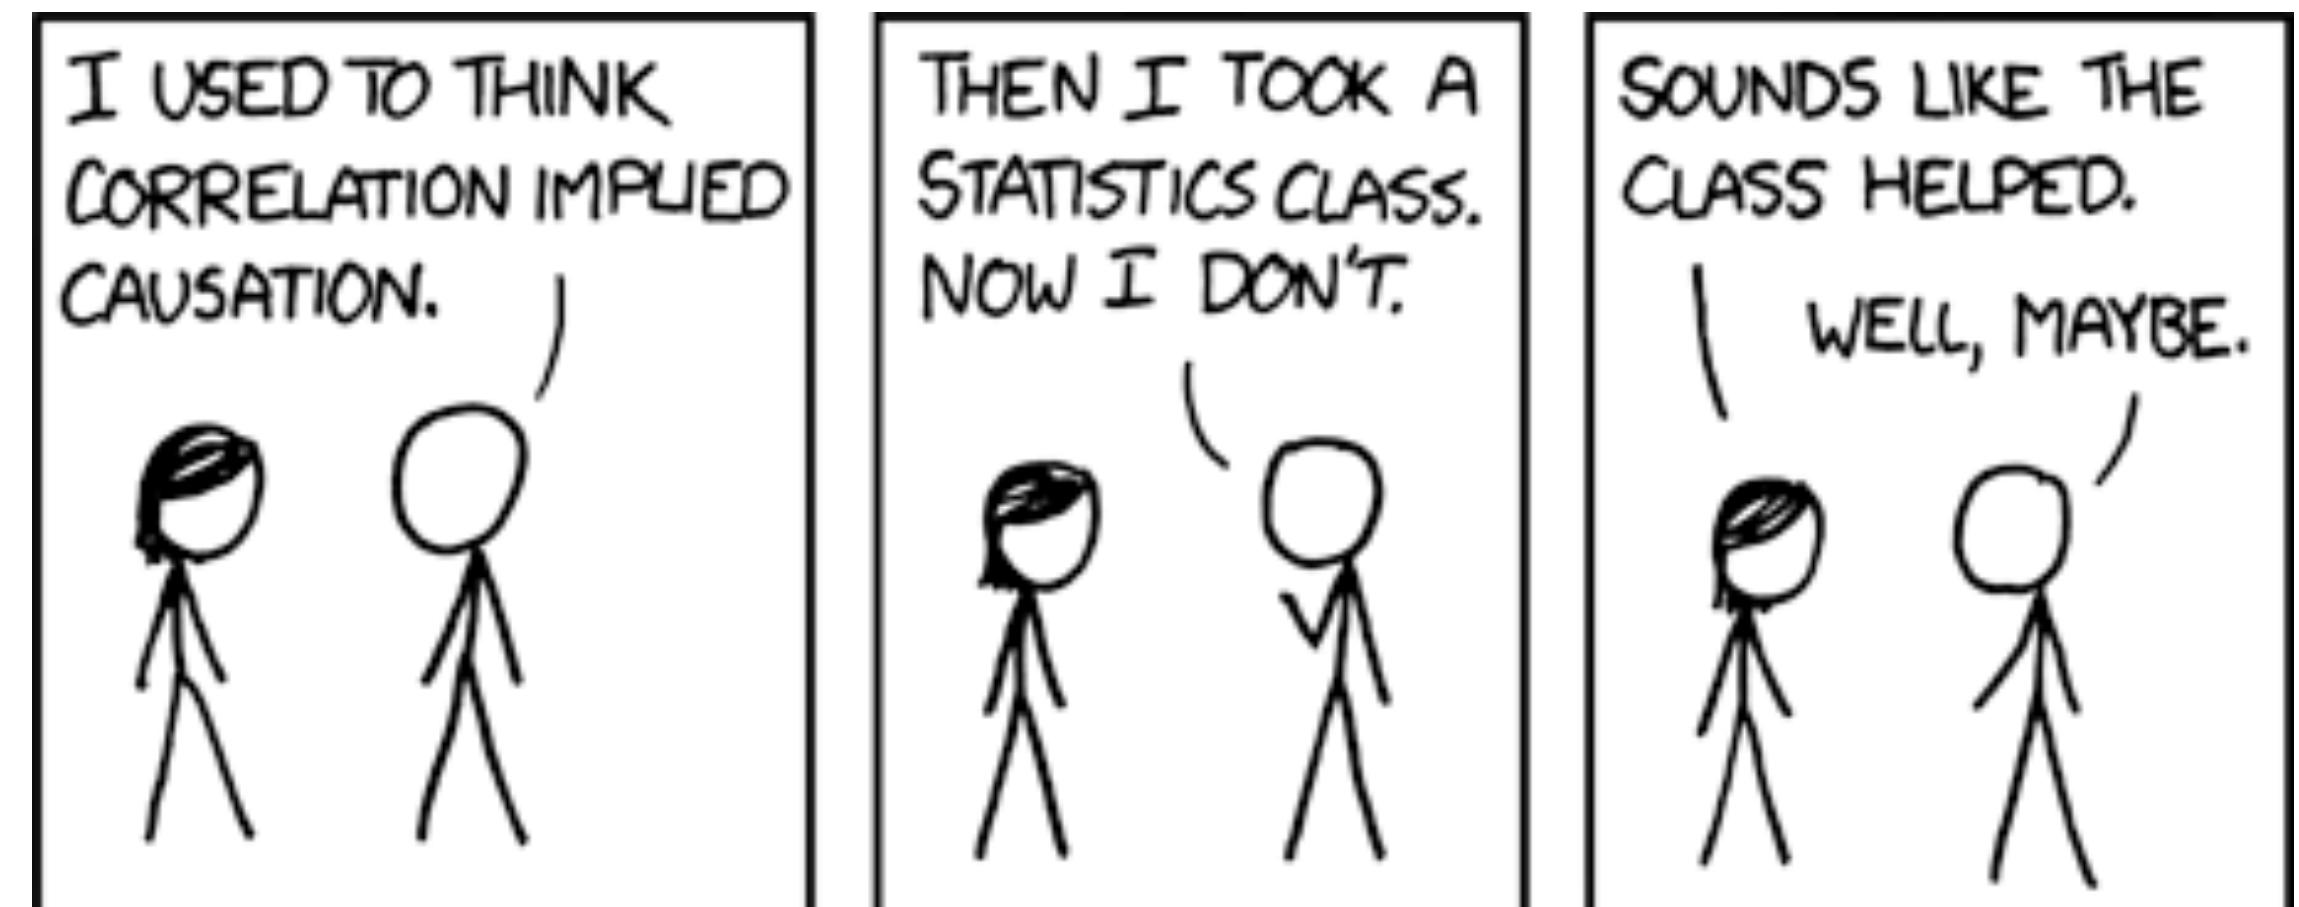


Beyond CS 256: CSCI 379: Causal Inference

Does X cause Y? If so, how? And what is the strength of this causal relation? Teasing apart causation from association in messy real-world data; Topics in causal inference including

- causal graphical models and unsupervised learning of the structure of these models
- expression of causal quantities as functions of observed data
- and robust/efficient estimation of these quantities using statistical and machine learning methods

Prerequisite:
CSCI 136, and either CSCI
256 or STAT 201/202



Course Evaluations

About Course Evaluations

Purpose.

- Your feedback will help **improve this course** for other students taking it in the future
- Help shape the CS curriculum

Who sees the evaluation results:

- Senior faculty in Dept & College
- Maybe Maud
- Me (only *after* I put in grades)
- Remember that the evaluations are **confidential**

Course Evaluations: the Script

- You may skip questions that you don't wish to answer, and there is no penalty for choosing not to participate.
- All of your answers are confidential and I will only receive a report on your responses after I have submitted all grades for this course.
- While evaluations are open, I will receive information on how many students have filled out the evaluations, but I won't know which of you have and haven't completed them.
- I won't know which responses are associated with which student unless you identify yourself in the comments.

Blue Sheets!

No longer blue or a sheet

"Due to the pandemic, physical blue comment sheets will not be distributed. Instead, students will see a digital equivalent at the conclusion of the evaluations."

Purpose. Comments or feedback directed only to me
(no one else reads these but me)

To access the online evaluations, log into **Glow** (glow.williams.edu) using your regular Williams username and password (the same ones you use for your Williams email account). On your Glow dashboard you'll see a course called "**Course Evaluations**." Click on this and then follow the instructions you see on the screen. If you have trouble finding the evaluation, you can ask a neighbor for help or reach out to ir@williams.edu.

All the Best
at

Williams and Beyond

pls fill out course evals

Acknowledgments

- Some of the material in these slides are taken from
 - Kleinberg Tardos Slides by Kevin Wayne (<https://www.cs.princeton.edu/~wayne/kleinberg-tardos/pdf/04GreedyAlgorithmsI.pdf>)
 - Jeff Erickson's Algorithms Book (<http://jeffe.cs.illinois.edu/teaching/algorithms/book/Algorithms-JeffE.pdf>)